

AI that computes with time and is already winning

A new family of learning machines, built for a world that happens as events, beats the best published models on public benchmarks with a fraction of their size.

5 of 5

datasets won on the public EasyTPP event-prediction leaderboard

0.639

sepsis early-warning AUPRC vs 0.583 for the best published model $\uparrow$

15×


smaller and cheaper than the state-of-the-art model on Retweet, while more accurate

34 / 35

independent runs ahead of the best published result on the seven benchmarks won

In one sentence. Sleeping Machines is building a general-purpose computing substrate for machine intelligence, one that learns through time and sparse events, and its first public results are wins against the strongest published models, achieved by small, efficient models under strict, pre-registered test protocols.

Every independent run ahead of the best published result




All results in this brief are completed runs on official benchmark data, scored once on sealed test sets. Each benchmark's unit, and whether higher or lower is better, is stated beside every comparison ($\uparrow$ higher is better, $\downarrow$ lower is better). Sources: experiments/B1_EASYTPP.md, experiments/B2_IRREGULAR_TS.md, report/I_SCIENCE.md.

One substrate for machine intelligence

Today's AI runs on a clock: every input is chopped into fixed steps, and the whole model works at every step whether or not anything happened. The real world is not like that. Purchases, heartbeats, sensor readings, messages, machine logs and words arrive as *events*, irregularly, with meaning in their timing and in their silences.

Sleeping Machines builds learning machines in which **time itself does part of the computation**. Our goal is one trainable substrate that spans:



What makes it different, in plain terms:

- **Memories that live through silence.** Each memory keeps evolving between events, so “how long since” and “what happened in between” are information, not missing data.
- **A race of clocks decides what happens next.** Possible next events race against each other; the one that fires first wins. This gives exact predictions of both *when* and *what*, and the same mathematics reproduces the attention mechanism inside today's large language models.
- **Only the relevant parts wake up.** Updates are sparse and addressed: the machine can hold far more capacity than it uses for any single event, which is where the name comes from.
- **Learning from roads not taken.** The model is also credited for the alternatives that lost the race, so it learns which routes to take.

Why this matters commercially. Most of the world's valuable data is timestamped events: transactions, clinical records, telemetry, security and IT logs, user behaviour. Models that understand time natively can be more accurate on that data and much cheaper to run, on CPUs and edge devices as well as in datacentres. The same principles point to hardware without a global clock, where idle capacity costs almost nothing.

The benchmarks, in plain language


We compete on established public benchmarks, using their official data splits and scoring rules, against the best results other teams have published. The headline event benchmark is **EasyTPP** (ICLR 2024), the standard test suite for predicting *when* the next event happens and *what* it will be. Its current leader is **S2P2** (NeurIPS 2025, GE HealthCare and UC Irvine), a modern state-space model.

Benchmark	What the data is	Why it matters	Score (unit, direction)
Taobao	Clicks and purchases of users on Alibaba's Taobao marketplace, by product category	E-commerce: anticipating what a shopper does next and when	log-likelihood, nats per event ↑
Taxi	New York taxi pick-ups and drop-offs across city districts	Mobility and logistics: demand and fleet planning	log-likelihood, nats per event ↑
StackOverflow	Badges earned by users of the StackOverflow Q&A site	User engagement and lifecycle prediction	log-likelihood, nats per event ↑
Retweet	Retweet cascades on Twitter, by type of user retweeting	Social media: how information spreads, and how fast	log-likelihood, nats per event ↑
Amazon	Product reviews by Amazon customers, by category	Retail behaviour across product lines	log-likelihood, nats per event ↑
P19 sepsis	PhysioNet 2019: 38,803 intensive-care stays, 34 vital signs and lab values measured at irregular times	Early warning of sepsis, a leading cause of hospital deaths; every hour of earlier treatment matters	AUPRC and AUROC, 0 to 1 ↑
PAM	PAMAP2 physical-activity monitoring: 5,333 recordings from wearable inertial sensors (17 channels), 8 activities	Wearables and health monitoring: recognising what a person is doing from body-worn sensors	accuracy and F1, 0 to 1 ↑
P12 mortality	PhysioNet 2012: 11,988 intensive-care stays, 36 irregular measurements	Predicting in-hospital mortality from the first 48 hours	AUROC and AUPRC, 0 to 1 ↑

Reading the scores.

- **Log-likelihood (nats per event), higher is better.** It measures how much probability the model gave to what actually happened, both its timing and its type. A gain of 0.08 nats per event means the model assigns about 8% more probability to every real event, on average, compounding over every event in a sequence. Values can be negative; less negative is better.
- **AUPRC, higher is better (0 to 1).** How well the model ranks the rare patients who develop sepsis (4% of stays) ahead of everyone else, with emphasis on raising true alarms without false ones. It is the clinically relevant metric when the condition is rare. **AUROC** (0 to 1, higher is better) measures overall ranking.
- **Size and cost, lower is better.** Parameters measure how big a model is; multiply-adds per event measure the work needed for each prediction, which drives running cost and energy.

Ahead of the best published models




Each blue dot is one independent run (EasyTPP: five training seeds; sepsis and wearables: the five official data splits), scored once on the sealed test set. The grey line is the best published result. Higher is better in every panel: log-likelihood in nats per event for the EasyTPP datasets, AUPRC (0 to 1) for sepsis, accuracy (0 to 1) for wearables.

Benchmark	Best published	Sleeping Machines	Margin	Our size and compute vs the leading model
Taobao nats/event ↑	1.318 (IFTTP)	1.399 ± 0.003	+0.081	0.92× S2P2's compute per event; best time and type predictions
Taxi nats/event ↑	0.522 (S2P2)	0.525 ± 0.001	+0.003	1/12 of S2P2's size and compute ; reproduced independently on separate hardware
StackOverflow nats/event ↑	-2.163 (S2P2)	-2.144 ± 0.004	+0.019	1.26× compute; a model of S2P2's size also wins (-2.153)
Retweet nats/event ↑	-6.348 (NHP)	-6.326 ± 0.001	+0.022	1/15 of S2P2's size and compute ; S2P2 scores -6.365
Amazon nats/event ↑	0.781 (S2P2)	0.803 ± 0.001	+0.022	0.29× S2P2's compute , 3.6× smaller; best time and type predictions
Wearables, PAM accuracy ↑	0.975 (MTM)	0.978 ± 0.007	+0.003	19× smaller than MTM (46,316 vs 873K parameters); F1 0.980 vs 0.976; 4 of 5 splits ahead
Sepsis, P19 AUPRC ↑	0.583 (MTM)	0.639 ± 0.039	+0.056	62,681 parameters (MTM reports 873K for its PAM configuration; its P19 size is not reported); AUROC 0.916 vs 0.903 ↑; not just summary statistics: without its temporal memory the same network drops to 0.572 on every split, and it beats trees on the same statistics on every split


± is the standard deviation across the five runs. Compute is multiply-adds per event, counted from S2P2's released code at its published configuration.

Why these wins matter. EasyTPP is the standard benchmark of the field, and S2P2 is a 2025 model from a strong industrial and academic team. We lead on all five of its datasets, and on two of them with a model one-twelfth to one-fifteenth its size. The same family also leads a public clinical benchmark and a public wearable-sensor benchmark. These are not one lucky run: 34 of 35 independent runs are ahead.

Winning while small



Learned parameters, in thousands; fewer means a smaller, cheaper model (lower is better). Grey: S2P2, the leading published state-space model on EasyTPP, at its published configuration for each dataset; for PAM, MTM (2025) at the parameter count it reports for PAM. Every Sleeping Machines model shown is also more accurate than the model beside it. *StackOverflow: the model matched to S2P2's size, which also wins.




Work per prediction, in thousands of multiply-adds per event (lower is cheaper). "ours: 1/15" means our model needs one fifteenth of S2P2's work.

What size buys.

- **Running cost and energy** scale with the work per prediction.
- **Deployment anywhere:** models of 20,000 to 60,000 parameters run on ordinary CPUs, phones and embedded devices, close to the data.
- **Privacy:** clinical and industrial data can stay on site.
- **Headroom:** accuracy at small size means the architecture, not brute scale, is doing the work.

One family across different kinds of data

S2P2 is shown on one kind of task. Our family wins on two very different ones with shared machinery: the sepsis classifier reuses *the same temporal memory layer, same code and same size*, as the event-prediction model.



One shared core across four domains (development data, first seed). We trained a *single* temporal memory of 14,080 parameters on taxi rides, online shopping, Q&A activity and product reviews at the same time; only the small input and output layers differ per dataset. It matched or beat the four separately trained models on three of the four (log-likelihood, nats per event, higher is better: Taxi 0.490 vs 0.487, Taobao 1.290 vs 1.284, StackOverflow -2.186 vs -2.177), with about 40% fewer parameters in total. Mobility, shopping and online-community timing share one learned temporal dynamics. Amazon trailed and was still improving when training stopped; per-dataset training budgets are the next test.

Why generality matters for value. A company with one specialised model has one product. A company with a core that transfers across data types can serve many markets from one technology, and every new domain makes the core stronger.

Each layer of evidence supports the next

Next: scale and transfer. Large-scale language on GPUs, skill transfer between data types, clockless hardware.


5. It generalises. One shared core across four event domains; one memory layer for events and clinical records.

4. It wins public benchmarks. All 5 EasyTPP datasets, P19 sepsis and PAM wearables, at up to 15× smaller size, every run ahead.

3. It learns language efficiently. On 10M characters it beats a Transformer: 1.888 vs 1.908 bits per character ↓ with 0.40× its training compute. At 90M characters it beats a tuned Transformer-256×4 (1.783 vs 1.811 ↓) at 1.06× its compute.

2. The mechanisms work. Event order learned from 2,000 examples seen once: 99.7% accuracy vs 33–41% for Transformers ↑. Timing-only discrimination: 95.3% where any order-only model is capped at 50% ↑. Learning from roads not taken improves language 2.507 → 2.371 bits per character ↓ for 0.3% extra training work; more idle capacity improves it further at the same active work.

1. Mathematics: it contains today's leading architectures. A race of clocks reproduces the attention of Transformers *exactly*, and Mamba-style selective state-space models are exact members of the family (verified numerically to 4×10^{-13}). Everything those models can represent, ours can, plus native time and sparsity.



Why the steps add up. The family is not a narrow trick for one benchmark. Its mathematics contains the architectures that power today's AI; its mechanisms solve temporal problems those architectures fail at; it already learns language competitively at small scale; and it wins public benchmarks in two domains with one shared memory design. Scale is the remaining work, and each layer beneath it is already demonstrated.

Hardware. In our cost model, the native language model moves 5.8× fewer bytes per character than a comparable Transformer at better quality (modelled, not yet measured on silicon; a funded milestone).

How we know the results are real

🔒 Sealed tests, rules fixed in advance. Every win uses the benchmark's official data splits. The reporting rule is written down before the test set is touched, and each model is scored on the test set exactly once.

🔄 Five independent runs. Every reported result is the average of five runs (seeds or official splits), with its spread. On every benchmark won, all five runs are ahead of the best published result.

🖨️ Reproduced on separate hardware. The Taxi win was reproduced independently on a separate desktop computer with separately downloaded data: 0.5252 vs 0.5250 nats per event ↑, every run above S2P2.

✅ The benchmark's own scorer agrees. Re-scoring our models with EasyTPP's own estimation method gives the same results within its sampling noise.

🔍 We found and closed a scoring pitfall. Timestamps in these datasets are rounded (for example to whole seconds). Flexible models can exploit the rounding for large, fake gains. We built an audit that removes the rounding; we withdrew our own early results that failed it, and every reported win passes it.

📊 Losses reported as plainly as wins. Where we are behind, we say so with the numbers and the next test (next page). Comparisons always use the same units and protocol as the published result.

Where we are not ahead yet, and the next test

Front	Current result	Next test
P12 mortality	AUROC 0.871 vs 0.880 ↑ (behind); AUPRC level, 0.585 vs 0.586 ↑	Ensemble of five runs, still far smaller than the leader (in progress)
Tokenised language (GPT-2 tokens, FineWeb)	Ahead of the classic count model: 6.009 vs 6.537 nats per token ↓ at 1M training tokens (3 runs); at 4M tokens 5.521 vs 6.100, so the lead grows with data; recall across irregular gaps solved (97.5%)	Further data scaling (16M tokens); a published small Transformer under the same protocol
Large-scale language	At 90M characters a tuned Transformer-192×4 leads: 1.704 vs our 1.783 bits per character ↓	Race attention at scale on GPU
Small-scale language	Tuned small LSTMs lead by 0.04–0.06 bits per character ↓ at 10M characters	Same programme as above
Real-world tables	Decision trees lead on the banknote dataset	Learned comparison discovery on real tables

WHAT COMES NEXT

From proof to platform

- **Publish** the complete event-leaderboard results, after the IP filing decision.
- **Scale language**: race attention against competent Transformers at scale on GPUs.
- **Transfer**: show that skill learned on one kind of data helps on another.
- **First commercial vertical**: event-native models for interleaved, timestamped logs (industry, IoT, IT operations, security, finance) on existing CPUs and GPUs, the shortest path to revenue.
- **Hardware**: feasibility on FPGA and neuromorphic silicon for clockless, learning-capable event processors.

The bottom line. A new model family, designed from first principles around time and sparse events, is already ahead of the best published models on seven public benchmarks across three kinds of data, including every dataset of the field's standard event benchmark, with models up to 15 times smaller. Its mathematics contains today's dominant architectures. The remaining work is scale, transfer and hardware, and each has a concrete, measurable next step.