

Sleeping Machines — investor pitch

A learning architecture and computing substrate for intelligence everywhere.

We are building AI that computes through time and sparse events, learns which work to perform, and maintains persistent context. The ambition is capable, adaptive intelligence from datacenters to mobile devices and robots, with models, execution software and eventually clockless hardware designed together.

Our core combines learned temporal races, addressed memory, separate keys and values, deep content-bearing messages and counterfactual credit to routes that did not win. This creates a path to more useful capacity than selected work, and to learning near the point of use. The complete construction and its quality/resource consequences are the proposed differentiation.

The research already has tangible positive results. On the completed single-seed 10M-character, one-pass text8 comparison, a small credited native model beats the saved Transformer control with about 15.4 times less estimated fitting work. A larger model slightly beats the saved LSTM's quality but uses more work. A separate causal CPU pilot improves predictions through full-backbone online adaptation. Against baselines retuned at our exact compute budgets (10 runs, selected by validation), the native model beats every tuned Transformer but loses to tuned small LSTMs by 0.04–0.06 bpc. Closing that local-modelling gap is the current language milestone. These are early scoped results, not modern frontier or hardware claims. The current report preserves their controls and limitations.

The first commercial target is useful quality at lower total datacenter cost. The next steps are actual-trained sparse-backend admission, measured serving, modern replication and a defined customer deployment. Software/model licensing can precede accelerator IP and custom silicon. Mobile and robotics extend the platform toward always-available contextual AI and adaptive machines.

The upside is ownership of a widely used AI infrastructure layer: models and runtime, licensed compute IP, and potentially a family of learning-capable event processors. Repeatable datacenter deployments or OEM design wins can support large businesses; successful frontier-scale models and cross-domain adoption could make this a foundational platform.

The current proposal is a **€3M raise at €50M priced pre-money**, with **€100M** as a stretch scenario. The case combines platform ambition, mechanisms with scoped empirical support, relevant founder experience and a staged execution plan. It does not price the company by adding application markets or claiming AGI. The €3M/18-month program buys team, quality replication, runtime proof and hardware feasibility and company-owned IP protection, including patent review and selective filings; attainable terms depend on investor conviction and diligence.

The expanded opportunity includes native irregular ingestion, adaptive asynchronous codecs, task-oriented communications and learned memory policies. Shared sensing/action/language paths could let embodied experience inform cognition. These are testable opportunities, not completed application wins. Existing learned codecs, irregular-stream models and VLA systems are relevant precedents and controls.

Prepared 3 October; revised 4 October 2026. Read the [current valuation rationale](#), [full investment case](#), [opportunity register](#) and [current research evidence](#). The [original pitch](#) preserves earlier pricing. This remains a fundraising thesis; scalable generalism, commercial advantage and clockless hardware economics require their own proofs.