

One substrate: three completed evidence anchors

The ambition spans language and reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. The evidence now connects token sequences, anonymous interleaved process logs and mixed-type comparisons within the temporal/sparse family.

Evidence front	Completed indication	Scope and next proof
Anonymous interleaved processes	FAS mean AUROC 0.5924 vs 0.5587 at 256 events; 0.7370 vs 0.7272 at 512 events. Every native seed beats the best generic control at both points.	Three native seeds; FAS v1 quality wins against six saved anonymous generic controls. Privileged diagnostics excluded. Sealed v2 and stronger neural references follow.
Mixed-type tables, small data	100% on 256 synthetic DEV rows from 64 FIT rows; 6,370 learned parameters. Typed numeric, categorical, Boolean and missingness comparisons feed the integrated temporal core.	One seed, fixed predicates. Relabeling, column order, missing-value semantics and actual alternative-write credit checks pass. Learned discovery and real tables versus trees follow.
Token language and useful memory	Two seeds at 1M TRAIN tokens: 7.2515 / 7.2636 DEV NLL, mean 7.2575. Memory erasure raises loss 0.0104 / 0.0174.	GPT-2 FineWeb, two passes, 2,040 DEV targets. Public Transformer comparison, complete larger fitting work and reserved 4M point follow. No fitted scaling law.

The reusable construction is the advantage: type-respecting interfaces produce addressed events; persistent memory, temporal computation, selective updates and counterfactual credit learn their consequences. Capacity, selected activity, timing and learning are separate design axes.

These instances were fitted separately. The next unification test shares trained core parameters across domains and measures transferred skills and retention. Source-bound headline packet:
report/unification_headline_evidence_20261006_v1.json. All historical evidence pages are retained.