

Sleeping Machines — Part III — Experiment record

Generated by `report/split_report.py` from `REPORT.md` (SHA-256 `ef5ac4d381e1689d...`) on 2026-10-09 17:00 UTC. Do not edit by hand: edit the record or `report/split_manifest.json` and rerun. [Part I — The science](#) is the reader-facing report; [Part II — Methods](#) holds protocols, accounting and implementation checks; [Part III — Record](#) holds every experiment entry. `REPORT.md` remains the complete combined record.

This part opens with the public benchmark record and the model-family chapter, then lists every experiment entry in record order, wins and losses alike. Part I is the reader-facing account; earlier versions are in git history.

Contents

A. Public benchmark record and model family (2 sections)

1. Public benchmark record — 9 October 2026
2. The model family and its place in the landscape — 4 October 2026

B. Experiment record (149 sections)

1. Appendix A. Deep event recognition
2. Appendix A (continued). Learned timing and transfer
3. Appendix A (continued). One temporal encoder
4. Appendix A (continued). Making depth useful
5. Appendix B. Breadth of the common implementation
6. Appendix B (continued). Total training cost
7. Appendix B (continued). Completed language references
8. Appendix B. New completed AWS language evidence
9. Appendix B. AWS hierarchy depth: a useful constraint
10. Appendix B. AWS early temporal mechanism screen
11. Appendix B. AWS early temporal mechanism screen
12. Appendix B. AWS early language mechanism screen
13. Appendix B. Protected state/shared rules: order
14. Appendix B. Protected state/shared rules: order
15. Appendix B. Protected state/shared rules: order
16. Appendix B. Protected state/shared rules: order
17. Appendix B. Protected state/shared rules: order
18. Appendix B. Protected state/shared rules: order
19. Appendix B. Protected state/shared rules: paired_timing
20. Appendix B. Protected state/shared rules: paired_timing
21. Appendix B. Native frozen confirmation
22. Appendix B. Native frozen confirmation
23. Appendix B. Persistent-write credit diagnosis
24. Appendix B. Ours and boosted trees: banknote
25. Appendix B. Ours and boosted trees: wine_red
26. Appendix B. Addressed-state write credit
27. Appendix B. Learned history and credit reach
28. Appendix B. Credit horizon with fixed context
29. Appendix B. Banknote: no confirmed advantage

30. Appendix B. Frozen banknote confirmation
31. Appendix B. Frozen banknote confirmation
32. Appendix B (continued). Ours: language work as scaling develops
33. Appendix B (continued). Ours and neural controls: accuracy versus FLOPs
34. Appendix B (continued). Later native language: quality versus fitting work
35. Appendix B (continued). Accuracy versus inference FLOPs
36. Appendix B (continued). Completed language variants and work
37. Appendix B (continued). Completed language variants and work
38. Appendix B (continued). Completed language variants and work
39. Appendix B (continued). Completed language variants and work
40. Appendix B (continued). Completed language variants and work
41. Appendix B (continued). Ours: per-position race KV memory
42. Appendix B (continued). Ours: per-position race KV memory
43. Appendix B (continued). Ours: per-position race KV memory
44. Appendix B (continued). Ours: per-position race KV memory
45. Appendix B (continued). Ours: cheaper learning updates
46. Appendix B (continued). Ours: shared-match temporal arrivals
47. Appendix B (continued). Ours: independent temporal heads
48. Appendix B (continued). Ours: completed head/data scaling
49. Appendix B (continued). Ours: longer temporal credit
50. Appendix B (continued). What most reduces research uncertainty
51. Appendix B (continued). Ours: content-gated temporal reception
52. Appendix B (continued). Ours: count-carrying native receivers
53. Appendix B. Longer credit versus learned count smoothing
54. Appendix B. Does the trained deep state contribute?
55. Appendix B. Restoring addressed value credit: completed fits
56. Appendix B. Fixed-feature credit and compiled inference
57. Appendix B. Frozen feature geometry and matched count calibration
58. Appendix B. Beyond query counts: balanced joint learning
59. Appendix B. Frozen retention versus interaction decoding
60. Appendix B. Protected outcomes and joint terminal races
61. Appendix B. Joint-credit work, capacity and activity
62. Appendix B. Fixed joint-credit confirmation
63. Appendix B. A terminal symmetry, not a learning impossibility
64. Appendix B. Practical headroom: a stronger joint-event table
65. Appendix B. Real gesture practical calibration
66. Appendix B. Integrated real-packet quality and work
67. Appendix B. Real-packet learning, capacity and total workflow
68. Appendix B. Frozen practical inference and storage
69. Appendix B. Real learned features versus the initial reservoir
70. Appendix B. Shared fitting noise: covariance hypothesis fails its gate
71. Appendix B. Causal persistent-state access: positive probe, qualified cause
72. Appendix B. Retained payloads, clocks, layers and regularization
73. Appendix B. Completed packet-scale clock comparison
74. Appendix B. Completed matched credit pilot: terminal pairs / seed6
75. Appendix B. Completed matched credit pilot: state choice / seed6

76. Appendix B. Completed matched credit pilot: state choice / seed7
77. Appendix B. Completed matched credit pilot: state clock / seed6
78. Appendix B. Actual-write choice: independent-seed gate fails
79. Appendix B. Why actual write and joint-clock credit need a test
80. Appendix B. Frozen variance audit after the failed clock pilot
81. Appendix B. Calibration: noisy routing changes global gradient balance
82. Appendix B. Exact probabilities do not remove winner-dependent credit bias
83. Appendix B. Evolution offset preserves signal-time coupling
84. Appendix B. Completed coupled-time offset pilot: seed6
85. Appendix B. Completed coupled-time offset pilot: seed7
86. Appendix B. When added representation freedom can help
87. Appendix B. Other-host completed native capacity comparisons
88. Appendix B. Strong compact controls rule out an easy storage claim
89. Appendix B (continued). Diagnostic: count receivers over the temporal carrier
90. Appendix B (continued). Ours: native temporal-core language
91. Appendix B (continued). Ours: native event/state capability
92. Appendix B (continued). Ours: occupied state and activity
93. Appendix B (continued). Ours: compact delayed learning
94. Appendix B (continued). Ours: historical write credit
95. Appendix B (continued). Ours: frozen information-flow diagnosis
96. Appendix B (continued). Ours: separate online neural learning
97. Appendix B (continued). Separate statistical language baseline
98. Appendix B (continued). Learned language and depth
99. Appendix B (continued). Persistent learned language
100. Appendix B (continued). Learned event language: training and inference work
101. Appendix B (continued). Statistical online-learning pilot
102. Appendix C. Preserved historical results and revisions
103. Completed AWS 90M native language comparisons
104. Completed public archive benchmark campaign
105. Appendix. Native language at 10M: the integrated core, segment-batched
106. Appendix. Tokenized256K: bounded learning-rate comparison
107. Appendix. Measured language quality, capacity and compute
108. Appendix. Tokenized language: 1M capacity/data comparisons
109. Appendix. Tokenized language: 256K capacity/data comparisons
110. Appendix. Tokenized language: 64K capacity/data comparisons
111. Appendix. Independent seed for selected 64K capacity
112. Appendix. Tokenized messages: replicated payload contribution
113. Appendix. Tokenized8K: replicated learning and memory-use decisions
114. Appendix. Language strengths: measured decisions and complete work
115. Appendix. Tokenized language: confirmed small-fit learning
116. Appendix. Tokenized language: integrated learning development
117. FAS native frozen-recipe replication
118. Appendix. Public benchmarks: NeuroBench and SHD (status from completed result files)
119. Appendix B. AWS coarse temporal screens: two seeds pass
120. Appendix B. Full coarse matrix: positive means, failed all-seed gates
121. Appendix B. Strong coarse controls and frozen readout evidence

122. Appendix B. Integrated quadratic head: contracts pass, quality fails
123. Appendix B. Receiving losing arrivals requires boundary credit
124. Appendix B. Native reception intervention, seed6: no smoke admission
125. Appendix B. Native reception intervention, seed7: no smoke admission
126. Appendix B. AWS frozen polynomial decoder: both nomination gates fail
127. Appendix B. Race-scaled reception: available arrival support
128. Appendix B. Memory-conditioned keys explain routing confidence
129. Appendix B. Clock-preserving route calibration, seed6
130. Appendix B. Clock-preserving route calibration, seed7
131. Appendix B. AWS frozen replay critics and parameter calibration
132. Appendix B. Replay allocation: fresh score gains, parameter failure
133. Appendix B. AWS compact context head: practical controls remain stronger
134. Appendix B. Integrated exact replay: first seed does not improve quality
135. Appendix B. Native conditional choice-clock parameter geometry
136. Appendix B. Bounded emitter bridge restores the missing direction
137. Appendix B. Completed fine-packet replay confirmations
138. Appendix B. AWS deeper all-race pilot: matched gates fail
139. Appendix B. Parameter-targeted allocation retains a scoped positive gain
140. Appendix B. AWS private state and shared maps: pilot gates fail
141. Appendix B. Restored native sensitivity has small, mixed route utility
142. Appendix B. AWS depth8 replay: seed7 positive pilot
143. Appendix B. AWS depth8 replay: seed8 confirmation fails
144. Appendix B. Persistent representation and missing delayed credit
145. Appendix B. Native text8 horizon return: four credit reversals
146. Appendix B. Horizon confirmation: downstream effects vary by context
147. Appendix B. Deep fitting diagnosis: looser clipping fails its prediction
148. Appendix B. Deep race sampling: preserve mixed quality and complete work
149. Appendix B. Live counterfactual messages: conditional content and route credit

A. Public benchmark record and model family

Public benchmark record — 9 October 2026

Sealed tests on the official splits, scored once per run; losses and pending verdicts in the same table. Own numbers are computed from the result files in the source-bound packet [public_wins_headline_evidence_20261009_v2.json](#). The reader-facing account of these results, their mechanisms and theory is [Part I](#); the appendices below are the experiment record. Earlier versions of this report are in git history.


Benchmark	Ours (sealed test)	Best published	Verdict and cost
EasyTPP Taxi (nats/event, higher better)	0.5250 $\pm$ 0.0010	0.522 (S2P2)	Win, 5 of 5 seeds; 1/12 of S2P2 parameters and compute; reproduced on separate hardware
EasyTPP Taobao	1.3991 $\pm$ 0.0025	1.318 (IFTTP)	Win, 5 of 5; +0.081 at 0.92 $\times$ S2P2 compute
EasyTPP StackOverflow	-2.1525 $\pm$ 0.0045	-2.163 (S2P2)	Win at matched size and compute, 5 of 5
EasyTPP Retweet	-6.3262 $\pm$ 0.0009	-6.348 (NHP)	Win, 5 of 5; 1/15 of S2P2 parameters and compute
EasyTPP Amazon	0.8028 $\pm$ 0.0007	0.781 (S2P2)	Win, 5 of 5; 0.29 $\times$ S2P2 compute
EasyTPP, one configuration for all five	0.526 / 1.397 / -2.144 / -6.324 / 0.802 (Taxi / Taobao / SO / Retweet / Amazon)	as above	Win on all five without per-dataset tuning; 24 of 25 seeds ahead (pre-registered)
P19 sepsis (ICU records)	AUPRC 0.639 $\pm$ 0.039; AUROC 0.916	0.583 / 0.903 (MTM)	Win, five official splits; 62,681 parameters
PAM wearable activity	accuracy 0.978 $\pm$ 0.007; F1 0.980	0.975 / 0.976 (MTM)	Win, 4 of 5 splits ahead; 46,316 vs 873K parameters
TGB tgbn-trade (node affinity, NDCG@10)	0.8680 $\pm$ 0.0005	0.863 (NAVIS, ICLR 2026)	Win, 3 of 3 pre-registered seeds; 2,107 parameters
P12 mortality (ICU records)	AUROC 0.871; AUPRC 0.585	0.880 / 0.586 (MTM)	Behind on AUROC, level on AUPRC
Bosser & Ben Taieb MOOC (total NLL, lower better)	-226.8 $\pm$ 6.3	bar -239.7 (pre-registered); best single model -233.8	Loss (frozen one-configuration model); six datasets running
TGB tgbl-wiki (link prediction, MRR)	validation 0.8518 (7,995 parameters)	TPNet validation 0.842, test 0.827	In development, ahead on validation; sealed seeds running
FAS v2 (anonymous interleaved logs)	validation 0.702	time-encoded Transformer reference 0.704	In development; sealed verdict pending the reference seeds

The model family and its place in the landscape — 4 October 2026

Sleeping Machines is a family of causal networks of stateful temporal programs. Content, computational time, addresses and persistent evidence jointly determine computation. One interface supports different local programs and heterogeneous compositions; the fitted native model is one member. The integrated target combines computational delays and races, hard selective state updates, small learned messages, separate keys and values, and credit to unrealized alternatives, with capacity beyond costly activity.

Sleeping Machines in the landscape of model families

A stateful temporal-program envelope; intersections and conditional constructions, not automatic superiority.



No required global periodic tick. Current emulators are clocked; hardware energy, general TTT and morphing remain open.

Known primitives are credited; the proposed contribution is their coherent integration and measured consequences.

Position among model families

The family spans recurrence and SSMs, attention and Transformers, sparse experts, statistical and historical memory, test-time training and event-driven temporal systems. Its contribution is their coherent temporal and selective integration, and the quality, learning and resource behavior that follows.

Selective and dense regions coexist in one model. Synchronous barriers are local schedules; the general semantics need no global periodic update tick. Rich interaction goes where a task requires it, and other regions stay economical. Online learning and batching use the same structure, and so do the next targets: native test-time training, structural adaptation, shared-world multimodal learning and clockless hardware.

Read in order: [overview](#) → [formal core](#) → [complete contrasting members](#) → [design choices](#) and [composition rules](#). The [visual atlas](#), [implementation review](#) and [claim/evidence map](#) supply detail. Computational power and work, representation and trainability are reviewed from atoms to complete systems.

The [opportunity register](#) covers direct irregular inputs, adaptive asynchronous codecs, task-oriented communication, learned memory policies and embodied/cognitive transfer, each with its precedents and first proof condition.


B. Experiment record

Appendix A. Deep event recognition

The strongest completed speech result in the model family is **79.69%** on 512 private held-speaker utterances (selected single six-block prefix; trained with twelve blocks). Published official-test results below use a different partition; they are reference targets.

Ours: private development configuration	Correct	Accuracy ↑
Original eight-layer parent	370/512	72.27%
Parent + six-block residual; pass 3	408/512	79.69%
Single six-block temporal encoder; pass 0	406/512	79.30%
Calibrated six-block continuation	400/512	78.12%
Bounded twelve-block encoder	400/512	78.12%
Directional twelve-block encoder	401/512	78.32%
Selected single six-block prefix; trained with twelve blocks	408/512	79.69%

Ours = Sleeping Machines



e143 temporal residual learning

A six-block width-128 temporal encoder learns corrections while the inherited eight-layer parent stays frozen. It adds 395,814 parameters to the parent's 53,296. Signed modal states, nonlinear gates and residual vectors learn from all source identities and original event times before causal pooling. This is a larger parallel model, not fourteen sequential layers; completed-utterance supervision does not yet teach calibrated early answers. On 657 disjoint utterances from the same held speakers, accuracy improves from 70.47% to 77.63%.


Published reference; official-test protocol	Reported accuracy ↑
Ours: full official SHD comparison	Pending
EventSSM: asynchronous learned state-space layers	95.9%
S7: input-dependent temporal state	96.3%

Our private sample uses training-file speakers 3/6; official test accuracy is unmeasured. The temporal residual uses three passes and a fresh optimizer; its larger capacity and budget are not a matched single-factor comparison. Its listed score selects the best private-development epoch; the curve shows all three. The original parent was fitted on 4,096 examples. Continuation timers exclude loading; the residual timer includes it. Both include evaluation and are not energy measurements. One seed. References: [EventSSM](#), [S7](#). These published scores were checked against the original papers; they are targets for a full official-test comparison, not scores on our private split.

Appendix A (continued). Learned timing and transfer

Learned delays contribute to the answer. Restoring the hidden clocks to their initial values, with source vectors, state/value maps and the trained classifier retained, loses 13 correct answers. Timing changes the temporal interactions used by the representation; it performs computation.

Ours = Sleeping Machines



e145 learned timing and transfer

Ours: same trained readout; subsystem reset	Correct	Accuracy ↑
All learned parameters retained	408/512	79.69%
Only hidden clocks reset	395/512	77.15%
Only source vectors reset	373/512	72.85%
Only state stack reset	358/512	69.92%
Only decay/frequency parameters reset	407/512	79.49%

The temporal stack and source vectors also learn useful coordinated representations. Resetting the stack loses 50 correct answers; resetting sources loses 35. Decay/frequency resets change one decision. These changes depend on the fitted solution's coordination; their effects cannot be added or treated as a matched comparison of retrained architectures.

The 657-utterance audit is disjoint from fitting and the development sample and uses the selected checkpoint unchanged. It gains 62 correct answers and loses 15, for a net improvement of 47. Both samples use the same two held training speakers. Official-test and additional-speaker generalization are the next evaluation targets.

A stronger mathematical account of routing

An affine packet summary can preserve both the final state and the average of raw temporal states, including their teachers. This permits richer pooling before expensive nonlinear maps. A second derivation shows why many almost-equal delays may offer little usable choice: their effects point in nearly the same direction. Diverse payloads and temporal modes, sufficient delay spread and downstream visibility determine useful route reserve.


One exploratory run: 449,110 total parameters, inherited parent plus a trained six-block encoder. Three passes use about 33 minutes including preparation/evaluation and 1.79 GiB peak RSS; the guarded host retains at least 10,361 MiB sampled available memory. Complete physical work and joules remain unmeasured. Formal statements and numerical contracts are in THEORY §§226–236.

Appendix A (continued). One temporal encoder

A six-block temporal encoder retains nearly all the combined model's development accuracy through one ordinary query head. Deployment removes the frozen parent: 395,814 parameters replace 449,110. Modal states, gated vector messages and winning delays remain. Its weights inherit earlier encoder training; combined teacher predictions are used only to initialize the head.

Ours: private development configuration	Correct	Accuracy ↑	NLL ↓
Combined model: parent + temporal correction	408/512	79.69%	0.684
Single encoder: clean head, before continuation	374/512	73.05%	0.818
Single encoder: paired head, selected pass zero	406/512	79.30%	0.631
Single encoder: selected trained prefix	408/512	79.69%	0.625

Ours = Sleeping Machines



e152 single encoder learning

Fitting the head on clean and transformed speech improves held accuracy by 32 answers with the temporal features frozen. Its covariance penalty suppresses class-visible nuisance variation. The right panel compares the heads on the same features; the left shows all subsequent unrestricted encoder passes. The selected trained prefix reaches 518/657 (78.84%) versus 510/657 (77.63%) for the combined model on the reused disjoint audit. Audit labels do not choose the checkpoint.


Training the directional twelve-block extension, then selecting its six-block prefix on development, retains the combined model's 408 correct answers with lower NLL and 395,814 deployed parameters. The extra training blocks are removed after their learned contributions reduce held accuracy. Selecting the six-block prefix improves deployment quality/work.

Seed 6, private train-file speakers 3/6; official-test parity remains unmeasured. Head fitting uses 6,144 unique fitting utterances: one clean view for the first arm, clean plus one transformed view for the paired arm. The arms also change regularization and use three/two encoder passes respectively, so total budgets are not matched. All continuation epochs are plotted; selection uses development accuracy, then NLL. Extra teacher/cache/head work and inherited fitting must be charged. Modal/vector maps are locally dense; no empty ticks or event-pair attention are added. The selected prefix additionally inherits the full twelve-block fitting pass; pruning does not erase that training cost. Prefix/full choice is post-hoc private-development selection. Formulae and numerical checks: THEORY §§237–264. One CPU forward evaluation takes 11.90 s for the selected prefix versus 23.71 s for the combined model. Packing/query included, loading excluded; one timing observation, not joules.

Appendix A (continued). Making depth useful

Identity growth preserves the classifier and old teachers while added output maps receive label credit. They must also change useful features. A tightly bounded twelve-block extension learns weights but changes no audit decisions when its six appended blocks are removed.

Ours: one matched fitting pass	Parameters	Fit accuracy ↑	Private accuracy ↑	NLL ↓
Six-block control	395,814	90.72%	78.12%	0.639
Twelve blocks: bounded outputs	696,888	90.72%	78.12%	0.639
Twelve blocks: directional units	696,120	90.77%	78.32%	0.644



e164 depth use and work

Directional conditioning normalizes temporal-state features before their output map. An invertible coordinate change rescales classifier-sensitive directions and preserves hidden null directions for later computation. The fixed transform folds into an ordinary map at deployment. Initial outputs and old teachers remain exact; fitting replays verify the actual proposed update.

The directional model reaches 509/657 on the reused audit. Removing its six appended blocks changes 11 predictions: 0 are correct only with the blocks and 9 only without them. This measures fitted contribution with the trained prefix/head retained; it is not a retrained architecture comparison.

All arms inherit the paired-head checkpoint and use 6,144 fitting utterances, the same order, channel/time transformations and one encoder pass. Old-group LR is 0.0000203125; directional new groups use 0.0001953125 from fitting-only replay. Changed normalization and update coordinates form one intervention. New blocks initially add 36 ms latency; labels supervise completed untimed utterances. Audit reuse is explicit; official-test parity remains unmeasured. CPU points are one warmed observation per model, packing/query included and loading excluded; energy is unmeasured. Weight-coordinate folding, all source work and added depth must be charged during training. Local maps remain dense, with no empty ticks or event-pair attention. Theory §§249–264.


Appendix B. Breadth of the common implementation

These small development screens test one implementation across tasks. Text and market variants include separately fitted statistical evidence. Ours denotes Sleeping Machines; TF is the saved Transformer. Accuracy improves upward; prediction loss and FLOPs improve downward.

Task / quality direction	Ours: quality	TF: quality	Ours: inference MFLOPs ↓	TF: inference MFLOPs ↓
Text8 bpc ↓	2.915 bpc	3.729 bpc	1.26	1.84
Market event prediction nats/event ↓	3.823 nats/event	4.208 nats/event	1.26	1.84
Temporal composition accuracy ↑	97.3%	90.2%	0.29	0.38
MNIST accuracy ↑	75.8%	69.9%	1.66	2.57

Task / quality direction	Ours: quality	TF: quality	Ours: inference MFLOPs ↓	TF: inference MFLOPs ↓
Event-camera gestures accuracy ↑	59.1%	15.9%	23.99	124.23

Ours = Sleeping Machines



breadth work ratios

Ours uses eight layers and TF two, both width 32, with the same neural-fitting examples, encoding, objective and eight epochs. Inference counts maps/scans (ours) and maps/attention (TF): two FLOPs per multiply-add, excluding padding, scalar nonlinearities and expert preparation. Training includes backward, clipping, Adam and our evidence/calibration; its ledger follows.

Lower core inference work in every screen. Gestures use **5.18x less** with 59.1% versus 15.9% accuracy. That 44-query screen has an underfitting TF control; a general vision claim requires complete benchmarks and stronger references.


Why training can cost more: the race core evaluates all three candidate vector payloads during training, versus only the winner during inference. Its eight layers also exceed the reference's two. With short contexts, that work outweighs the saved attention cost; these rows do not show a training efficiency advantage. On the longer event-camera prefixes, the common model's estimated total uses 49.0% of the reference training arithmetic, including calibration.

Seed 6; neural fit/development counts: text 2,048/256, market 512/256, temporal 1,024/256, MNIST 1,024/256, gestures 88/44. The common text model also has a separately fitted 32,768-character evidence bank; market evidence is fitted on a prior day. The references have no such bank. MNIST uses pooled training-set images; gestures use first-second prefixes and disjoint users. Batch 16, or four for gestures. No official real-data test. Market fixed evidence: 3.670 nats/event.

Appendix B (continued). Total training cost

The ledger estimates the entire completed fitting budget for each reported model. It includes prediction and loss, backpropagation, gradient clipping and Adam across all eight epochs. The common model also pays for its evidence bank and initial readout calibration. These models were trained from initialization; there is no inherited neural fitting to omit.

Ours = Sleeping Machines



e172 complete training work

Model/task	Forward + loss	Backward	Clip	Adam	Evidence + calibration	Total
Language: Ours	62.633	130.873	0.165	0.658	2.930	197.26
Language: TF	31.246	63.990	0.083	0.334	0	95.65
Market: Ours	15.657	32.716	0.041	0.164	0.733	49.31
Market: TF	7.811	15.996	0.020	0.081	0	23.91
Temporal: Ours	7.725	16.546	0.081	0.324	0.364	25.04
Temporal: TF	4.773	9.540	0.041	0.162	0	14.52
Mnist: Ours	40.394	84.376	0.083	0.333	1.886	127.07
Mnist: TF	33.026	68.489	0.049	0.197	0	101.76
Dvs: Ours	51.580	107.386	0.028	0.112	2.404	161.51
Dvs: TF	99.023	230.261	0.014	0.057	0	329.35

All table values are estimated GFLOPs for the whole fitting run, not per query. A multiply-add counts as two operations. Forward/loss and backward use the saved E172 four-query operator trace scaled by recorded map/scan work. Ours uses its logged candidate-map/scan counts. Reference padding is reconstructed from all fitting prefix lengths, the original shuffle seed and batch sizes, including fused attention products. Clipping and Adam are charged once per original step. Other arithmetic and the small calibration eigensolver are estimates.

Event-target arithmetic excludes padding and simulator dispatch/allocation; required candidate maps, losing-value teaching, scans and learning remain charged. References use the same FLOP convention. Evaluation, search, encoding, special functions, integer/index work, comparisons and memory traffic are outside these totals. These single-seed screens have different depths/quality and do not measure event hardware, matched-quality cost or energy. Ledger: [estimate_training_work.py](#).

Appendix B (continued). Completed language references

The Transformer and LSTM benchmarks have already been run. Their completed result files remain in the repository and are reused as reference targets for the full learned-event benchmark. Lower bits per character (bpc) means better prediction.

Model	Fitting characters / passes	Test bpc ↓	Full training FLOPs ↓
Ours: learned event-state model (planned)	10M / four passes	Pending	Pending
LSTM; width 512, one recurrent layer	10M / six passes	1.799	432.59T

Model	Fitting characters / passes	Test bpc ↓	Full training FLOPs ↓
Transformer; width 256, four layers	10M / four passes	1.908	888.78T
LSTM; width 512, one recurrent layer	90M / 6 passes	1.661	3.89P
Transformer; width 256, four layers	90M / 4 passes	1.604	8.00P

The 10M references score exactly the same 999,999 text8 targets in [95M,96M), with frozen validation-selected weights and cold initial context. The 90M LSTM uses the same test interval and its saved recurrent scoring protocol. All use the historical 27-character alphabet and 200,000-character validation selection; data budgets, capacities and fitting passes differ.

Training estimates include every fitting step, forward/loss, backpropagation, gradient clipping and Adam. Backward is approximated as twice forward; multiply-add counts as two FLOPs. G/T/P mean billion/trillion/quadrillion. Validation/test evaluation, memory traffic and runtime are outside these arithmetic totals.

Ours: learned-language benchmark status

The planned comparison uses 10M fitting characters, four passes, 200,000 validation characters and the same 1M test interval. Priority is the integrated sparse/timed architecture documented in the next appendix. Completed dense-carrier fits remain diagnostic controls; their quality does not establish sparse-model performance. Numerical contracts and progressively larger integrated development fits precede promotion. The aligned full-test result and its fitting work remain pending. The preserved 28,403-parameter pilot's 3.351 development bpc uses a smaller fitting budget and different split; it is not a comparable test result.

Earlier 1M-character references also remain saved: LSTM **2.179** and Transformer **2.367 test bpc**, each with twenty fitting passes. The separate count/copy baseline and the cross-task Transformer/retrieval LSTM comparisons remain in their labeled sections and Appendix B.

Saved evidence: [10M LSTM aligned result](#); [10M Transformer aligned result](#); [90M LSTM saved result](#); [90M TF saved result](#). An earlier 90M Transformer attempt was interrupted by its RSS watchdog before producing a completed test result; its provenance is preserved.

Appendix B. New completed AWS language evidence

The six-block sparse receiver model reaches 3.106 development bpc on 32K fitting characters. It learns meaningful prediction, but this is the older single-head, observed-character-pool variant, not the newer eight-block native/content-gated reception model. The large dense controls are completed reference targets with much better held-out quality and much larger fitting budgets.

Model	Fit chars / passes	Dev bpc ↓	Test bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓
Ours sparse d32/L6	32,768/4	3.106	Not scored	114.25	0.872
LSTM512	90M/6	1.616	1.661	3,893,396.04	7.210
Transformer256/L4	90M/4	1.580	1.604	8,000,253.35	22.223

Model	Parameters	Guarded wall h	Peak RSS MiB	Infer MFLOPs / char ↓
Ours sparse	1,388,871	0.39	558.4	0.0885
LSTM512	1,199,323	10.64	2246.9	2.4024
Transformer256/L4	3,238,427	31.18	3822.5	14.8104


All rows use the same units and whole-fit/per-target denominators within each column. Sparse work is a representative full-step arithmetic estimate with unit-weight special functions; dense work uses shapes and backward = twice forward, including clipping/Adam. Both exclude evaluation and traffic. Dense inference uses the saved shape convention, including Transformer window overlap; sparse inference is a winner-only trace. Guarded wall and RSS include different simulator/runtime overheads and are not energy measurements.

The sparse development set has 8,191 targets; dense selection uses 200,000 validation characters and the saved 1M test interval. Neither their development scores nor their fitting budgets are matched. Raw resource gaps cannot establish comparable-quality or iso-FLOP supremacy. The interrupted Transformer attempt remains preserved; only its completed retry appears as a quality point. All completed points, including the 90M Transformer, are retained in the common quality/work and inference figures.

Appendix B. AWS hierarchy depth: a useful constraint

These are older race-network diagnostics on the fixed depth-three Random Hierarchy Model. They test composition and depth, not the newer integrated language/event construction. Every shown model uses width 200, 64K fitting examples, ten passes, architecture seed 0 and rule seed 0.

Ours = Sleeping Machines



aws hierarchy depth

Model	Depth	Final held-out accuracy	Guarded wall s	Peak RSS MiB
Ours plain race	1	70.12%	70.4	129.3
Ours plain race	2	83.72%	227.4	202.4
Ours plain race	3	85.06%	382.2	249.7
Ours plain race	4	84.12%	536.5	252.1
Ours residual-2	2	72.72%	266.1	217.4
Ours residual-2	3	74.54%	610.4	369.3
Ours residual-2	4	72.56%	1094.7	574.0


Plain depth 1/2/3/4 gives 70.12/83.72/85.06/84.12%: depth helps to three blocks, then regresses slightly. The completed residual-2 depth 2/3/4 variants give 72.72/74.54/72.56%, below their plain counterparts. Depth-four residual accuracy rises throughout its ten passes; its longer-budget convergence is untested. This constrains the tested implementation and budget, not all residual paths.

Only completed provenance files produce rows. Plain and residual depth-four runs are distinct. Final epoch is reported; evaluation curves are visible throughout fitting, so this is exploratory held-out evidence, not independent confirmation. Single rule/model seed; no 64K matched dense-control or complete FLOP/energy supremacy is inferred.

Appendix B. AWS early temporal mechanism screen

Small completed integrated pilots from the committed fast matrix. These answer mechanism questions before larger scaling; their development budgets differ from the saved main language and native512-target results. Contracts and accounting smokes are excluded.

Ours = Sleeping Machines



aws fast screen temporal

Model	Fit / passes / dev	Dev quality	Whole fit GFLOPs	Fit MFLOPs / target	Infer MFLOPs / target
Ours timing full	128/4/64	78.12%	0.265	0.518	0.1101
Ours timing rank	128/4/64	79.69%	0.265	0.518	0.1100
Ours order full	128/4/64	51.56%	0.265	0.518	0.1104
Ours order pathwise	128/4/64	51.56%	0.263	0.514	0.1104

Model	Parameters	Receivers/commits/matches per event	Wall s	Seed
Ours timing full	42,898	128/16/32	791.8	6
Ours timing rank	42,898	128/16/32	799.8	6
Ours order full	42,932	128/16/32	789.3	6
Ours order pathwise	42,932	128/16/32	768.4	6

FLOPs count MAC as two and special functions once; all losing-value credit and Adam are charged. Temporal inference per query includes intervening input events; language per target is per character. Capacity/activity counts are per event, not per query. CPU simulation/audit overhead, RNG, traffic and physical energy are separate; AWS wall observations may include authorized CPU concurrency. Temporal tasks and altered source counts are explicitly named; different tasks do not form a single accuracy scaling curve. Single-seed development evidence, not supremacy. Paired independent seeds and frozen held-out confirmation precede benchmark promotion.


Interpretation: observed-time accuracy 78.12% versus refitted rank-time 79.69% does not yet demonstrate an elapsed-time advantage. State-clearing and stretched-gap probes are diagnostic interventions, not refitted controls.

Next experimental questions: state clearing and stretched silent gaps damage order predictions; private source rules lose training exposure as capacity grows. The next integrated battery tests protected memory, shared rules with private state, and paired timing whose labels cannot be inferred from rank alone.

Appendix B. AWS early temporal mechanism screen

Small completed integrated pilots from the committed fast matrix. These answer mechanism questions before larger scaling; their development budgets differ from the saved main language and native512-target results. Contracts and accounting smokes are excluded.

Ours = Sleeping Machines



aws fast screen temporal

Model	Fit / passes / dev	Dev quality	Whole fit GFLOPs	Fit MFLOPs / target	Infer MFLOPs / target
Ours order sources16	128/4/64	48.44%	0.280	0.548	0.1103
Ours order sources64	128/4/64	32.81%	0.341	0.666	0.1104
Ours order shallow control	128/4/64	51.56%	0.134	0.262	0.0569

Model	Parameters	Receivers/commits/matches per event	Wall s	Seed
Ours order sources16	157,940	512/16/32	829.1	6
Ours order sources64	617,972	2048/16/32	980.5	6
Ours order shallow control	21,684	64/8/16	403.3	6


FLOPs count MAC as two and special functions once; all losing-value credit and Adam are charged. Temporal inference per query includes intervening input events; language per target is per character. Capacity/activity counts are per event, not per query. CPU simulation/audit overhead, RNG, traffic and physical energy are separate; AWS wall observations may include authorized CPU concurrency. Temporal tasks and altered source counts are explicitly named; different tasks do not form a single accuracy scaling curve. Single-seed development evidence, not supremacy. Paired independent seeds and frozen held-out confirmation precede benchmark promotion.

Capacity interpretation: commits and matches remain fixed while available state grows, but fixed queries reduce per-source training exposure. The 64-source development set has one population; its collapsed bootstrap interval is not useful uncertainty.

Appendix B. AWS early language mechanism screen

Small completed integrated pilots from the committed fast matrix. These answer mechanism questions before larger scaling; their development budgets differ from the saved main language and native512-target results. Contracts and accounting smokes are excluded.

Ours = Sleeping Machines



aws fast screen language

Model	Fit / passes / dev	Dev quality	Whole fit GFLOPs	Fit MFLOPs / target	Infer MFLOPs / target
Ours native	512/4/1023	4.2003	0.290	0.142	0.0301
Ours reception	512/4/1023	4.1949	0.325	0.159	0.0348
Ours late	512/4/1023	4.1962	0.319	0.156	0.0341
Ours waiting	512/4/1023	4.1955	0.304	0.149	0.0318

Model	Parameters	Receivers/commits/matches per event	Wall s	Seed
Ours native	14,971	32/16/32	166.1	6
Ours reception	16,123	32/16/32	239.2	6
Ours late	16,123	32/16/32	205.1	6
Ours waiting	16,123	32/16/32	207.6	6

FLOPs count MAC as two and special functions once; all losing-value credit and Adam are charged. Temporal inference per query includes intervening input events; language per target is per character. Capacity/activity counts are per event, not per query. CPU simulation/audit overhead, RNG, traffic and physical energy are separate; AWS wall observations may include authorized CPU concurrency. Temporal tasks and altered source counts are explicitly named; different tasks do not form a single accuracy scaling curve. Single-seed development evidence, not supremacy. Paired independent seeds and frozen held-out confirmation precede benchmark promotion.

Appendix B. Protected state/shared rules: order

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S64 private/P0/observed/s6	95.02	0.1889	1.364	0.666	0.110
Ours S64 private/P0/observed/s7	98.34	0.0558	1.364	0.666	0.110
Ours S64 private/P0/observed/s8	98.83	0.0430	1.364	0.666	0.110
Ours S64 shared/P0/observed/s6	99.90	0.0279	1.054	0.515	0.110
Ours S64 shared/P0/observed/s7	99.90	0.0869	1.054	0.515	0.110

Construction	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours S64 private/P0/observed/s6	512/1024/4	616,964	2048	16/32

Construction	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours S64 private/P0/observed/s7	512/1024/4	616,964	2048	16/32
Ours S64 private/P0/observed/s8	512/1024/4	616,964	2048	16/32
Ours S64 shared/P0/observed/s6	512/1024/4	14,180	2048	16/32
Ours S64 shared/P0/observed/s7	512/1024/4	14,180	2048	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Protected state/shared rules: order

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S64 shared/P0/observed/s8	97.95	0.0723	1.054	0.515	0.110
Ours S16 shared/P0/observed/s7	44.53	1.3721	0.262	0.511	0.110
Ours S16 shared/P0/observed/s8	70.31	0.8630	0.262	0.512	0.110
Ours S16 shared/P2/observed/s7	26.17	1.3999	0.253	0.494	0.107
Ours S16 shared/P2/observed/s8	51.56	1.0933	0.253	0.494	0.107

Construction	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours S64 shared/P0/observed/s8	512/1024/4	14,180	2048	16/32
Ours S16 shared/P0/observed/s7	128/256/4	14,180	512	16/32
Ours S16 shared/P0/observed/s8	128/256/4	14,180	512	16/32
Ours S16 shared/P2/observed/s7	128/256/4	13,988	512	16/32
Ours S16 shared/P2/observed/s8	128/256/4	13,988	512	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Protected state/shared rules: order

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S16 private/P0/observed/s7	25.78	1.4432	0.280	0.548	0.110
Ours S16 private/P0/observed/s8	42.97	1.3150	0.281	0.548	0.110
Ours S16 private/P0/observed/s6	75.39	0.7592	0.280	0.548	0.110
Ours S16 private/P0/observed/s7	51.17	1.0313	0.280	0.548	0.110
Ours S16 private/P0/observed/s8	69.14	0.8689	0.280	0.548	0.110

Construction	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours S16 private/P0/observed/s7	128/256/4	157,940	512	16/32
Ours S16 private/P0/observed/s8	128/256/4	157,940	512	16/32
Ours S16 private/P0/observed/s6	128/256/4	157,700	512	16/32
Ours S16 private/P0/observed/s7	128/256/4	157,700	512	16/32
Ours S16 private/P0/observed/s8	128/256/4	157,700	512	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Protected state/shared rules: order

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S16 shared/P0/observed/s6	40.23	1.2576	0.262	0.512	0.110
Ours S16 shared/P0/observed/s7	32.03	1.4567	0.262	0.512	0.110
Ours S16 shared/P0/observed/s8	25.00	1.5001	0.262	0.512	0.110
Ours S16 private/P0/observed/s6	44.14	1.1194	0.280	0.548	0.110
Ours S16 private/P2/observed/s6	44.92	1.2147	0.271	0.529	0.107

Construction	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours S16 shared/P0/observed/s6	128/256/4	14,420	512	16/32
Ours S16 shared/P0/observed/s7	128/256/4	14,420	512	16/32
Ours S16 shared/P0/observed/s8	128/256/4	14,420	512	16/32
Ours S16 private/P0/observed/s6	128/256/4	157,940	512	16/32
Ours S16 private/P2/observed/s6	128/256/4	155,828	512	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Protected state/shared rules: order

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S16 shared/P0/observed/s6	75.39	0.8327	0.262	0.511	0.110
Ours S16 shared/P2/observed/s6	56.25	1.0359	0.253	0.494	0.107
Ours S4 private/P0/observed/s6	54.30	0.9787	0.265	0.518	0.110
Ours S4 private/P2/observed/s6	47.27	1.1060	0.256	0.500	0.107

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S4 shared/P0/observed/s6	70.70	0.8556	0.261	0.511	0.110

Construction	Fit/dev/passess	Parameters	State slots	Updates/scores per event
Ours S16 shared/P0/observed/s6	128/256/4	14,180	512	16/32
Ours S16 shared/P2/observed/s6	128/256/4	13,988	512	16/32
Ours S4 private/P0/observed/s6	128/256/4	42,932	128	16/32
Ours S4 private/P2/observed/s6	128/256/4	42,356	128	16/32
Ours S4 shared/P0/observed/s6	128/256/4	14,180	128	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Protected state/shared rules: order

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S4 shared/P2/observed/s6	56.64	1.0367	0.252	0.493	0.107

Construction	Fit/dev/passess	Parameters	State slots	Updates/scores per event
Ours S4 shared/P2/observed/s6	128/256/4	13,988	128	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Protected state/shared rules: paired_timing

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S4 private/P0/observed/s7	90.23	0.3483	0.265	0.518	0.110
Ours S4 private/P0/observed/s8	91.80	0.1989	0.265	0.518	0.110
Ours S4 private/P0/rank/s7	50.00	0.7153	0.265	0.518	0.110
Ours S4 private/P0/rank/s8	50.00	0.7151	0.265	0.518	0.110
Ours S4 private/P0/observed/s6	95.31	0.1850	0.265	0.518	0.110

Construction	Fit/dev/passess	Parameters	State slots	Updates/scores per event
Ours S4 private/P0/observed/s7	128/256/4	42,898	128	16/32
Ours S4 private/P0/observed/s8	128/256/4	42,898	128	16/32
Ours S4 private/P0/rank/s7	128/256/4	42,898	128	16/32
Ours S4 private/P0/rank/s8	128/256/4	42,898	128	16/32

Construction	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours S4 private/P0/observed/s6	128/256/4	42,898	128	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Protected state/shared rules: paired_timing

Completed integrated pilots only. Protected modes retain information during silence; temporal modes still evolve. Shared learned rules retain private addressed state and remove private source embeddings. Paired timing keeps marks/order identical while labels differ; rank-only prediction has an exact 50% paired ceiling under coupled noise.

Construction	Dev accuracy%	Dev NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours S4 private/P0/rank/s6	50.00	0.7035	0.265	0.518	0.110
Ours S4 private/P2/observed/s6	82.42	0.5131	0.256	0.500	0.106

Construction	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours S4 private/P0/rank/s6	128/256/4	42,898	128	16/32
Ours S4 private/P2/observed/s6	128/256/4	42,322	128	16/32

Exact full fitting includes producer graphs, losing proposals, backward, clipping and Adam; specials have unit weight. Independent population/pair uncertainty is distinct from seed uncertainty. Protected-prefix initialization also removes faster initial temporal modes; any timing change is not isolated spectral evidence. Scope remains synthetic pilot quality, not physical energy.

Appendix B. Native frozen confirmation

Completed 1,024-query synthetic holdout evaluations only. All fitted seeds and matched controls remain visible; selected checkpoints use development NLL before confirmation. Existing seed6 and AWS-replication weights are reused with their original whole fitting work charged.

Construction	Dev accuracy %	Holdout accuracy %	Holdout NLL	Whole fit GFLOPs	Fitting lineage
Ours: paired timing observed/seed 6	95.31	93.07	0.2230	0.265	Reused; charged
Ours: paired timing observed/seed 7	90.23	86.52	0.3979	0.265	Reused; charged
Ours: paired timing observed/seed 8	91.80	86.04	0.3296	0.265	Reused; charged
Ours: paired timing rank/seed 6	50.00	50.00	0.7039	0.265	Reused; charged
Ours: paired timing rank/seed 7	50.00	50.00	0.7162	0.265	Reused; charged
Ours: paired timing rank/seed 8	50.00	50.00	0.7140	0.265	Reused; charged

Primary gains require the complete three-seed crossed population/pair analysis with correction across two contrasts. Partial scores cannot pass a gate. Synthetic mechanism confirmation is distinct from time-aware dense controls, real-data supremacy and physical energy.

Appendix B. Native frozen confirmation

Completed 1,024-query synthetic holdout evaluations only. All fitted seeds and matched controls remain visible; selected checkpoints use development NLL before confirmation. Existing seed6 and AWS-replication weights are reused with their original whole fitting work charged.

Construction	Dev accuracy %	Holdout accuracy %	Holdout NLL	Whole fit GFLOPs	Fitting lineage
Ours: order S16 shared/seed 6	75.39	79.59	0.8253	0.262	Reused; charged
Ours: order S16 shared/seed 7	44.53	41.80	1.3791	0.262	Reused; charged
Ours: order S16 shared/seed 8	70.31	70.80	0.8565	0.262	Reused; charged
Ours: order S16 private/seed 6	44.14	47.75	1.1264	0.280	Reused; charged
Ours: order S16 private/seed 7	25.78	31.15	1.4322	0.280	New fit
Ours: order S16 private/seed 8	42.97	35.84	1.3393	0.281	New fit

Primary gains require the complete three-seed crossed population/pair analysis with correction across two contrasts. Partial scores cannot pass a gate. Synthetic mechanism confirmation is distinct from time-aware dense controls, real-data supremacy and physical energy.

Appendix B. Persistent-write credit diagnosis

A frozen selected native checkpoint is replayed with alternative delivered content, alternative persistent write, and both. The exact four-corner decomposition separates those effects from their interaction and the current message-linearization residual. The forward architecture and fitted weights remain unchanged.

Ours: frozen audit component	Mean absolute loss effect
Full alternative branch	0.020980
Persistent-write effect	0.012694
Delivered-value linearization residual	0.000288
Delivery/write interaction	0.000066

The two opposed directions in the earlier audit are explained by persistent writes: at event0/block4 the value-only change is +0.001504 but the write-only change is −0.016749; at event7/block7 they are −0.001280 and +0.033918. The route chooses a memory address as well as a message. Training must teach that future state effect.

Twelve fixed probes, one population/address/noise seed, fixed selected-node time. Mean absolute write effect 0.012694 versus value residual 0.000288; these absolute summaries are not additive percentages. Hybrids are diagnostic interventions, not legal proposed routes. Value residual includes nonlinear response and downstream route switches. This is conditional fidelity evidence, not an expected-gradient failure rate.

Charged replay counts:48 forwards,12 backwards,12,288 races; 7.107s wall, 390.3MiB peak RSS. Arithmetic is uninstrumented. Checkpoint weights and outer RNG are preserved. Any new state-aware teacher needs integrated contracts, full accounting and a matched small fit.

Source: [completed factorial audit](#); [joint state/time theory](#).

Appendix B. Ours and boosted trees: banknote

Independent feature-ID rows, state reset between rows, train-only scaling and duplicate-feature group isolation. Ours uses eight native event blocks with parallel heads and content/state mixing; R2 adds temporal reception. Static processing coordinates are not physical asynchronous samples.

Model	Fit/dev rows	Dev NLL / RMSE ↓	Dev accuracy / MAE	Whole fit GFLOPs	Fit MFLOPs/row	Infer MFLOPs/row
Ours R0	128/128	0.1553	0.9531	0.335	0.655	0.139
Ours R2	128/128	0.2700	0.8906	0.379	0.740	0.163
Boosted trees	128/128	0.2319	0.9297	Not counted	Not counted	Not counted

Model	All fit wall s	Peak RSS MiB	Tree nodes / bytes
Ours R0	981.79	445.8	Not applicable
Ours R2	1452.02	453.7	Not applicable
Boosted trees	0.09	404.0	890/49,840

Four development checkpoints or four separately fitted tree candidates; all candidate tree fitting wall time is charged. Neural fit arithmetic is an actual forward/loss/backward/clipping/Adam trace, with specials counted once; preprocessing, evaluation and RNG are separate. Tree FLOPs are unavailable and are not manufactured. Neural wall time includes CPU simulation/audit instrumentation. These are small exploratory development results; reserved test labels are not scored. Strong tabular/frontier superiority requires larger frozen protocols and independent seeds.

Reception ablation: ours R2 reaches 89.06% accuracy / 0.270 NLL versus native R0 95.31% / 0.155. Whole fitting work increases from 0.335 to 0.379 GFLOPs. The added reception capacity has not earned its cost in this single-seed static-data screen.

Confirmation protocol: native checkpoint reuse plus seeds 7/8, original trees, CatBoost and logistic regression; four development selection opportunities per family, then frozen reserved-test scoring. See experiments/AWS_BANKNOTE_CONFIRMATION.md. No pending test score is reported.

Appendix B. Ours and boosted trees: wine_red

Independent feature-ID rows, state reset between rows, train-only scaling and duplicate-feature group isolation. Ours uses eight native event blocks with parallel heads and content/state mixing; R2 adds temporal reception. Static processing coordinates are not physical asynchronous samples.

Model	Fit/dev rows	Dev NLL / RMSE ↓	Dev accuracy / MAE	Whole fit GFLOPs	Fit MFLOPs/row	Infer MFLOPs/row
Ours R0	128/128	0.8238	0.6422	0.816	1.594	0.341
Ours R2	128/128	0.7579	0.6205	0.919	1.795	0.398
Boosted trees	128/128	0.6489	0.4890	Not counted	Not counted	Not counted

Model	All fit wall s	Peak RSS MiB	Tree nodes / bytes
Ours R0	2553.48	468.1	Not applicable
Ours R2	3547.22	485.4	Not applicable
Boosted trees	0.11	404.3	464/25,984

Four development checkpoints or four separately fitted tree candidates; all candidate tree fitting wall time is charged. Neural fit arithmetic is an actual forward/loss/backward/clipping/Adam trace, with specials counted once; preprocessing, evaluation and RNG are separate. Tree FLOPs are unavailable and are not manufactured. Neural wall time includes CPU simulation/audit instrumentation. These are small exploratory development results; reserved test labels are not scored. Strong tabular/frontier superiority requires larger frozen protocols and independent seeds.

Reception helps this regression pilot: ours R2 RMSE 0.758 versus R0 0.824 (8.0% lower), for 0.919 versus 0.816 whole-fit GFLOPs (12.6% more). Trees retain lower RMSE 0.649. This positive within-model effect contrasts with banknote/language reception failures; it is not a cross-family win.

Appendix B. Addressed-state write credit

Completed integrated pilots only: private S4/P0, eight blocks, two independent heads, d8/pool2, 128 fitting queries per pass/four passes, 256 development queries, seed 6. Both retain hard temporal races and winner-only inference. The zero-credit model exactly nests the parent; added memory/time credit is a local surrogate, not an arbitrary unbiased sequence-gradient estimator.

Ours = Sleeping Machines



state credit quality work

Model	Accuracy %	NLL	Whole fit GFLOPs	Fit MFLOPs/query	Infer MFLOPs/query
Ours: write credit	55.08	0.9778	0.277	0.542	0.110
Ours: baseline	54.30	0.9787	0.265	0.518	0.110


Model	Fit/dev/passes	Parameters	State slots	Updates/scores per event
Ours: write credit	128/256/4	42,932	128	16/32
Ours: baseline	128/256/4	42,932	128	16/32

All fitting forward/loss/backward/normalization/clipping/Adam and losing proposals are charged. Special functions have unit weight beside arithmetic; integer discovery/traffic/energy remain separate. Training-only auxiliary state views and whole-process RSS are recorded in each result. Numerical/optimizer prerequisites and accounting smokes are excluded from benchmark plots. Matched gains: 0.78 percentage points / 0.00091 NLL, at 1.045× fitting work. The predeclared larger-fit gate (5 points/.02 NLL/at most2× work) fails. This reuses exploratory development populations; independent seeds and fresh confirmation remain required.

Appendix B. Learned history and credit reach

Frozen checkpoint audit on32 fixed development positions. Both learned models depend on history; the native predictions change even when identical16-character suffixes receive the same race noise. The native model is therefore not strictly a bigram predictor. Equal average loss to a count model calibrates predictive quality; it does not identify learned features or context dependence.

Ours = Sleeping Machines



language learning context

Saved model	Layers with gradients	Layer norm range	Replay tokens	Backwards
Ours: carrier,131K	6	0.419-0.692	3120	1
Ours: native,8K	8	1.027-3.487	3120	1

The actual current count-composition logit gradient matches responsibility-weighted cross entropy to $1.9\text{e-}09$; all six carrier layers receive gradients. Mean responsibility is 3.43% on this 64-target initialization probe. This supports investigating attenuated task signal and conditioning, rather than assuming a general gradient disconnect.

No optimizer steps, weight changes or official-test access. History controls reset state and replay the retained suffix at its absolute positions; per-position native noise is coupled. This 32-position slice is not the full saved development quality or a matched-data model comparison. Gradient norms aggregate different parameter groups; they show reach, not superior conditioning or unbiased hard-route credit. The 64-character history is not uniformly better than 16 on this slice; useful long-range/semantic features remain open. All replays/backwards are counted; arithmetic is uninstrumented. Wall 33.21s, peak RSS 325.5MiB. See [experiments/LANGUAGE_LEARNING_DIAGNOSIS_20261002.md](#).

Appendix B. Credit horizon with fixed context

Completed frozen native8K/H2/d16/depth8 diagnostic. All three arms retain the same 128-token development context, score the same last 16 targets and receive identical per-position race noise. Only graph reach changes: 16, 32 or 64 tokens. Predictions match exactly; the gradient changes, separating retained information from the credit used to learn how to retain it.

Graph horizon	Slice bpc	Gradient norm	Graph tokens	Replay tokens
16	4.016190	7.0827	16	128
32	4.016190	6.3785	32	128
64	4.016190	6.6635	64	128

Short/long credit	Difference norm	Relative difference %	Gradient cosine
16 versus 32	1.7577	27.56	0.9713
16 versus 64	2.0581	30.89	0.9570
32 versus 64	1.0092	15.15	0.9890


The 16-versus-64 gradient difference has norm 30.89% of the 64-token gradient, with cosine 0.9570. The 32-versus-64 difference is 15.15%. Every layer receives credit, and short-credit norms can be larger because omitted contributions can cancel retained ones. Credit truncation materially changes the gradient on this probe. The next test refits matched models at each credit horizon.

Frozen weights; zero optimizer steps and no official-test access. One 16-target slice, 384 replay tokens, 112 graph tokens, 3 backwards. Wall 11.08s, peak RSS 518.0MiB. Arithmetic is uninstrumented. Native hard-route credit is a surrogate; 64 tokens is a comparison, not an all-history unbiased reference. Relative difference divides the norm of the gradient difference by the longer-credit gradient norm, not a percentage of predictive quality or retained features.

Appendix B. Banknote: no confirmed advantage

Competitive accuracy, no confirmed advantage: the three-seed reserved test does not sustain the development lead. Ours averages 91.8% versus 94.0% for the original trees. Logistic regression has lower test log loss. The full four-family confirmation remains incomplete. The reserved-test accuracy point estimates are close, but no equivalence margin was specified. Parity is therefore a descriptive reading, not a proven equivalence claim.

Ours = Sleeping Machines



banknote reserved test

Model	Completed seeds	Mean test accuracy %	Mean test NLL
Ours native	3	91.81	0.2353
Boosted trees	3	93.95	0.2065
CatBoost	2/3	Pending	Pending
Logistic	3	94.66	0.0906

For the completed ours/tree comparison, the paired accuracy difference is -2.14 percentage points (descriptive 95% crossed seed/feature-group interval -6.90 to $+2.43$). The control-minus-ours NLL difference is -0.0287 (98.33% interval -0.1738 to $+0.1277$). Logistic regression improves NLL by 0.1447 (98.33% interval 0.0338 to 0.2660). These three-seed intervals are approximate and share one test split.

Partial analysis: local_banknote_partial_confirmation_20261002T013000Z.json; 11/12 final cells, 4000 bootstrap draws, 270 feature groups, three seeds. NLL intervals allow for three control comparisons; accuracy intervals are descriptive. This does not replace the incomplete full four-family gate. The original 95.3% versus 93.0% development screen and all per-seed work remain below.

Appendix B. Frozen banknote confirmation

Completed test scores only. Checkpoints/candidates were selected on 128 development rows after fitting 128 rows. Reserved feature groups were scored after choices were frozen; all rows start with cold state. The per-seed ledger preserves completed scores; the summary identifies families with all three seeds.

Model/seed	Dev NLL	Test NLL	Test accuracy%	Whole fit GFLOPs	Fit MFLOPs/row	Infer MFLOPs/row
Ours native/s6	0.1553	0.1426	94.31	0.335	0.655	0.139
Ours native/s7	0.2594	0.2494	90.75	0.335	0.655	0.139
Ours native/s8	0.3270	0.3138	90.39	0.335	0.655	0.139
Boosted trees/s6	0.2319	0.2065	93.95	Not counted	Not counted	Not counted
Boosted trees/s7	0.2319	0.2065	93.95	Not counted	Not counted	Not counted
Boosted trees/s8	0.2319	0.2065	93.95	Not counted	Not counted	Not counted

Model/seed	Charged fit wall s	Test wall s	Peak RSS MiB	Fit provenance
Ours native/s6	981.79	4.54	334.3	Historical fit reused
Ours native/s7	976.56	4.65	447.0	New fit
Ours native/s8	989.25	4.61	447.0	New fit
Boosted trees/s6	0.09	0.00	434.3	New fit

Model/seed	Charged fit wall s	Test wall s	Peak RSS MiB	Fit provenance
Boosted trees/s7	0.09	0.00	434.9	New fit
Boosted trees/s8	0.09	0.00	434.9	New fit

Native forward/loss/backward/clipping/Adam are traced. Seed6 reuse retains its original full fitting charge; it adds no optimizer steps. Control fitting includes all four independent candidates. Their FLOPs are unavailable. Audit instrumentation, preprocessing and physical energy are separate. Repeated seeds share test rows and must not be pooled as independent observations. Paired seed/feature-group analysis and all three prespecified seeds are required for the confirmation claim.

Appendix B. Frozen banknote confirmation

Completed test scores only. Checkpoints/candidates were selected on 128 development rows after fitting 128 rows. Reserved feature groups were scored after choices were frozen; all rows start with cold state. The per-seed ledger preserves completed scores; the summary identifies families with all three seeds.


Model/seed	Dev NLL	Test NLL	Test accuracy%	Whole fit GFLOPs	Fit MFLOPs/row	Infer MFLOPs/row
CatBoost/s6	0.1023	0.1207	93.95	Not counted	Not counted	Not counted
CatBoost/s7	0.1042	0.1078	95.37	Not counted	Not counted	Not counted
Logistic/s6	0.0832	0.0906	94.66	Not counted	Not counted	Not counted
Logistic/s7	0.0832	0.0906	94.66	Not counted	Not counted	Not counted
Logistic/s8	0.0832	0.0906	94.66	Not counted	Not counted	Not counted

Model/seed	Charged fit wall s	Test wall s	Peak RSS MiB	Fit provenance
CatBoost/s6	0.35	0.00	461.6	New fit
CatBoost/s7	0.35	0.00	460.8	New fit
Logistic/s6	0.01	0.00	429.6	New fit
Logistic/s7	0.01	0.00	429.7	New fit
Logistic/s8	0.01	0.00	429.9	New fit

Native forward/loss/backward/clipping/Adam are traced. Seed6 reuse retains its original full fitting charge; it adds no optimizer steps. Control fitting includes all four independent candidates. Their FLOPs are unavailable. Audit instrumentation, preprocessing and physical energy are separate. Repeated seeds share test rows and must not be pooled as independent observations. Paired seed/feature-group analysis and all three prespecified seeds are required for the confirmation claim.

Appendix B (continued). Ours: language work as scaling develops

This ledger updates from completed integrated-model stages. It shows the emerging work advantage alongside its quality and data budget. Per-target fitting work removes the difference in the number of presentations. Quality and data budgets are shown alongside work for each row.



integrated language work progress

Model	Fit / passes	bpc / split ↓	Whole fit GFLOPs ↓	Fitting MFLOPs / target ↓	Forward MFLOPs / position ↓
Ours / d32 / p2	32,768 / 4	3.106 / dev	114.247	0.872	0.089
Ours / d16 / p2	32,768 / 4	3.121 / dev	31.053	0.237	0.026
Ours / d16 / p2	8,192 / 4	3.398 / dev	7.788	0.238	0.025
Ours / d16 / p4	8,192 / 4	3.426 / dev	14.721	0.449	0.032
LSTM / width 512	10M / six	1.799 / test	432,592.997	7.210	2.402
Transformer / width 256	10M / four	1.908 / test	888,775.443	22.223	7.405

Compare within a column: whole-fit totals use GFLOPs for every model; per-target and forward work use MFLOPs for every model. One GFLOP is 1,000 MFLOPs. Whole-fit totals also depend on the number of training presentations; the per-target column divides that out.

All table values, figures and ratios use arithmetic plus one operation per special function, matching the historical neural estimate convention. This is not a physical energy cost. Ours arithmetic-only whole-fit totals (GFLOPs): 32,768 / pool 2: 110.088; 32,768 / pool 2: 29.883; 8,192 / pool 2: 7.492; 8,192 / pool 4: 14.152. Separate special-function counts are preserved in each result.

The raw work gap is substantial. The completed 32,768-character integrated stage's representative forward estimate is **84× smaller** than the larger saved Transformer estimate; fitting work per target is **25× smaller**. These are configuration-level work ratios. Our development score and the reference official test score use different targets and data budgets. The gap is not a matched-quality supremacy claim.


Ours: d denotes payload width and p pool size; fixed character pools and event depths. Each result records its validation interval and credit horizon. References: width-512 LSTM or four width-256 Transformer layers, 256-position fitting chunks and 999,999 aligned official test targets. Ours uses representative operator traces including counterfactual credit, backward, clipping and Adam; neural references use shape formulas and backward $\approx$ twice forward. RNG, indexing, memory traffic and evaluation passes are additional. Same-quality and iso-FLOP conclusions await comparable completed runs.

Appendix B (continued). Ours and neural controls: accuracy versus FLOPs

Each point is a completed model, not a projected scaling law. Left: ours on cold development characters, with integrated models and earlier carrier controls labelled separately. The new 2K screens score 2,047 development targets; the earlier ladders score 8,191. Right: saved neural test results INCLUDING the later segment-batched native width, depth, pool and credit models (purple diamonds, NB labels). Lower bpc means better prediction; lower fitting work means fewer estimated operations. No curve is drawn between different model families or scoring splits.

Ours = Sleeping Machines

Quality versus fitting work; scoring protocols and data budgets differ



language quality vs work


Model type	Fitting budget	bpc / split ↓	Whole fit GFLOPs ↓	Fitting MFLOPs / target ↓
Ours: integrated d32/p2	32,768 / 4 passes	3.106 / dev	114.247	0.872
Ours: batched native p96/d4 + route credit, 6 passes	10,000,000 / 6 passes	1.888 / test	352,078.096	5.868
Ours: carrier w128g	1,048,576 / 4 passes	2.210 / dev	8,373.302	1.996
LSTM: 512	90,000,000 / 6 passes	1.661 / test	3,893,396.042	7.210
Transformer: 256x4	90,000,000 / 4 passes	1.604 / test	8,000,253.349	22.223

The table selects the largest fitting budget currently completed for each family; the best score breaks ties. Point numbers refer to the following variant ledger, which lists all plotted variants. Variant labels: I = ours integrated payload/pool/data; IKV adds per-position race memory (S uses the content index); NB = later batched native, T128/T256 evaluation shown; C = ours carrier width/data (g means content gates); L = LSTM width/data; T = Transformer width x layers/data; s denotes seed. K is 1,024 characters in ours labels; M is decimal million in neural labels.

Estimates include learning, clipping and Adam, with unit-weight special functions. Ours uses representative operator traces; neural controls use shape formulas and backward approximately twice forward. Scoring splits, data, passes, capacity and credit differ; these panels are evidence inventories, not an iso-FLOP or equal-quality benchmark.

Appendix B (continued). Later native language: quality versus fitting work

Focused view of the later credited width/depth/capacity models. Same completed T256 scores as the common inventory; nominal10M fitting characters, one pass and saved one-pass controls. Blue is alternative-value route credit; light blue is timing-only credit; gray is a dense control.



latest native language fitting

Model	T256 test bpc	Whole fit TFLOPs est.	Fit MFLOPs / input position est.
p32/d4 + route credit	2.3715	7.24	0.72
p32/d4/pool4 + route credit	2.3452	11.95	1.20
p64/d4 + route credit	2.1833	26.79	2.68
p64/d4/pool4 + route credit	2.1795	44.07	4.41
p96/d4 + route credit	2.1625	58.65	5.87
LSTM-256	2.1706	20.31	2.03
Transformer-256x2	2.4269	111.26	11.13
Transformer-256x4, 4 passes	1.9083	888.78	22.22
LSTM-512, 6 passes	1.7993	432.59	7.21


Native quality comes from the compiled training evaluator, fitting work from representative full-step traces; controls use shape estimates. Every column has the same units and denominator for ours and controls. Native random-segment fitting and evaluation tail coverage differ from controls. Single seeds; trained sparse-backend rescore and modern replications remain open. p96 now slightly exceeds LSTM quality with more fitting work. All13 later native fits, including timing-only/pool1/write-credit history, remain in the common graph and following ledger.

Appendix B (continued). Accuracy versus inference FLOPs

Inference predicts with frozen weights: no backward pass, clipping or optimizer update. These are the same completed checkpoints, quality scores and point IDs as the fitting graph. Ours uses saved forward operator traces; the integrated models read only winning values. Later native NB points use emulator traces, charged for evaluated warm positions per scored target; winner-only estimates remain in the native appendix pending trained parity/rescore. LSTM and Transformer costs use shape estimates. Development and test evidence remain separate.

Ours = Sleeping Machines

Quality versus inference work; estimates, scoring protocols and data differ



language quality vs inference

Model type	bpc / split ↓	Inference MFLOPs / character ↓	Cost boundary
Ours: integrated d32/p2	3.106 / dev	0.0885	Winner-only inference trace
Ours: batched native p96/d4 + route credit, 6 passes	1.888 / test	3.8929	Emulator trace × evaluated positions/scored targets
Ours: carrier w128g	2.210 / dev	0.6389	Saved forward operator trace
LSTM: 512	1.661 / test	2.4024	Recurrent shape estimate
Transformer: 256x4	1.604 / test	14.8104	Overlapping-window shape estimate

Solid Transformer points estimate its saved 256-position scorer: full windows advanced by 128 positions, approximately two forward positions per scored character (boundary/tail overhead omitted). Hollow points show a hypothetical one-step decode with cached keys/values and 256 available positions, using $L(24d^2 + 4Td + 30d + 20T) + 54d + 135$ unit-weight operations, $T = 256$. No cached decoder was run. Its plotted bpc belongs to the saved window scorer; learned positions reset between windows, so cache reuse has not been shown to preserve those scores.

Two FLOPs per multiply-add; special functions count as one operation. Ours traces include numerical clocks and loss scoring; neural elementwise overhead is approximate. Traces are representative warm-state costs, not full-stream measurements; growing KV candidate occupancy can change work. The separate KV pages also show projected event-architecture costs that remove numerical clock simulation. RNG, indexing, memory traffic and physical race energy are additional. These are work estimates, not latency or joules, and differing data, quality and evaluation protocols prevent a supremacy conclusion.

Appendix B (continued). Completed language variants and work

Variant	Params K	Fit / passes	bpc / split ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
1. Ours: AWS I32/p2/32K/s6	1,388.9	32,768 / 4	3.106 / dev	114.247	0.872	0.0885
2. Ours: I16/p2/32K/s6	361.4	32,768 / 4	3.121 / dev	31.053	0.237	0.0255
3. Ours: I16/p2/8K/s6	361.4	8,192 / 4	3.398 / dev	7.788	0.238	0.0255
4. Ours: I16/p4/8K/s6	720.0	8,192 / 4	3.426 / dev	14.721	0.449	0.0321
5. Ours: IHR2x32/b64/u64@0.002D8/2K/s6	3,819.5	2,048 / 4	3.740 / dev	23.019	2.811	0.5060
6. Ours: IHR2x32D8/2K/s6	3,819.5	2,048 / 4	3.786 / dev	27.731	3.387	0.5059
7. Ours: IHR2x32/u128@0.004D8/2K/s6	3,819.5	2,048 / 4	3.779 / dev	19.999	2.442	0.5058
8. Ours: IHR2x32/u64@0.002D8/2K/s6	3,819.5	2,048 / 4	3.733 / dev	22.753	2.779	0.5059
9. Ours: IHR2x32/u64@0.004D8/2K/s6	3,819.5	2,048 / 4	3.800 / dev	22.750	2.779	0.5060
10. Ours: IHR2x32/u128@0.004D8/8K/s6	3,819.5	8,192 / 4	3.485 / dev	79.953	2.440	0.5059
11. Ours: IHR2x32/u64@0.002D8/8K/s6	3,819.5	8,192 / 4	3.490 / dev	89.999	2.747	0.5058
12. Ours: IHR4x32/u128@0.004D8/2K/s6	7,778.3	2,048 / 4	3.981 / dev	48.431	5.915	1.2898

Each row retains its original architecture, fitting budget and score. The selected 10M LSTM/Transformer rows use the aligned 999,999-target scores; other neural rows retain their original E64 test scorers. The 90M LSTM uses its saved recurrent scoring protocol. Carrier and integrated development scores use frozen evaluation; integrated official scores appear only after their full test completes. Validation/test work, RNG and physical traffic are outside fitting totals. Sources: E64/E174, saved AWS E64 results and the completed parallel_language, episodic_language and language_batched JSON records. Later NB rows use T256 when completed (first v1 stays T128), actual fitting presentations and native window overlap charged per scored target; different tail coverage is retained. The global ledger uses emulator floating arithmetic consistently; fitting work per target divides by actual training target presentations. The separate KV page reports architectural projections. No new dense model was trained.

Appendix B (continued). Completed language variants and work

Variant	Params K	Fit / passes	bpc / split ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
13. Ours: IHR4x32/u128@0.004D8/8K/s6	7,778.3	8,192 / 4	3.543 / dev	193.751	5.914	1.2882
14. Ours: IKV32D6/2K/s6	1,413.6	2,048 / 4	3.620 / dev	9.035	1.103	0.1466
15. Ours: IKV32D8/2K/s6	1,883.9	2,048 / 4	3.539 / dev	11.992	1.465	0.1941
16. Ours: IKVS32D8/2K/s6	1,883.9	2,048 / 4	3.554 / dev	11.999	1.465	0.1968
17. Ours: IKVS32D8/8K/s6	1,883.9	8,192 / 4	3.357 / dev	50.006	1.526	0.1969
18. Ours: I32D6/2K/s6	1,388.9	2,048 / 4	3.633 / dev	7.207	0.880	0.0890
19. Ours: I32D8/2K/s6	1,850.9	2,048 / 4	3.542 / dev	9.557	1.167	0.1173
20. Ours: I32D8/8K/s6	1,850.9	8,192 / 4	3.311 / dev	40.243	1.228	0.1173
21. Ours: IHR2x32/m2/u128@0.004D8/2K/s6	3,819.5	2,048 / 4	3.820 / dev	20.495	2.503	0.5165
22. Ours: IHR2x32/m4/u128@0.004D8/2K/s6	3,819.5	2,048 / 4	3.771 / dev	20.914	2.554	0.5268
23. Ours: IHR2x32/wc0.25/u64@0.002D8/2K/s6	3,819.5	2,048 / 4	3.722 / dev	23.469	2.866	0.5059
24. Ours: IHR2x32/wc1/u64@0.002D8/2K/s6	3,819.5	2,048 / 4	3.724 / dev	23.464	2.866	0.5058

Each row retains its original architecture, fitting budget and score. The selected 10M LSTM/Transformer rows use the aligned 999,999-target scores; other neural rows retain their original E64 test scorers. The 90M LSTM uses its saved recurrent scoring protocol. Carrier and integrated development scores use frozen evaluation; integrated official scores appear only after their full test completes. Validation/test work, RNG and physical traffic are outside fitting totals. Sources: E64/E174, saved AWS E64 results and the completed parallel_language, episodic_language and language_batched JSON records. Later NB rows use T256 when completed (first v1 stays T128), actual fitting presentations and native window overlap charged per scored target; different tail coverage is retained. The global ledger uses emulator floating arithmetic consistently; fitting work per target divides by actual training target presentations. The separate KV page reports architectural projections. No new dense model was trained.

Appendix B (continued). Completed language variants and work

Variant	Params K	Fit / passes	bpc / split ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
25. Ours: Native H2d16/p2/2K/s6	54.9	2,048 / 4	3.765 / dev	3.778	0.461	0.0974
26. Ours: Native H2d16/p2/8K/s6	54.9	8,192 / 4	3.557 / dev	15.116	0.461	0.0974
27. Ours: R2/uniform/reception/2K/s6	57.1	2,048 / 4	3.796 / dev	4.032	0.492	0.1056
28. Ours: R4/uniform/reception/2K/s6	59.3	2,048 / 4	3.795 / dev	4.183	0.511	0.1109
29. Ours: R4/late/reception/2K/s6	57.1	2,048 / 4	3.795 / dev	3.979	0.486	0.1041
30. Ours: R4/uniform/waiting/2K/s6	59.3	2,048 / 4	3.764 / dev	3.897	0.476	0.1007
31. Ours: C128/128K	308.0	131,072 / 4	2.643 / dev	1,032.197	1.969	0.6296
32. Ours: C32/128K	21.7	131,072 / 4	2.858 / dev	77.197	0.147	0.0470
33. Ours: C64/128K	80.3	131,072 / 4	2.727 / dev	274.735	0.524	0.1675
34. Ours: C128g/128K	309.6	131,072 / 4	2.587 / dev	1,046.656	1.996	0.6389
35. Ours: C256g/128K	1,208.9	131,072 / 4	2.572 / dev	4,025.494	7.678	2.4571
36. Ours: C128g/1024K	309.6	1,048,576 / 4	2.210 / dev	8,373.302	1.996	0.6389

Each row retains its original architecture, fitting budget and score. The selected 10M LSTM/Transformer rows use the aligned 999,999-target scores; other neural rows retain their original E64 test scorers. The 90M LSTM uses its saved recurrent scoring protocol. Carrier and integrated development scores use frozen evaluation; integrated official scores appear only after their full test completes. Validation/test work, RNG and physical traffic are outside fitting totals. Sources: E64/E174, saved AWS E64 results and the completed parallel_language, episodic_language and language_batched JSON records. Later NB rows use T256 when completed (first v1 stays T128), actual fitting presentations and native window overlap charged per scored target; different tail coverage is retained. The global ledger uses emulator floating arithmetic consistently; fitting work per target divides by actual training target presentations. The separate KV page reports architectural projections. No new dense model was trained.

Appendix B (continued). Completed language variants and work

Variant	Params K	Fit / passes	bpc / split ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
37. L256/1M	338.4	1,000,000 / 20	2.179 / test	40,628.875	2.032	0.6770
38. L256/10M	338.4	10,000,000 / 1	2.171 / test	20,306.115	2.032	0.6770
39. L512/10M	1,199.3	10,000,000 / 6	1.799 / test	432,592.997	7.210	2.4024
40. T112x8/1M	1,250.6	1,000,000 / 5	2.352 / test	51,107.144	10.223	6.7999
41. T256x2/1M	1,658.9	1,000,000 / 20	2.367 / test	222,614.402	11.133	7.4192
42. T256x2/10M	1,658.9	10,000,000 / 1	2.427 / test	111,261.602	11.133	7.4192
43. T256x4/10M	3,238.4	10,000,000 / 4	1.908 / test	888,775.443	22.223	14.8104
44. L512/90M	1,199.3	90,000,000 / 6	1.661 / test	3,893,396.042	7.210	2.4024
45. T256x4/90M	3,238.4	90,000,000 / 4	1.604 / test	8,000,253.349	22.223	14.8104
46. Ours: NB p16/d8, v1 (610 updates)/T128	54.9	10,000,000 / 1	2.899 / test	4,042.822	0.405	0.2629
47. Ours: NB p16/d8, skip2/T256	54.9	10,000,000 / 1	2.719 / test	4,047.897	0.405	0.2629
48. Ours: NB p32/d4/T256	108.9	10,000,000 / 1	2.506 / test	7,220.754	0.722	0.4792

Each row retains its original architecture, fitting budget and score. The selected 10M LSTM/Transformer rows use the aligned 999,999-target scores; other neural rows retain their original E64 test scorers. The 90M LSTM uses its saved recurrent scoring protocol. Carrier and integrated development scores use frozen evaluation; integrated official scores appear only after their full test completes. Validation/test work, RNG and physical traffic are outside fitting totals. Sources: E64/E174, saved AWS E64 results and the completed parallel_language, episodic_language and language_batched JSON records. Later NB rows use T256 when completed (first v1 stays T128), actual fitting presentations and native window overlap charged per scored target; different tail coverage is retained. The global ledger uses emulator floating arithmetic consistently; fitting work per target divides by actual training target presentations. The separate KV page reports architectural projections. No new dense model was trained.


Appendix B (continued). Completed language variants and work

Variant	Params K	Fit / passes	bpc / split ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
49. Ours: NB p32/d8, skip2/T256	210.0	10,000,000 / 1	2.456 / test	14,107.373	1.412	0.9253
50. Ours: NB p32/d8/pool4, skip2/T256	346.3	10,000,000 / 1	2.498 / test	23,603.624	2.362	1.5409
51. Ours: NB p32/d4/pool1 (control: no selection)/T256	74.8	10,000,000 / 1	2.439 / test	4,878.223	0.488	0.3253
52. Ours: NB p32/d4 + route credit/T256	108.9	10,000,000 / 1	2.371 / test	7,242.661	0.725	0.4792
53. Ours: NB p32/d4/pool4 + route credit/T256	177.0	10,000,000 / 1	2.345 / test	11,949.789	1.196	0.7870
54. Ours: NB p32/d4 + read and write credit/T256	108.9	10,000,000 / 1	2.384 / test	7,279.888	0.728	0.4792
55. Ours: NB p64/d4 + route credit/T256	422.5	10,000,000 / 1	2.183 / test	26,787.462	2.680	1.7765
56. Ours: NB p32/d8, skip2 + route credit/T256	210.0	10,000,000 / 1	2.326 / test	14,151.188	1.416	0.9253
57. Ours: NB p64/d4/pool4 + route credit/T256	689.8	10,000,000 / 1	2.179 / test	44,073.829	4.410	2.9157
58. Ours: NB p96/d4 + route credit/T256	940.9	10,000,000 / 1	2.162 / test	58,647.635	5.868	3.8929
59. Ours: NB p64/d4 + route credit, 4 passes/T256	422.5	10,000,000 / 4	1.955 / test	107,193.761	2.680	1.7765
60. Ours: NB p96/d4 + route credit, 6 passes/T256	940.9	10,000,000 / 6	1.888 / test	352,078.096	5.868	3.8929

Each row retains its original architecture, fitting budget and score. The selected 10M LSTM/Transformer rows use the aligned 999,999-target scores; other neural rows retain their original E64 test scorers. The 90M LSTM uses its saved recurrent scoring protocol. Carrier and integrated development scores use frozen evaluation; integrated official scores appear only after their full test completes. Validation/test work, RNG and physical traffic are outside fitting totals. Sources: E64/E174, saved AWS E64 results and the completed parallel_language, episodic_language and language_batched JSON records. Later NB rows use T256 when completed (first v1 stays T128), actual fitting presentations and native window overlap charged per scored target; different tail coverage is retained. The global ledger uses emulator floating arithmetic consistently; fitting work per target divides by actual training target presentations. The separate KV page reports architectural projections. No new dense model was trained.

Appendix B (continued). Ours: per-position race KV memory

Both integrated models fit 2,048 characters for 4 passes, with payload 32, 6 sparse receiver depths, seed 6 and 2,047 identical cold development targets. The KV arm retains separate historical keys and values at every depth; learned queries select one value through time. Incoming content is retained and gated with the retrieved message. No dense carrier is added.



episodic language comparison D2048 depth6 character

Ours: memory	Dev bpc ↓	Projected fit GFLOPs ↓	CPU fit GFLOPs ↓	Projected forward MFLOPs/char ↓
receiver	3.633	7.206	7.207	0.0890
kv	3.620	9.034	9.035	0.1464


Completed KV improvement over receiver memory: +0.013 bpc (positive is better). The index admits up to 8 recent matching-character entries plus 4 recent positions. Older entries outside these tails cannot be addressed by this index. Duplicates are removed. Development averages 11.36 keys scored and one value delivered per retrieval query. All 12,282 entries remain stored (3.00 MiB raw keys/values); the oldest selected entry is 1,829 characters old. This bounds reads, not stored history.

Physical clock competition replaces explicit numerical rate exponentiation and noise/rate division plus the bounded-delay simulation in the projected ledger. Query/key/value maps, scored candidates, gated content, backward, counterfactual teaching, clipping and actual Adam remain charged. Counts are representative first/mature/partial traces; special functions have unit weight here and are separate in JSON. Physical rate setting, clock circuits, index/address operations, RNG and traffic need their own implementation costs; FLOPs do not certify energy.

Temporal races avoid the explicit normalizing reduction/division and deliver one value at inference; training reads all admitted values for route credit. The orange bar is an analytical same-shortlist aggregation comparison, not another trained model. Candidate coverage is approximate and does not guarantee full-bank attention equivalence. Random-hyperplane indexing is an established primitive (Charikar, STOC 2002); novelty is not claimed for this index. Historical activations are detached at the credit boundary and are not recomputed after parameter updates. One seed and a small data budget; no equal-quality Transformer or frontier claim.

Appendix B (continued). Ours: per-position race KV memory

Both integrated models fit 2,048 characters for 4 passes, with payload 32, 8 sparse receiver depths, seed 6 and 2,047 identical cold development targets. The KV arm retains separate historical keys and values at every depth; learned queries select one value through time. Incoming content is retained and gated with the retrieved message. No dense carrier is added.



episodic language comparison D2048 depth8 character

Ours: memory	Dev bpc ↓	Projected fit GFLOPs ↓	CPU fit GFLOPs ↓	Projected forward MFLOPs/char ↓
receiver	3.542	9.556	9.557	0.1172
kv	3.539	11.990	11.992	0.1939


Completed KV improvement over receiver memory: +0.003 bpc (positive is better). The index admits up to 8 recent matching-character entries plus 4 recent positions. Older entries outside these tails cannot be addressed by this index. Duplicates are removed. Development averages 11.36 keys scored and one value delivered per retrieval query. All 16,376 entries remain stored (4.00 MiB raw keys/values); the oldest selected entry is 1,829 characters old. This bounds reads, not stored history.

Physical clock competition replaces explicit numerical rate exponentiation and noise/rate division plus the bounded-delay simulation in the projected ledger. Query/key/value maps, scored candidates, gated content, backward, counterfactual teaching, clipping and actual Adam remain charged. Counts are representative first/mature/partial traces; special functions have unit weight here and are separate in JSON. Physical rate setting, clock circuits, index/address operations, RNG and traffic need their own implementation costs; FLOPs do not certify energy.

Temporal races avoid the explicit normalizing reduction/division and deliver one value at inference; training reads all admitted values for route credit. The orange bar is an analytical same-shortlist aggregation comparison, not another trained model. Candidate coverage is approximate and does not guarantee full-bank attention equivalence. Random-hyperplane indexing is an established primitive (Charikar, STOC 2002); novelty is not claimed for this index. Historical activations are detached at the credit boundary and are not recomputed after parameter updates. One seed and a small data budget; no equal-quality Transformer or frontier claim.

Appendix B (continued). Ours: per-position race KV memory

Both integrated models fit 2,048 characters for 4 passes, with payload 32, 8 sparse receiver depths, seed 6 and 2,047 identical cold development targets. The KV arm retains separate historical keys and values at every depth; learned queries select one value through time. Incoming content is retained and gated with the retrieved message. No dense carrier is added.



episodic language comparison D2048 depth8 semantic

Ours: memory	Dev bpc ↓	Projected fit GFLOPs ↓	CPU fit GFLOPs ↓	Projected forward MFLOPs/char ↓
receiver	3.542	9.556	9.557	0.1172
kv	3.554	11.997	11.999	0.1965

Completed KV improvement over receiver memory: -0.012 bpc (positive is better). The content index uses three random-hyperplane bits of learned keys/queries, with up to 8 recent/full-history samples in the query bucket and its one-bit neighbors, plus 4 recent positions. Duplicates are removed. Development averages 10.07 keys scored and one value delivered per retrieval query. All 16,376 entries remain stored (4.00 MiB raw keys/values); the oldest selected entry is 2,026 characters old. This bounds reads, not stored history.


Physical clock competition replaces explicit numerical rate exponentiation and noise/rate division plus the bounded-delay simulation in the projected ledger. Query/key/value maps, scored candidates, gated content, backward, counterfactual teaching, clipping and actual Adam remain charged. Counts are representative first/mature/partial traces; special functions have unit weight here and are separate in JSON. Physical rate setting, clock circuits, index/address operations, RNG and traffic need their own implementation costs; FLOPs do not certify energy.

Temporal races avoid the explicit normalizing reduction/division and deliver one value at inference; training reads all admitted values for route credit. The orange bar is an analytical same-shortlist aggregation comparison, not another trained model. Candidate coverage is approximate and does not guarantee full-bank attention equivalence. Random-hyperplane indexing is an established primitive (Charikar, STOC 2002); novelty is not claimed for this index. Historical activations are detached at the credit boundary and are not recomputed after parameter updates. One seed and a small data budget; no equal-quality Transformer or frontier claim.

Appendix B (continued). Ours: per-position race KV memory

Both integrated models fit 8,192 characters for 4 passes, with payload 32, 8 sparse receiver depths, seed 6 and 8,191 identical cold development targets. The KV arm retains separate historical keys and values at every depth; learned queries select one value through time. Incoming content is retained and gated with the retrieved message. No dense carrier is added.

Ours = Sleeping Machines



episodic language comparison D8192 depth8 semantic

Ours: memory	Dev bpc ↓	Projected fit GFLOPs ↓	CPU fit GFLOPs ↓	Projected forward MFLOPs/char ↓
receiver	3.311	40.241	40.243	0.1173
kv	3.357	49.998	50.006	0.1967

Completed KV improvement over receiver memory: -0.047 bpc (positive is better). The content index uses three random-hyperplane bits of learned keys/queries, with up to 8 recent/full-history samples in the query bucket and its one-bit neighbors, plus 4 recent positions. Duplicates are removed. Development averages 10.40 keys scored and one value delivered per retrieval query. All 65,528 entries remain stored (16.00 MiB raw keys/values); the oldest selected entry is 8,135 characters old. This bounds reads, not stored history.


Physical clock competition replaces explicit numerical rate exponentiation and noise/rate division plus the bounded-delay simulation in the projected ledger. Query/key/value maps, scored candidates, gated content, backward, counterfactual teaching, clipping and actual Adam remain charged. Counts are representative first/mature/partial traces; special functions have unit weight here and are separate in JSON. Physical rate setting, clock circuits, index/address operations, RNG and traffic need their own implementation costs; FLOPs do not certify energy.

Temporal races avoid the explicit normalizing reduction/division and deliver one value at inference; training reads all admitted values for route credit. The orange bar is an analytical same-shortlist aggregation comparison, not another trained model. Candidate coverage is approximate and does not guarantee full-bank attention equivalence. Random-hyperplane indexing is an established primitive (Charikar, STOC 2002); novelty is not claimed for this index. Historical activations are detached at the credit boundary and are not recomputed after parameter updates. One seed and a small data budget; no equal-quality Transformer or frontier claim.

Appendix B (continued). Ours: cheaper learning updates

Credit still propagates over 16-character segments. Gradients are summed over U targets, normalized by their actual count, clipped once and used for one Adam update. All models here use two independent heads, payload 32/head, eight blocks, seed 6, four passes over 2,048 fitting characters and 8,191 frozen development targets.

Ours = Sleeping Machines



parallel optimizer work

Ours: U / learning rate	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Adam + clip GFLOPs ↓
U16 / 0.001	3.786	27.731	3.387	10.521
U64 / 0.002	3.733	22.753	2.779	5.225
U64 / 0.004	3.800	22.750	2.779	5.225
U128 / 0.004	3.779	19.999	2.442	2.645


The cheapest schedule within the declared 0.05 bpc tolerance of the best pilot uses U128, lr 0.004: 3.779 bpc and 27.9% less whole fitting work than reference. The best quality is 3.733 bpc. These completed results support optimizer amortization in this configuration, not language-model supremacy.

Learning rates and warmup differ across configurations, so this is not an isolated optimizer-interval ablation. Selected checkpoints minimize frozen development loss over the fixed four passes; all fitting work remains charged. Each result is one seed. Candidate scoring, losing-value credit, backward and gradient accumulation/normalization remain in the operator ledger. Larger-data and repeat-seed comparisons must establish transfer of the selected schedule.

Appendix B (continued). Ours: shared-match temporal arrivals

2 independent spatial heads, payload 32/head, 8 blocks. All rows fit 2,048 characters for 4 passes and score 8,191 cold development targets. Each query forms its candidate matches once. Multiple temporal marks reuse those rates; only the winning emitter renews its clock. The receiver and each selected message evolve until the last local read, then the messages are averaged and gated.

Ours = Sleeping Machines



repeated arrival quality work 1df4adaf89

Ours delivers 4 times as many historical winner messages for 4.57% additional whole fitting arithmetic in this completed screen. Quality changes from 3.779 to 3.771 bpc. The additional deliveries reuse the expensive match.

Ours: arrivals / head	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
m=1	3.779	19.999	2.442	0.5058
m=2	3.820	20.495	2.503	0.5165
m=4	3.771	20.914	2.554	0.5268


Candidate discovery and independent Q/K/V projections are retained. Multiple marks reuse one rate setting within each query; trained scores and candidate trajectories can differ across runs. More marks can retrieve the same value; they do not create extra learned spatial heads or discover absent candidates. Training reads all admitted values once and aggregates the conserved per-arrival teacher in $O(Cd + md)$. All delivered messages, temporal transports, backward, clipping and Adam remain charged.

Completed single-seed development screens; m=1 reuses the saved reference under exact nesting contracts. Unit-weight special functions are included; CPU minimum comparisons/RNG and memory traffic are separate counters. The bounded numerical time encoding is not a demonstrated homogeneous physical Poisson clock. This local counterfactual teacher is a declared surrogate, not an exact gradient through nonlinear route changes. No matched-quality dense-model, physical-energy or frontier superiority is inferred.

Appendix B (continued). Ours: independent temporal heads


Each head has its own receiver pool, historical bank and query/key/value/gate matrices. A winning content vector evolves through learned rotation and decay until its channel is read. The next block reads at the latest parallel arrival, preserves each channel and learns their mix. Heads need not arrive simultaneously. This is implemented in the CPU emulator; execution there is serial.

Ours = Sleeping Machines



parallel temporal heads

Ours = Sleeping Machines



parallel temporal head pilots


Ours: heads / update	Fit / passes	Dev bpc ↓	Whole fit GFLOPs ↓
H2 / U128 / lr 0.004	8,192 / 4	3.485	79.953
H2 / U64 / lr 0.002	8,192 / 4	3.490	89.999
H4 / U128 / lr 0.004	2,048 / 4	3.981	48.431
H4 / U128 / lr 0.004	8,192 / 4	3.543	193.751

Payload 32 per head: H2 total width 64, H4 total width 128; eight blocks. More heads also increase capacity, and source/channel dynamics differ from the old single-head model. The baseline H2 pilot selects epoch 2 and overfits later; no head-count quality benefit is established. Adam interval U is separate from 16-character credit. Training reads admitted losing values and charges gradients, clipping and optimizer work. Contracts pass for causality, independent projections, evolving channels, all-head gradients and exact next-update recovery. All completed variants remain in the ledger; this table shows the latest four records.

Appendix B (continued). Ours: completed head/data scaling

Fixed d32 per head, eight blocks, pool2, credit16, U128/lr.004, four fitting passes, seed6. Every point scores the same 8,191 frozen development targets. More data improves these configurations, while four heads increase both width/capacity and fitting cost.

Ours = Sleeping Machines



parallel head data work

Ours: heads / fit	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓
H2 / 2,048	3.779	19.999	2.442
H2 / 8,192	3.485	79.953	2.440
H4 / 2,048	3.981	48.431	5.915
H4 / 8,192	3.543	193.751	5.914

The earlier 8K single-head controls reach receiver 3.311 and indexed KV 3.357 bpc. The multihead construction also changes source/channel dynamics and total width, so this is not a pure head-count ablation. The completed parallel-head 8K results missed the declared 0.10 bpc tolerance of the indexed control; the campaign stopped before 32K/131K promotion. Route-credit fidelity, recurrent/channel conditioning, candidate coverage and optimization are diagnosis targets. These results constrain this implementation rather than the whole substrate.

One seed and small fitting budgets. Whole fitting includes all four passes, backward, admitted losing-value credit and optimizer work. Development selection uses the lowest full development loss over those passes. Logical FLOPs and unit-weight special functions do not measure wall time, physical traffic or energy; no language supremacy follows from these points.

Appendix B (continued). Ours: longer temporal credit

Independent H2 heads, d32/head, 8 event blocks, pool2. 2,048 fitting characters / 4 passes; 8,191 frozen development targets, seed6. Adam uses U64 / lr0.002. Graphs remain live for 64 targets before detachment. Forward stored history and the inference architecture are retained.

Ours: credit	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
16	3.733	22.753	2.779	0.5059
64	3.740	23.019	2.811	0.5060

Sealed historical keys and values still affect predictions, but detachment removes later loss paths to their old producers. Longer credit restores those paths for more writes inside each optimizer window; it does not backpropagate through unlimited history. Temporal races, sparse receiver commits, separate Q/K/V and losing-route credit remain active.

Matched 16-credit records are shown when completed under identical settings. Additional backward/normalization/clip/Adam work is counted and peak memory is guarded. Inference traces average different representative spans; the operation definitions and inference architecture are the same. One seed, development selection, no frontier claim. Forward-partition equality and full 64-credit gradient/update contracts precede fitting.

Appendix B (continued). What most reduces research uncertainty

The main direction now tests native content-and-time computation directly. Episodic race attention remains a preserved comparison. The next comparison fits native state computation and an attention scaffold at equal total work.

Priority	Experiment	Doubt resolved
1	Integrated order/time learning; refitted credit/time controls; three seeds	Can deep sparse temporal state learn useful representations?
2	Native-core text8 adapter and small-to-larger data ladder	Does the native construction learn economically without a KV attention bank?
3	Occupy 4, 16, 64 stream states at fixed event/query budgets	Does useful state grow without proportional per-event activity?
4	Chronological real streams and predict-before-update adaptation	Does the advantage survive real data and online change?
5	Timestamp-aware AWS controls, equal-budget/quality curves	Is the quality/resource advantage reproducible?
6	Whole-system FPGA/ASIC timing, traffic and energy measurement	Does the physical substrate deliver the projected savings?

The first native branch uses eight event blocks, two independent receiver heads, observed source addresses, persistent content/state, analytic temporal evolution and counterfactual learning. It has no per-position KV attention. Other sources keep independent progress; source-local causal dependencies and internal joins remain charged.


The same native core receives token content for the language test. Sharing receiver maps across tokens changes capacity and parameter exposure; it is a whole-construction comparison, not an isolated attention-removal ablation. Persistent state is not a full-cache equivalence claim.

Contracts/smokes precede fixed-budget pilots, conditional capacity/data scaling and independent replication. Whole fitting, per-target work, inference, occupancy, memory and confidence intervals are published from completed files. Synthetic learning, real-data Pareto advantage and physical joules are separate milestones. The executable protocol, gates and current host limitations are documented in experiments/RESEARCH_VALUE_PLAN.md.

Appendix B (continued). Ours: content-gated temporal reception

Native eight-block independent-head models reuse key/query matches for two/four additional scalar clock policies. Local rotating/decaying clock vectors gate content-dependent projections and compose the next message. The waiting control retains the same clocks and joins but removes temporal reception/readout. The native parent has no extra clock branch.

Ours = Sleeping Machines



delay language quality work

Ours	Fit chars / passes	Dev bpc ↓	CPU fit GFLOPs ↓	Fit MFLOPs / target ↓	Infer MFLOPs / char ↓
Native	2,048/4	3.765	3.778	0.461	0.0974
R2 uniform	2,048/4	3.796	4.032	0.492	0.1056
R4 uniform	2,048/4	3.795	4.183	0.511	0.1109
R4 late	2,048/4	3.795	3.979	0.486	0.1041
R4 uniform wait	2,048/4	3.764	3.897	0.476	0.1007

Ours	Projected whole fit GFLOPs	Receivers / commits per token	Matches / clocks per token	Parameters
Native	3.777	32/16	32/16	54,907
R2/uniform/on	4.029	32/16	32/48	57,083
R4/uniform/on	4.178	32/16	32/80	59,259
R4/late/on	3.976	32/16	32/48	57,083
R4/uniform/waiting	3.891	32/16	32/80	59,259

Completed fits only; identical frozen 8,191-target development protocol, four passes, U64/lr.002/warm512, ordinary credit16. Different fitting sizes are explicitly marked. All scalar policies, projections, temporal bases, gates, counterfactual content teachers and actual Adam remain charged. Projected arithmetic removes only numeric clock simulation. Physical rate setting, clock circuits, traffic, precision and measured joules remain separate. Exploratory development results; these do not alone establish comparable-quality Transformer superiority.

Current matched-fit interpretation: best completed added-clock row is 0.0010 bpc better than native, with 3.13% more fitting work. Pending allocations and waiting controls cannot establish a benefit yet.

Appendix B (continued). Ours: count-carrying native receivers

The unchanged native eight-block core supplies the base predictive; addressed context-suffix receivers of orders 1..K carry sufficient statistics and deliver by an escape-race cascade with learned discount and concentration (Theory §§376–380, 387). The escape-gate variant makes discount and concentration per-position functions of the native predictive and count evidence; the count-message variant adds the counts to the base logits (Theory §389). Fitting counts are leave-one-out; development counts are prequential persistent state with frozen weights. Count tables are capacity; each target touches K addresses.

Model	Fit chars / passes	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Infer MFLOPs / char ↓
Native alone 2,048	2,048/4	3.765	3.778	0.461	0.0974

Model	Fit chars / passes	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Infer MFLOPs / char ↓
Count-carrying native K4 2,048	2,048/4	2.734	3.806	0.465	0.0982
Same, untrained base 2,048	2,048/0	2.741	Not trained	Not trained	0.0982
Count-carrying native K4 + count message 2,048	2,048/4	2.744	4.032	0.492	0.1102
Count-carrying native K4 + escape gate 2,048	2,048/4	2.695	3.824	0.467	0.0993
Count-carrying native K4 + escape gate [credit64] 2,048	2,048/4	2.694	3.843	0.469	0.0993
Count-carrying native K4 + escape gate [minimal core p2/d1] 2,048	2,048/4	2.694	0.072	0.009	0.0031
Count-carrying native K4 + escape gate [minimal core p2/d1] [credit64] 2,048	2,048/4	2.694	0.072	0.009	0.0031
Count-carrying native K4 + gate + message 2,048	2,048/4	2.703	4.053	0.495	0.1113
KN counts, frozen o4	2,048/1	3.615	Not FLOPs	Not FLOPs	Not FLOPs
Calibration ceiling: adaptive interpolated KN o6	2,048/1	2.521	Not FLOPs	Not FLOPs	Not FLOPs
Native alone 8,192	8,192/4	3.557	15.116	0.461	0.0974
Count-carrying native K4 8,192	8,192/4	2.671	15.227	0.465	0.0982
Same, untrained base 8,192	8,192/0	2.682	Not trained	Not trained	0.0982
Count-carrying native K4 + escape gate 8,192	8,192/4	2.595	16.853	0.514	0.1155
Count-carrying native K4 + escape gate 8,192	8,192/4	3.474	16.789	0.512	0.1155
Count-carrying native K4 + escape gate 8,192	8,192/4	2.601	15.314	0.467	0.0993
Count-carrying native K4 + escape gate [minimal core p16/d1] 8,192	8,192/4	2.592	4.181	0.128	0.0358
Count-carrying native K4 + escape gate [minimal core p16/d1] 8,192	8,192/4	3.492	4.161	0.127	0.0358
Count-carrying native K4 + escape gate [minimal core p2/d1] 8,192	8,192/4	2.588	0.288	0.009	0.0031
Count-carrying native K4 + escape gate [minimal core p2/d1] 8,192	8,192/4	2.588	5.995	0.183	0.0632
Count-carrying native K4 + escape gate [minimal core p2/d1] 8,192	8,192/4	2.593	1.744	0.053	0.0184
Count-carrying native K4 + escape gate [minimal core p2/d1] 8,192	8,192/4	3.471	1.744	0.053	0.0184
KN counts, frozen o4	8,192/1	3.081	Not FLOPs	Not FLOPs	Not FLOPs
Calibration ceiling: adaptive interpolated KN o6	8,192/1	2.414	Not FLOPs	Not FLOPs	Not FLOPs
Count-carrying native K4 32,768	32,768/4	2.560	60.790	0.464	0.0982
Same, untrained base 32,768	32,768/0	2.593	Not trained	Not trained	0.0982
Count-carrying native K4 + escape gate 32,768	32,768/4	2.447	60.790	0.464	0.0993
Count-carrying native K4 + escape gate [minimal core p2/d1] 32,768	32,768/4	2.401	1.151	0.009	0.0031
KN counts, frozen o5	32,768/1	2.704	Not FLOPs	Not FLOPs	Not FLOPs
Calibration ceiling: adaptive interpolated KN o8	32,768/1	2.271	Not FLOPs	Not FLOPs	Not FLOPs

Same 8,191 development targets for every row; one seed. Count increments/lookups are integer table work reported in the result files, not FLOPs. The initialized-base/escape row measures whole-model fitting benefit; it does not isolate the native base. Count rows are dev-selected-order references, not neural controls. How to read these rows (Theory §§393–394): at a few thousand to tens of thousands of characters, smoothed counting is a near-optimal estimator, and no learner (Transformers included) is expected to exceed it. At 10M characters the repository's dense Transformer control is still .12 bpc worse than frozen counts (§381). The strongest such reference, stream-adaptive interpolated Kneser–Ney, scores 2.521 / 2.414 / 2.271 bpc at 2K / 8K / 32K and 2.101 at 131K (results/count_reference/curie_adaptive_kn_language_reference_20261002T131500Z.json). It is a calibration ceiling, not a competitor. Distances to it measure remaining smoothing, and the rows here are mechanism diagnostics, not a verdict on the architecture, whose claims are tested on tasks where learning matters (§394). Exploratory development evidence.

At 2,048 fitting characters, fitting the native base and escape parameters improves 0.0071 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 2,048 fitting characters, fitting the native base and escape parameters improves 0.0542 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 2,048 fitting characters, fitting the native base and escape parameters improves 0.0455 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 2,048 fitting characters, fitting the native base and escape parameters improves 0.0381 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 2,048 fitting characters, fitting the native base and escape parameters improves -0.0032 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 32,768 fitting characters, fitting the native base and escape parameters improves 0.0329 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 32,768 fitting characters, fitting the native base and escape parameters improves 0.1930 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 32,768 fitting characters, fitting the native base and escape parameters improves 0.1454 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves 0.0108 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves 0.0988 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this

smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves 0.0988 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves 0.0922 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves 0.0890 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves 0.0915 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves -0.2996 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves -0.2977 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves -0.3213 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 8,192 fitting characters, fitting the native base and escape parameters improves 0.0809 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 2,048 fitting characters, fitting the native base and escape parameters improves 0.0464 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

At 2,048 fitting characters, fitting the native base and escape parameters improves 0.0543 bpc over their untrained initialization. The complete composed predictor improves over native-alone, while this smaller learning contribution is the relevant comparison for the cost of fitting the base. The integer count path remains charged separately.

Appendix B. Longer credit versus learned count smoothing

Completed four-arm integrated comparison, seed6: same2K fitting characters/four passes, 8,191 development targets, K4 escape gate, U64/lr.002/warmup512. Each fit processes 8,188 targets and 128 optimizer updates. Full H2/d16/depth8 and minimal H2/d2/depth1 each compare 16 versus 64 tokens of graph reach; persistent state and forward mechanisms remain unchanged within each core.

Model/credit	Dev bpc	Whole fit GFLOPs	Fit MFLOPs/target	Infer MFLOPs/char
Ours: full core /credit16	2.6951	3.824	0.467	0.0993
Minimal core control /credit16	2.6938	0.072	0.009	0.0031
Ours: full core /credit64	2.6943	3.843	0.469	0.0993
Minimal core control /credit64	2.6937	0.072	0.009	0.0031

Model/credit	Parameters	Core state slots	Updates/scores/teacher values per target
Ours: full core /credit16	54,955	32	16/32/32
Minimal core control /credit16	467	4	2/4/4
Ours: full core /credit64	54,955	32	16/32/32
Minimal core control /credit64	467	4	2/4/4

Full-core longer-credit gain 0.0008bpc; minimal-core gain 0.0001bpc. Full versus minimal advantage at 64: -0.0006bpc, with 1.005× full-core fitting work versus 16. The predeclared follow-up gate fails.

All neural fitting forward/loss/backward/normalization/clipping/Adam and admitted losing-value credit are charged in CPU emulator units; projected clockless work is a separate result ledger. Each target also looks up 4 count addresses; integer counts, discovery, traffic and energy stay separate. This reuses development data and is a single-seed screen, not confirmation, semantic-feature proof or supremacy. Passing requires at least 0.02bpc full-core credit gain AND at least 0.02bpc advantage over the matched minimal 64 core at no more than 2× full-core fitting work. No automatic larger fit.

Appendix B. Does the trained deep state contribute?

Frozen saved full/minimal 64-credit models from the matched 2K fits above. Same 32 development targets at positions 128–159, same race noise and causal count vectors. Erase stored content once at the slice start, retaining arrival times and absolute count cursor; state may rebuild. The learned escape gate responds to the changed neural base.

Model/intervention	Composed bpc	Raw base bpc	Change in composed bpc
Full core/intact	3.4541	4.9093	+0.0000
Full core/Erase source context	3.4916	5.0308	+0.0375
Full core/Erase receiver content	3.4495	4.9933	-0.0046
Full core/Erase all content	3.5022	5.1477	+0.0481
Minimal control/intact	3.4904	4.9638	+0.0000
Minimal control/Erase source context	3.4885	4.9424	-0.0018
Minimal control/Erase receiver content	3.4901	4.9649	-0.0002
Minimal control/Erase all content	3.4882	4.9401	-0.0022

Full-core stored content helps this slice by 0.0481bpc, mainly through the carried source context; erasing receiver content alone does not hurt. Useful recurrence is present, but this does not establish hierarchical semantics, generalization across slices, or a full-development lead. Per-layer receiver erasures are preserved in the diagnostic JSON.

Model	Mean base responsibility	Median base responsibility	Dynamic/base-only gradient norm
Full core	0.0717	0.0045	0.207
Minimal control	0.0571	0.0031	0.195

Credit diagnostic uses 64 fitting targets at positions 64–127. With gate values fixed, the exact base-logit derivative is responsibility × (q − onehot); numerical errors are below 5e−9. Actual dynamic-gate gradients reach every layer. Responsibility suppression is correct mixture credit, not a demonstrated autograd bug; Adam can compensate for uniform scaling. The distribution and

usefulness of the residual signal, information retention and retrieval remain hypotheses.


No optimizer steps or weight changes. Raw base is trained as a conditional residual, so its standalone bpc is diagnostic. Audit12.85s/522,736KiB; no FLOP or superiority claim. Source: local_deep_core_attribution_20261002T075000Z.json. The synthetic long-range queue was retired before launch after unsupported chance claims and credit/update confounding were found; repaired-driver contracts require matched optimizer windows and separately measured local controls.

Appendix B. Restoring addressed value credit: completed fits

Five matched integrated arms:1,024 fitting characters/four passes,4,092 targets/64 Adam updates,2,047 dev targets, H2/pool2/c16/U64/warm512/lr.002/seed6. Full native temporal core versus original projected-value slots, late-projected raw-feature slots and two shallow controls. Core initialization and race RNG are matched. Minimum cold dev bpc over fixed passes selects each checkpoint; frozen fit replay is secondary and never selects weights.

Model	Cold dev bpc	Fit-replay dev bpc	Selected pass
Native full p16/L8	3.9681	3.9674	4
Stored projection p16/L8	3.9030	3.8967	4
Late projection p16/L8	3.9421	3.9809	4
Late p16/L1	4.0862	4.0801	4
Late p2/L1	4.5396	4.5274	4

Ours = Sleeping Machines



value credit learning 20261002T090800Z

Late-full gain over native:+0.0260bpc; over stored projection:-0.0391; over minimal:+0.5975; over same-width shallow:+0.1441. The predeclared nomination gate fails.

Late projection preserves fixed-weight computation by linearity and restores fixed-feature projection credit from old detached slots. It does not restore historical core-producer credit or add linear-reader expressivity. Fixed hash addresses are not learned pooling/KV race attention. The same-width shallow arm isolates depth better than the width-changing minimal control. Reused dev and one seed: exploratory quality/learning evidence, not semantic proof or supremacy.

Appendix B. Fixed-feature credit and compiled inference

Frozen-reader compilation folds $A=R_vW$ once and removes the unused W matrix. Coupled predictions and raw-slot state preserve the unfused construction within numerical tolerance. Under2FLOPs/MAC, fold2d³ replaces2d² per occupied read, breaking even after d occupied reads. This changes deployment work, not the fitting optimizer; compiled models refuse training.

Late model	Unfused infer KFLOPs/target	Fused infer KFLOPs/target	Removed weights	Fold KFLOPs	Break-even occupied reads
late_full	100.180	99.540	1024	65.536	32
late_shallow	20.562	19.922	1024	65.536	32
late_minimal	1.224	1.214	16	0.128	4

Inference samples include 16 actual targets after 128 warm tokens; the saved projection work equals $2d^2$ times observed occupied reads. Matrix-fold work is charged above; constructor initialization/copy/RNG are separate. This is arithmetic/weight reduction with frozen quality, not measured latency, energy or competitive superiority.

Text-trained model	Prefix-pair KL, gap8	Prefix-pair KL, gap32	Prefix-pair KL, gap64
native_full	2.14e-04	1.84e-03	3.46e-04
addressed_full	1.94e-03	5.28e-04	1.10e-03
late_full	1.87e-03	4.80e-04	5.50e-04
late_shallow	1.88e-06	3.18e-07	2.90e-07
late_minimal	8.10e-14	0.00e+00	0.00e+00

Each parity diagnostic balances all four input bit pairs with identical noise/query suffix. Actual causal query count vectors are identical for orders 1–8, but targets are opposite across paired prefixes. A predictor restricted to those suffix/count inputs has at least 1bit target logloss; the full prefix determines parity exactly. Cue-free prefixes make the count contract explicit, replacing the retired synthetic all-orders independence assertion.

These frozen text checkpoints were never trained on parity. KL and top-context differences diagnose prefix dependence/retention, not parity learning or semantic abstraction. Zero optimizer steps, parameter fingerprints unchanged. A later task fit needs matched full/shallow controls, new-prefix generalization and complete prefix computation/credit accounting. Historical feature-producer credit, learned address pooling and KV races remain open. Source:

local_value_credit_frozen_20261002T090800Z.json; theory63 records exact scope.

Appendix B. Frozen feature geometry and matched count calibration

Model	Frozen fit bpc	Cold dev bpc	Fit head gradient norm	Dev feature participation rank
native_full	3.6920	3.9681	0.597	2.31
addressed_full	3.6595	3.9030	0.536	1.95
late_full	3.7272	3.9421	0.637	1.60
late_shallow	3.8785	4.0862	0.249	4.53
late_minimal	4.5775	4.5396	0.323	1.21

At the selected weights, replay 1,023 fitting and 2,047 development targets from cold state with the primary evaluation RNG. Capture actual head inputs, then compute fixed-feature linear-head loss derivatives and covariance in float64. No head fitting, optimizer step or parameter changes. Frozen dev scores reproduce the primary results within $1e-5$ bpc.

The participation rank is $(\text{trace } C)^2 / \text{trace}(C^2)$, a measure of feature geometry. Finite joint training need not make the final head stationary; its gradient norm and covariance rank alone do not identify a bug or semantic abstraction. Frozen fitting loss helps separate training fit from development generalization without using online training scores.

Count order	Frozen KN bpc	Frozen WB bpc	Frozen AD bpc	Adaptive WB bpc	Adaptive AD bpc
1	3.8263	3.8008	3.8692	3.4806	3.5024

Count order	Frozen KN bpc	Frozen WB bpc	Frozen AD bpc	Adaptive WB bpc	Adaptive AD bpc
2	3.7986	3.9596	3.9576	3.1259	3.1566
3	3.7530	4.1254	4.0371	3.0285	3.0271
4	3.7515	4.2123	4.0771	3.0873	3.0308

Same 1,024 fit characters and 2,047 dev targets, reused count-reference implementation, fixed discount.75, all orders/methods shown separately. Count construction uses one pass; learned models use four gradient passes. Counts are fit-prefilled; adaptive variants also update from observed dev outcomes after prediction. These state/history policies differ from cold neural scoring; frozen neural fit replay is reported separately above.


KN: interpolated Kneser–Ney; WB: Witten–Bell; AD: absolute discount. Integer represented count increments/lookups and measured construction/scoring wall time are in the completed audit, without conversion to neural FLOPs. Equal scalar loss calibrates prediction quality; it does not establish identical features or a same-quality compute advantage. No per-target order selection or official-test access. Full spectra, data/checkpoint hashes and diagnostic boundaries are retained in `local_value_credit_frozen_20261002T090800Z.json`.

Appendix B. Beyond query counts: balanced joint learning

Each episode observes a bit followed by its complement, then another such pair, eight shared noise symbols and a unique query cue. The target is the relation between the earlier bits. All four bit combinations occur for each noise suffix. Observed symbol totals and actual order-0..8 query count vectors match; all 36 declared WB/AD controls score one bit.

Integrated model	Dev bits/query	Dev accuracy	Selected pass
native_full	0.9995	50.00%	8
tapped_full	1.0002	48.44%	14
tapped_shallow	1.0002	50.00%	16

Ours = Sleeping Machines



local balanced joint 20261002T121500Z analysis learning

The native and learned-delay variants fail the predeclared 75%/.8-bit dependency gate. Eight fitting suffix groups give 32 targets/pass; 16 passes, 512 target presentations, 64 Adam updates, 64 held-out dev targets. All arms have query-only binary supervision and full 15-event credit; the target is never an input. Noise couples the quartet and variants.

The one-bit bound restricts the predictor to query suffix/count inputs, including root marginals. It excludes arbitrary inspection of the full prefix or other count addresses. Counts should lead in their supported local regime; this test asks for useful nonlocal prediction. One selected seed/development set, not semantic language or supremacy.

Appendix B. Frozen retention versus interaction decoding

Encoder / degree	Fit bits/query	Fresh bits/query	Fresh accuracy
native full selected / 1	0.994	1.063	50.78%
native full selected / 2	0.920	12.138	50.78%
native full initial / 1	0.999	1.026	50.00%
native full initial / 2	0.876	25.333	50.78%
tapped full selected / 1	0.999	1.011	49.22%
tapped full selected / 2	0.943	21.694	50.00%
tapped shallow selected / 1	1.000	1.011	50.00%
tapped shallow selected / 2	0.980	7.876	50.00%

Frozen encoder	First-bit probe accuracy	Second-bit probe accuracy
native_full_selected	50.78%	51.56%
native_full_initial	50.78%	48.44%
tapped_full_selected	50.78%	53.12%
tapped_shallow_selected	50.78%	51.56%

Freeze the selected encoders, plus native initial weights as a reservoir control. Fit zero-initialized affine or standard degree-2 polynomial residuals on actual query features. Fitting standardization only; four race-noise views of the same 32 fitting episodes. Regularized full-batch L-BFGS, fixed $\lambda=1e-5/100$ iterations; reported derivatives do not establish convergence. New seed 73001 has 32 paired suffix groups/128 targets and never selects readout settings or weights.


Neither readout generalizes; polynomial heads become badly overconfident. Separate one-bit probes also remain near chance. They use extra bit supervision only for diagnosis; their outputs never enter the relation predictor. These failures motivate evidence-access tests rather than assuming a larger decoder solves the current model.

Zero core optimizer steps and encoder fingerprints preserved. Eight predeclared readout arms on one fresh synthetic distribution, not general confirmation or a semantic claim. Failed probes do not prove all information was erased or a universal learning ceiling.

Appendix B. Protected outcomes and joint terminal races

Each observed token writes its observed successor at its predecessor address, never a target-derived or hand-selected bit address. All occupied addresses are candidates. Two learned query/key exponential races deliver small interpreted symbol values to a generic bilinear query decoder. Native temporal computation, sparse vector receivers and earlier counterfactual learning remain. No XOR extraction enters prediction.

Integrated arm	Selected dev bits	Reserved bits	Reserved accuracy	Selected pass
joint_full	0.077	0.102	96.09%	16
local_full	0.933	0.750	78.12%	13
joint_shallow	0.062	0.119	95.31%	16



local joint outcome 20261002T125700Z analysis learning

Reserved seed74001/32 groups/128 queries scores fixed selected weights after all fits; no retuning. Joint-minus-local loss improvement: +0.648 bits; full-minus-shallow improvement: +0.017 bits. The joint-credit gate passes; all arms remain visible.

One fitted seed and structured generator. Sixteen fit groups/64 queries, 16 passes, U4/lr.01/p4/H2/pool2,fullL2/shallowL1. The one-bit bound applies to identical query suffix/count inputs including root marginals, not arbitrary inspection of other count addresses. No natural-language, dense-control or resource-supremacy claim.

Appendix B. Joint-credit work, capacity and activity

Arm	Whole fit GFLOPs est.	Fit MFLOPs/query est.	Inference MFLOPs/query
joint_full	0.2379	0.2323	0.0534
local_full	0.2319	0.2265	0.0534
joint_shallow	0.1534	0.1498	0.0361

Arm	Core / raw address capacity	Mean raw occupied	Core commits / raw writes per query	Terminal keys / values per query	Training loss pairs per query
joint_full	8 / 27	10.56	60 / 14	21.12 / 2	112.06
local_full	8 / 27	10.56	60 / 14	21.12 / 2	0.00
joint_shallow	4 / 27	10.56	30 / 14	21.12 / 2	112.06

All fitting columns use 1,024 query presentations, 256 Adam updates and all 15 observed prefix events per query. Whole-fit first/last-window estimates include native candidates, losing proposals, protected-state discovery, terminal decoder/loss, backward, clipping and Adam. Inference scores every occupied key and delivers only two values. Available addresses, occupied state, scored keys, commits, raw writes and value deliveries are distinct; sparse activity does not imply zero key or learning cost.

Joint training enumerates C^2 terminal losses and differentiates their categorical expected risk, giving exact conditional terminal content-choice credit. Local training uses a sampled pair and the existing value-linearized race surrogate. Their initial forward predictions and inference policy match; learning estimators differ. Earlier native route surrogates, raw fixed-address writes and downstream timing credit remain separate limitations. No exact whole-sequence gradient is claimed.

Joint/local whole-fitting work ratio: 1.026. 2FLOPs/MAC plus unit specials; traffic, raw integer state, Python objects, RNG and energy separate. Five contracts and three accounting smokes precede fits. The same-width shallow comparison changes initialization too. A successful protected-state read does not prove learned context pooling, arbitrary-distance KV retrieval or useful deep producer credit.

Appendix B. Fixed joint-credit confirmation

Fit seed / credit	Selected dev bits	New suffix bits	New suffix accuracy	Pass	Zero-value accuracy
s7 / joint	0.055	0.050	99.22%	16	98.83%
s7 / local	0.009	0.003	100.00%	15	100.00%
s8 / joint	0.000	0.000	100.00%	16	75.00%
s8 / local	0.001	0.004	100.00%	16	99.22%

Fit seed	Joint improvement bits/query	Paired suffix interval	Joint/local fitting work
7	-0.046	[-0.093, -0.016]	1.025
8	+0.003	[+0.000, +0.008]	1.026

Two additional fitted seeds 7/8 use the unchanged full p4/L2 configuration and 16 passes. All four declared fits are shown. Development seed 72001 selects the minimum across fixed passes; seed 75001/64 suffix groups/256 targets is reserved for fixed selected models. No parameter, learning rate, stopping rule or evaluation setting is returned on this set.

0 of 2 declared joint-versus-local gates pass. A seed must reach at least 75%/.8bits and improve local credit by .05bits. Paired suffix bootstrap intervals are conditional on the fitted seed and synthetic generator; they are not a broad confidence interval over learning algorithms or tasks.

Initial predictions, race choices and probabilities match within each joint/local pair. Inference mechanisms are identical. The differing training risk/credit objectives therefore test this terminal learning intervention, while protecting all earlier negative evidence.

Zero-value evaluation removes delivered content while keeping native context, query/key races and contextual decoder. No fitting or model selection; every parameter is restored and verified afterwards. Three arms retain 98.8–100% accuracy: native context has learned useful nonlocal prediction. Joint seed 8 drops to 75%. Value delivery is therefore not necessary in every selected fit; neither its training benefit nor a depth premium is isolated.

Query-suffix/count bound only. The protected bank has fixed observed predecessor addresses and no learned writes; the standard bilinear decoder and exact terminal content credit do not establish deep core learning, clock-gradient accuracy or natural-text gains.

Appendix B. A terminal symmetry, not a learning impossibility

Numerical contract	Completed result
Actual outcome bank	Uniform two-value law identical across second-bit alternatives
Fixed zero-context risk	1.000000bits
Maximum decoder gradient	6.939e-18
Maximum policy gradient	0.000e+00
Analytic marker-policy witness	6.921e-09expected bits

At query time the four distinguished outcome addresses contain $b_1, b_2, 1 - b_2, n_0$. Their multiset is $b_1, 0, 1, n_0$. Remaining entries depend only on the shared noise suffix. Uniform independent value reads therefore have the same full joint distribution for both second-bit alternatives. If additional context also lacks that bit, no decoder of those values can predict the balanced relation better than one bit.

At zero residual, every candidate pair has the parent logit. Exact terminal key credit is initially zero. With fixed zero-logit uninformative context, uniform policies and balanced examples, decoder gradients are zero too. The actual bank and decoder contracts verify this conditional stationary point

in float64 on four suffix groups. This is not a stationary-point proof for the whole trainable native core.

An explicit target-independent existence witness selects marker addresses24/25 and interprets symbols0/1 as signed values. A bilinear interaction can predict their relation with very small expected loss. The witness is analytic and never enters the fitted model or benchmark score. Finite score clamps leave a nonzero routing error floor, so unlimited decoder confidence need not reduce unconditional expected loss.

The useful solution and terminal symmetry coexist. Random key contrast may break the alias, which is compatible with late fitted transitions but does not prove their cause. A new controlled conditioning study must retain all seeds and mechanisms; these fixed confirmation fits are neither extended nor tuned after observation.

Theory70 states the information conditions and proof. Numerical result records source hashes, targets, wall and RSS. The obstruction concerns addressed value pooling and terminal initialization; it does not imply that all counting, temporal computation or architectural adaptation faces a mathematical ceiling.

Appendix B. Practical headroom: a stronger joint-event table

Fixed table evaluation	Queries	Accuracy	NLL	Maximum accuracy headroom
development	256	99.609%	0.051632	0.391pp
confirmation	1024	99.512%	0.055258	0.488pp

This control consumes observed timestamps as well as text. It retains the complete question string and the latest observed time for each of four marks. A generic fit learns one mark-age split per question: which mark, threshold, split direction and both leaf probabilities. The generator word/mark mapping and recency threshold are not supplied.

Same512 distinct fitting episodes/seed1301,one fitting pass,20 question strings; fit wall0.027026s,maxRSS67,940KiB. Fixed learner; dev2301/256 and confirmation3301/1,024 episodes do not tune it. Contracts verify label mutation, appended future observations and common clock shifts cannot change causal prediction features.

Less than half a percentage point of accuracy headroom remains. The proposed20-point learned-control advantage gate is impossible against this reference on these episodes, even for a perfect learner. A gain over the time-blind table would therefore measure extra supplied information rather than demonstrate the requested practical advantage.

The joint recency task remains a capability and mechanism diagnostic. Likelihood may still improve; neither saturated accuracy nor tiny fitting cost supports general supremacy. Strong real-stream calibration with causal common inputs takes priority. Old successful mechanism evidence and negative results are preserved.

Logical content inspections, timestamp writes, age subtraction, sorting comparisons, threshold candidates, count/probability work and wall are saved. They are not converted to neural FLOPs. Full question lookup is a task-specific conventional control, not a general language model or the integrated research architecture. Theory72 records scope.

Appendix B. Real gesture practical calibration

A meaningful practical region:984 first-second fitting gestures from users1–19 and192 development gestures from users20–23. The common causal representation has4×4 spatial cells, two polarities and20 observed50ms count closures. No official test was opened.

Control	Dev accuracy %	Dev NLL	Fit seconds	Dev inference seconds
Calibrated counts	58.85	1.3734	0.042	0.000

Control	Dev accuracy %	Dev NLL	Fit seconds	Dev inference seconds
linear_C0.1	61.46	1.2036	0.742	0.001
linear_C1	59.38	1.5903	0.439	0.001
linear_C10	57.81	2.4610	0.191	0.001
rbf_C1_g0.25	67.19	0.8804	0.860	0.068
rbf_C1_g1	74.48	0.7358	1.142	0.078
rbf_C1_g4	73.44	0.8643	1.816	0.077
rbf_C10_g0.25	68.23	0.8122	0.799	0.073
rbf_C10_g1	73.44	0.7065	1.259	0.072
rbf_C10_g4	72.40	0.8338	1.808	0.086

The lowest development-NLL control is rbf_C10_g1; its accuracy exceeds calibrated counts by 14.58 points. Maximum grid accuracy is 74.48%, so this observation protocol has headroom rather than the 99.5% saturation of the cheap recency table. This admits an integrated learning test, not a native advantage claim. The older weak dense gesture screens are not the practical ceiling.

All nine fixed learned cells are preserved. Linear/RBF controls use fitting-only log-count centering/scaling. Counts use all 640 time bins and 3-fold fitting-only alpha/temperature calibration; this avoids treating correlated camera events as independent label evidence. SVM probability calibration is included in measured fitting.

Common raw preprocessing costs 123.976s. The complete control campaign takes 15.926s with 411.1MiB peak process RSS. Table inference uses the complete 192-example batch. Third-party solver FLOPs are unmeasured, not zero; CPU wall, arithmetic, storage and energy are distinct. Development selection and one fitted seed are exploratory evidence, not independent confirmation.

Appendix B. Integrated real-packet quality and work

Model	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
Calibrated counts	58.85	1.3734	Unmeasured	Unmeasured	Unmeasured
linear_C0.1	61.46	1.2036	Unmeasured	Unmeasured	Unmeasured
linear_C1	59.38	1.5903	Unmeasured	Unmeasured	Unmeasured
linear_C10	57.81	2.4610	Unmeasured	Unmeasured	Unmeasured
rbf_C1_g0.25	67.19	0.8804	Unmeasured	Unmeasured	Unmeasured
rbf_C1_g1	74.48	0.7358	Unmeasured	Unmeasured	Unmeasured
rbf_C1_g4	73.44	0.8643	Unmeasured	Unmeasured	Unmeasured
rbf_C10_g0.25	68.23	0.8122	Unmeasured	Unmeasured	Unmeasured
rbf_C10_g1	73.44	0.7065	Unmeasured	Unmeasured	Unmeasured
rbf_C10_g4	72.40	0.8338	Unmeasured	Unmeasured	Unmeasured
Ours p16/L2/H2	65.10	0.9632	20.075	2.550	0.592


Same 984 distinct fitting gestures and 192 subject-disjoint development targets. Ours fits eight fixed passes: 7,872 target presentations and 496 Adam updates, U16, lr.003, seed6. Conventional solvers have their own convergence/calibration policies; equal pass or fitting-work protocols are not claimed. Every learned baseline is retained.

Selected native pass8: accuracy difference versus minimum-NLL control -8.33 points, NLL improvement -0.2567. Development quality dominance flag: False. Independent practical advantage remains unproved; no pending, best-seed or official-test prediction fills this table.

Native estimates include complete prefix/query forward/backward, all scored keys, candidate values, counterfactual value credit, normalization/clipping and Adam. First/last windows are sampled by actual16/8 target size;2FLOPs/MAC and unit-weight special functions. Solver arithmetic is unmeasured. Preprocessing, validation, traffic and measured CPU latency must be charged separately; isolated inference estimates cannot establish total resource or energy advantage.

Appendix B. Real-packet learning, capacity and total workflow

Ours = Sleeping Machines



local dvs native analysis 20261002T151000Z learning

Completed native quantity	Value
Parameters / available receivers	15,523 / 8
Persistent state tensor bytes	720
Keys / selected updates / counterfactual values per fitting prefix	168.0 / 84.0 / 168.0
Fit plus nonfitting forward GFLOPs estimate	21.256
Native workflow / whole control grid seconds	1557.590 / 15.926
Common raw preprocessing seconds	123.976
Native / grid workflow plus preprocessing seconds	1681.566 / 139.902
Native peak process RSS MiB	347.8

The unchanged native core computes through temporal races, separate keys/values, sparse persistent receiver updates and local mixing. Nonempty observed packets arrive at their physical closure and an observed query at1s; state resets per gesture. Race noise depends on pass/fitting seed, never clip identity, index or label. Counts are observed camera content rather than fitted statistical prediction experts.

Workflow wall includes fitting, validation, checkpoints and operation profiling; the complete control-grid wall includes all fixed solvers, calibration, scoring and serialization. Nonfitting neural forward work is estimated using completed prefix samples; raw preprocessing is common and added once to each pipeline. Small tensor state does not mean small process RSS, zero scoring cost, useful extra depth or zero optimizer work. Earlier numerical admission/research costs remain separately saved.

Appendix B. Frozen practical inference and storage

Model	Dev accuracy %	Dev NLL	Sequential ms / prefix	Saved model KiB
Calibrated counts	58.85	1.3734	0.013	55.4
linear_C0.1	61.46	1.2036	0.163	66.1
linear_C1	59.38	1.5903	0.161	66.1
linear_C10	57.81	2.4610	0.167	66.1
rbf_C1_g0.25	67.19	0.8804	0.717	4632.4
rbf_C1_g1	74.48	0.7358	0.670	4730.5
rbf_C1_g4	73.44	0.8643	0.711	5081.4
rbf_C10_g0.25	68.23	0.8122	0.608	4451.8
rbf_C10_g1	73.44	0.7065	0.671	4766.6
rbf_C10_g4	72.40	0.8338	0.680	5081.4
native	65.10	0.9632	55.076	103.9

All models consume the same saved observed first-second packet counts, one prefix at a time on one CPU thread. Fit-only feature transformation and native race simulation are timed. Three deterministic repeats measure execution variation; they are not independent fits. Warmup, loading and common raw event coalescing are reported separately.

Every frozen probability matches the completed development result and every repeated prediction is identical. Native weights are the fixed minimum-NLL selected checkpoint; no model is refitted, temperature-adjusted or selected during this audit. All conventional cells remain visible, including faster or better alternatives.

Storage is uncompressed joblib serialization of each fitted model plus its necessary fitting-only transform. It is not resident process memory or memory traffic. Native parameter tensor bytes and persistent state bytes are separate quantities. CPU emulator latency cannot be relabelled as event-hardware latency or measured energy.

Subject-disjoint development evidence only. A quality/resource tradeoff here requires frozen independent confirmation before promotion. No test leakage, broad supremacy, useful-depth premium or dormant-capacity advantage is inferred from this audit.

Appendix B. Real learned features versus the initial reservoir

Frozen encoder	Original accuracy %	Fresh-head accuracy %	Fresh-head NLL	Head	Replay seconds	Head grid seconds
initial	12.50	57.81	1.1311	rbf	66.606	2.949
selected	65.10	67.19	0.9999	linear	66.168	1.915

Both encoders expose the same 32-dimensional query feature and use the same 984 fitting/192 development gestures and race draws. Every original prediction is reconstructed exactly by its saved head; replay and new head fitting preserve every encoder parameter. The initial encoder is the exact pre-fitting reservoir, not a separate tuned control.

Three linear C values and six RBF C/gamma cells are evaluated by 3-fold fitting-only decoder NLL. The selected head is then fitted once on all fitting features. No head hyperparameter is selected on development. Scaling is fitting-fold only. Conditional decoder CV is not unbiased end-to-end validation because the selected encoder already saw all fitting labels during its original fit.

The fitted representation gives a better selected readout than the initial reservoir: 67.19% versus 57.81%. This supports useful feature learning under this probe protocol. Replacing the fitted

head increases accuracy only 2.08 points and worsens NLL relative to the original 65.10%/.963161. The final decoder alone does not close the strong raw-input control gap; failed finite heads also do not prove information is absent.

This is a frozen diagnostic, not an architectural substitution or practical advantage. The selected encoder still pays its original 20.075 GFLOPs/1557.590s workflow, plus full replay and every readout fit shown here. Solver arithmetic remains unmeasured. Useful extra depth, semantic language features, whole-route gradient accuracy and independent confirmation are separate questions. No official test was opened.

Theory 75 states admission and scope. Selected initial/fitted readout artifacts, all fitting-CV cells, native/checkpoint hashes, probability arrays, wall and RSS are saved.

Appendix B. Shared fitting noise: covariance hypothesis fails its gate

Saved seed	All variance ratio	Route ratio	Time ratio	Message ratio	Gate
6	0.99787	1.00166	1.04860	0.99099	FAIL
7	1.00178	0.99119	0.97407	0.99542	FAIL

Four fixed fitting prefixes per saved native seed 6/7, 32 independent whole-history draws and zero optimizer updates. Each clip gradient is recorded under the same stream. Shared-batch trace covariance is measured; independent-coupling covariance is estimated from those same marginal samples. Their difference equals the cross-clip covariance sum.

Both route-map ratios miss the declared 1.20 admission gate. Proposed fresh_shared/fresh_independent fits are stopped. Neither model shows a large cross-clip covariance penalty on these prefixes. This is not a corpus-wide result or a test of repeated-noise adaptation across optimizer updates. Small-group variance ratios retain their finite-sample scope.

Theory 93 derives the effective-batch covariance law and separates per-window freshness from cross-row decorrelation. Independent noise need not help when cross-covariance is negative. Flipout is a primary-paper analogy about shared perturbation correlation, not an implementation or transferred quality guarantee for these event races.

Two algebra contracts pass. Audit 56.443s/396.2MiB; 0.637545 known diagnostic GFLOPs est. Complete first-draw per-clip forward/backward traces times draws; reporting reductions separate. Frozen weights preserved. No training, official test or superiority claim.

Appendix B. Causal persistent-state access: positive probe, qualified cause

Frozen encoder	Feature access	Dimensions	Dev accuracy %	Dev NLL
initial	Query	32	57.81	1.1311
initial	Query+state	176	63.54	0.8862
selected	Query	32	67.19	0.9999
selected	Query+state	176	71.35	0.8591

Encoder	Original fit GF est.	Fit MF / presentation	Core replay GF est.	Replay MF / prefix	Decoder/grid GF
initial	0.000000	0.000000	1.735341	1.475630	Unknown
selected	20.075193	2.550202	1.735341	1.475630	Unknown

Same 984 fitting/192 development gestures. Augmentation adds all eight pre-query receiver memories, their ages at source query admission and occupancy flags to the actual 32-dimensional query feature. No future or target-dependent feature. Three contracts reproduce serial state, original

probabilities and saved query-probe probabilities while keeping every encoder parameter bitwise fixed. Nine-cell three-fold fitting-only head selection.


Initial augmentation improves NLL by .244819; trained augmentation by .140774. Both pass the declared diagnostic access signal. The trained augmented probe reaches 71.35%/.859084, versus original native 65.10%/.963161, compact prototype 66.67%/.902951 and strong full RBF 73.44%/.706478. Decoder/grid costs are unknown, so these are quality comparisons and diagnostic signals, not practical advantage.

Important qualification from the completed partition controls: stronger query-only regularization alone reaches .889842; adding full state at that same C contributes only .030759 further NLL improvement. At the weaker C, full-state NLL worsens to 1.278955. Most apparent trained-state gain is therefore compatible with regularization rather than uniquely missing information. Early/late state partitions retain smaller gains.

Theory94/95. Dense all-state probing is a diagnostic, not sparse inference or a new main architecture. Selected encoder saw all fitting labels and was dev-selected; decoder CV is not unbiased end-to-end validation. Known original fit uses 7,872 presentations; current two-pass core replay uses 1,176 prefixes. Solver/materialization/ traffic/energy and total probe fit/inference remain unmeasured rather than zero.

Appendix B. Retained payloads, clocks, layers and regularization

Ours = Sleeping Machines



local dvs state partition probe 20261002T214000Z learning

Encoder	Added state/control	Dims	Accuracy %	NLL	Kind/C
initial	Payloads	160	63.54	0.8862	rbf/10.0
initial	Clocks	48	55.21	1.1245	rbf/10.0
initial	Layer0	104	66.15	0.9048	rbf/10.0
initial	Layer1	104	64.58	0.9545	rbf/10.0
initial	Query@state setting	32	57.81	1.1311	rbf/10.0
initial	State@query setting	176	63.54	0.8862	rbf/10.0
selected	Payloads	160	71.35	0.8603	linear/0.1
selected	Clocks	48	67.71	1.0178	linear/1.0
selected	Layer0	104	70.83	0.8383	linear/0.1
selected	Layer1	104	69.79	0.8403	linear/0.1
selected	Query@state setting	32	67.19	0.8898	linear/0.1
selected	State@query setting	176	68.23	1.2790	linear/1.0

Four partitions retain the query: both payload layers, ages/occupancy only, layer0 state or layer1 state. Same nine-cell fitting-only decoder selection per partition, both initial and trained encoders. Two fixed configuration swaps per encoder test regularization without another grid or development selection.

Trained payload .860321, layer0 .838315 and layer1 .840329; clocks-only 1.017809. The query C.1 control .889842 explains most full-state gain. These finite probe differences do not isolate a causal depth failure or population conditional information. All outcomes remain visible; core producers/routes were never retrained.

Theory95. Cached features add zero CORE replay, not zero fitting work. Prior encoder fit/core replay retained; additional solver arithmetic unknown. Conditional folds share the label-trained encoder. Producer-held fitting examples are the next decoder-selection check. No official test or supremacy claim.

Appendix B. Completed packet-scale clock comparison

Model	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
Calibrated counts	58.85	1.3734	Unmeasured	Unmeasured	Unmeasured
linear_C0.1	61.46	1.2036	Unmeasured	Unmeasured	Unmeasured
linear_C1	59.38	1.5903	Unmeasured	Unmeasured	Unmeasured
linear_C10	57.81	2.4610	Unmeasured	Unmeasured	Unmeasured
rbf_C1_g0.25	67.19	0.8804	Unmeasured	Unmeasured	Unmeasured
rbf_C1_g1	74.48	0.7358	Unmeasured	Unmeasured	Unmeasured
rbf_C1_g4	73.44	0.8643	Unmeasured	Unmeasured	Unmeasured
rbf_C10_g0.25	68.23	0.8122	Unmeasured	Unmeasured	Unmeasured
rbf_C10_g1	73.44	0.7065	Unmeasured	Unmeasured	Unmeasured
rbf_C10_g4	72.40	0.8338	Unmeasured	Unmeasured	Unmeasured
Ours original	65.10	0.9632	20.075	2.550	0.592
Ours packet-scale clock	66.15	1.0420	20.075	2.550	0.592
Compact prototype33	66.67	0.9030	Unmeasured	Unmeasured	Unmeasured


Both native fits use 984 gestures/eight fixed passes, 7,872 presentations and 496 updates; 192 subject-disjoint dev targets select pass 8 by minimum NLL. Only initial decay rates and rotation frequencies are scaled to the observed 50ms packets. Temporal races, key/value separation, sparse updates and counterfactual learning remain. Parameters 15,523; available receivers 8; per-prefix 168 key scores, 84 selected updates, 168 candidate values.

Clock fit wall 1565.128s; peak RSS 347.2MiB. Accuracy rises 1.04 points versus original while NLL worsens. Both trail the selected full kernel and compact prototype in both quality measures. This initialization change does not establish advantage; the original lower-NLL result and all control cells are retained.

Same units and native target denominators in every work column. Solver FLOPs remain unmeasured; unequal solver policies and development tuning preclude iso-FLOP superiority. Raw preprocessing, validation, checkpoints, traffic and energy remain separate. One seed; no official-test access. Source: local_dvs_clock_full_20261002T153000Z.json; theory 76.

Appendix B. Completed matched credit pilot: terminal pairs / seed6

Ours = Sleeping Machines



local dvs credit comparison pairs 20261002T174000Z learning

Credit	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
local	54.17	1.3059	2.285696	2.2321	0.5917
terminal pairs	56.77	1.3498	2.296461	2.2426	0.5918

Credit	Workflow seconds	Peak RSS MiB	Keys / fit target	Commits / fit target	Values / fit target
local	94.464	338.9	168	84	168
terminal pairs	90.585	339.4	168	84	168


Fixed256 fit/192 development gestures, four passes/1,024 target presentations/64 Adam updates; same initialization, causal packets, draws and minimum-devNLL selection. Unchanged p16/L2/H2/pool2 architecture:15,523 parameters/eight available receivers; inference168 scored keys/84 commits/168 candidate values per21-event prefix. Fitting activity includes shadow replay where present.

Promotion gate FAILS: NLL improvement -0.043869, accuracy decline-2.604pp, whole-fit work ratio1.0047. Required gain>=.02, decline<=1pp, work ratio<=1.10, peak RSS<900,000KiB. A failed gate stops unchanged confirmation/scale-up; all passes and negative findings retained.

All forward/replay, backward, normalization/clipping and Adam paid;2FLOPs/MAC plus unit specials. Strong984-fit RBF73.44%/.7065 and compact66.67%/.9030 references have unequal fitting data; their solver FLOPs remain unmeasured. See common-unit full-fit tables above. One seed, no official test or superiority claim. local_dvs_credit_comparison_pairs_20261002T174000Z.json; theory77/80-84.

Appendix B. Completed matched credit pilot: state choice / seed6

Ours = Sleeping Machines



local dvs credit comparison state choice s6 20261002T192100Z learning

Credit	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
local	54.17	1.3059	2.285696	2.2321	0.5917
state choice	58.85	1.1940	3.059796	2.9881	0.5917

Credit	Workflow seconds	Peak RSS MiB	Keys / fit target	Commits / fit target	Values / fit target
local	94.464	338.9	168	84	168
state choice	103.948	339.8	336	168	336


Fixed256 fit/192 development gestures, four passes/1,024 target presentations/64 Adam updates; same initialization, causal packets, draws and minimum-devNLL selection. Unchanged p16/L2/H2/pool2 architecture:15,523 parameters/eight available receivers; inference168 scored keys/84 commits/168 candidate values per21-event prefix. Fitting activity includes shadow replay where present.

Promotion gate PASSES: NLL improvement 0.111931, accuracy decline-4.688pp, whole-fit work ratio1.3387. Required gain $\geq$.02, decline $\leq$ 1pp, work ratio $\leq$ 1.50, peak RSS<900,000KiB. A failed gate stops unchanged confirmation/scale-up; all passes and negative findings retained.

All forward/replay, backward, normalization/clipping and Adam paid;2FLOPs/MAC plus unit specials. Strong984-fit RBF73.44%/.7065 and compact66.67%/.9030 references have unequal fitting data; their solver FLOPs remain unmeasured. See common-unit full-fit tables above. One seed, no official test or superiority claim. local_dvs_credit_comparison_state_choice_s6_20261002T192100Z.json; theory77/80-84.

Appendix B. Completed matched credit pilot: state choice / seed7

Ours = Sleeping Machines



local dvs credit comparison state choice s7 20261002T193500Z learning

Credit	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
local	55.21	1.3065	2.285696	2.2321	0.5919
state choice	55.21	1.3362	3.059796	2.9881	0.5918

Credit	Workflow seconds	Peak RSS MiB	Keys / fit target	Commits / fit target	Values / fit target
local	96.754	338.9	168	84	168
state choice	101.906	339.5	336	168	336


Fixed256 fit/192 development gestures, four passes/1,024 target presentations/64 Adam updates; same initialization, causal packets, draws and minimum-devNLL selection. Unchanged p16/L2/H2/pool2 architecture:15,523 parameters/eight available receivers; inference168 scored keys/84 commits/168 candidate values per21-event prefix. Fitting activity includes shadow replay where present.

Promotion gate FAILS: NLL improvement -0.029693, accuracy decline0.000pp, whole-fit work ratio1.3387. Required gain>=.02, decline<=1pp, work ratio<=1.50, peak RSS<900,000KiB. A failed gate stops unchanged confirmation/scale-up; all passes and negative findings retained.

All forward/replay, backward, normalization/clipping and Adam paid;2FLOPs/MAC plus unit specials. Strong984-fit RBF73.44%/.7065 and compact66.67%/.9030 references have unequal fitting data; their solver FLOPs remain unmeasured. See common-unit full-fit tables above. One seed, no official test or superiority claim. local_dvs_credit_comparison_state_choice_s7_20261002T193500Z.json; theory77/80-84.

Appendix B. Completed matched credit pilot: state clock / seed6

Ours = Sleeping Machines



local dvs credit comparison state clock 20261002T174600Z learning

Credit	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
local	54.17	1.3059	2.285696	2.2321	0.5917
state clock	47.40	1.4507	3.060638	2.9889	0.5918

Credit	Workflow seconds	Peak RSS MiB	Keys / fit target	Commits / fit target	Values / fit target
local	94.464	338.9	168	84	168
state clock	99.314	339.5	336	168	336

Fixed256 fit/192 development gestures, four passes/1,024 target presentations/64 Adam updates; same initialization, causal packets, draws and minimum-devNLL selection. Unchanged p16/L2/H2/pool2 architecture:15,523 parameters/eight available receivers; inference168 scored keys/84 commits/168 candidate values per21-event prefix. Fitting activity includes shadow replay where present.

Promotion gate FAILS: NLL improvement -0.144750, accuracy decline6.771pp, whole-fit work ratio1.3390. Required gain $\geq$.02, decline $\leq$ 1pp, work ratio $\leq$ 1.50, peak RSS<900,000KiB. A failed gate stops unchanged confirmation/scale-up; all passes and negative findings retained.

All forward/replay, backward, normalization/clipping and Adam paid;2FLOPs/MAC plus unit specials. Strong984-fit RBF73.44%/.7065 and compact66.67%/.9030 references have unequal fitting data; their solver FLOPs remain unmeasured. See common-unit full-fit tables above. One seed, no official test or superiority claim. local_dvs_credit_comparison_state_clock_20261002T174600Z.json; theory77/80-84.

Appendix B. Actual-write choice: independent-seed gate fails

Seed	NLL improvement	Accuracy gain pp	Fit work ratio	Declared gate
6	+0.111931	+4.6875	1.33867	Pass
7	-0.029693	-0.0000	1.33867	FAIL

Same256 fitting/192 development gestures, four fixed passes and1,024 presentations per arm/seed. Seed6 improves54.17%/1.305937 to58.85%/1.194006. Seed7 gives55.21%/1.306508 local versus 55.21%/1.336201 corrected credit. Every pass and both completed comparisons are retained.

The unchanged method fails its independent confirmation gate. No full984-fit campaign, extra pass extension or best-seed promotion is admitted. Numerical credit correctness remains established for the isolated categorical component; repeatable predictive benefit is not. Averaging seeds does not

override the predeclared gate.

Next distinguish information access, branch/content credit exposure, common-clock noise and private-map dilution on fitting-only frozen probes. Alternative no-grad forwards teach choices without directly teaching losing content on that realization. Winner-sampled branch derivatives can nevertheless be correct in expectation; rare exposure and other approximate teachers require separate tests. A larger pool is not an automatic repair.

Theory87/88 prioritize one diagnosed repair, numerical contracts, a tiny integrated fit and a prespecified matched-work comparison before independent confirmation. Joint versus alternating route/message updates is a conditional schedule experiment, not a change to the coupled temporal forward computation or a currently demonstrated advantage.

Appendix B. Why actual write and joint-clock credit need a test

Frozen model	Opposed directions	Mean value effect	Mean write effect	Mean interaction
native_full_20261002T145300Z	4/15	0.01699	0.006021	0.00681
batched_local_pilot_s6_20261002T164700Z	5/24	0.00693	0.001304	0.00004
batched_pairs_pilot_s6_20261002T164700Z	5/24	0.01225	0.001653	0.00003

Replay32 fixed event/head sites on the first four previously used development clips/model. Both legal winner choices change delivered content and actual persistent writes, with full suffix replay at unchanged current first time and future random draws. Two additional value/write hybrids are diagnostic, not legal routes. Nonzero-comparable direction counts exclude tiny/zero gradients. Absolute effects are not additive attribution percentages.

In the original full model, one alternative improves value-only loss by.01016 but worsens write-only loss by.10514; its legal complete-route effect is+.09193 NLL. The averaged local value teacher favors that harmful alternative. These four original-model clips are correctly classified: this is a conditional credit defect, not attribution of the overall error rate.

Theory78 independently proves a convex affine cross-entropy example where the two-route local teacher reverses exact expected-loss descent. Actual reference/backward contracts pass. The full-state correction replaces one earlier score derivative by $\pi_i(F_i - b)$ minus $\lambda_i T$ times the probability-weighted centered loss. It retains sampled content derivatives and includes common-clock credit through later timing jumps. The isolated-node joint likelihood component is exact in expectation; other route teachers remain local, and variance is unresolved.

Routes, messages and time parameters do learn: frozen original/256-fit local/pair models change35.8%/26.2%/25.2% of16,128 audited choices versus initialization, with nonzero updates in every parameter group. Removing routing score-gradient paths leaves the factual forward unchanged but removes incoming-content/message gradients. This does not prove their fitted utility or diagnose the full quality gap.

Learned-window gradients were tested as smooth primitives; fixed-count repeated-arrival fits are a separate mechanism. New silence-burst/popcorn contracts cover scheduled deadlines, causal EOF, fixed-partition gradients and exact finite timeout-bank risk. No integrated learned window or popcorn fit is established. Hard merge/split credit and an irregular adapter remain prerequisites; fixed50ms packet schedules do not test natural silence.

Sources: completed frozen route-content audit171500Z, suffix audit172000Z, curvature contracts165900Z, silence-burst contracts170500Z and joint-clock contracts173000Z; theory78-81. Numerical diagnostics are not held-out superiority evidence.

Appendix B. Frozen variance audit after the failed clock pilot

Selected checkpoint	Choice RMS	Common-clock RMS	Clock / choice RMS	Summed variance
Local	0.000290	0.385649	1332	2.385095
Joint-clock treatment	0.000469	0.423377	904	3.252833

First four previously used fitting prefixes/model, event9/both layers/heads. Replay both legal delivered values and actual persistent writes through the complete suffix at8/16 exponential-time quadrature nodes, with fixed future draws. Each model pays192 complete legal shadow forwards. No optimizer or new quality score. Magnitudes and summed variances are in score space with the same per-clip/probe convention for both rows.

Decompose the conditional-winner joint score into choice credit $\pi_i(F_i-R)$ and common-clock credit $\pi_i(1-\text{Lambda} \cdot T)(R-b)$. If branch losses are independent of T , the common-clock mean is zero and baseline mismatch alone produces variance $\|\pi_i\|^2(R-b)^2$. Choosing $b=R$ removes it exactly; the constant-outcome numerical contract passes. Time-dependent later route jumps can make the common-clock mean useful, so deleting it universally is not justified.

The current intermediate decoder baseline differs substantially from actual suffix risk. Across these limited probes the common-clock RMS is900–1,300 times the choice RMS; an oracle baseline computed with all quadrature replays removes more than99.9999% of estimated variance. That oracle is a diagnostic ceiling with all replay cost paid, not a deployable cheap baseline, optimizer-noise measurement or fitted improvement.

Eight-versus16-node mean-gradient L2 differences range4.3e-8 to7.3e-7. This agrees numerically on these probes; convergence across discontinuous histories requires a separate test. Per-pass shared draws also prevent assuming ordinary minibatch variance reduction; actual batch covariance was not measured. Neither audit proves the cause of the held-out regression or the primary bottleneck on other benchmarks.

Decision: stop the failed joint-clock campaign. Next isolate exact conditional actual-write choice credit while retaining the native pathwise clock derivative and all temporal/sparse mechanisms. This retains a known approximation for downstream timing jumps; numerical contracts and one small smoke precede any matched fit. A later independent-noise/control-variate comparison needs fresh matched local controls and full recovery/work accounting. No automatic wider capacity or pass extension.

Completed diagnosticlocal_dvs_joint_credit_variance_20261002T175100Z.json;23.205s/307.3MiB peak RSS, one guarded CPU job. Theory82/83 and COUNTERFACTUAL_CREDIT_PLAN.md preserve proof, negative fit and limited diagnostic scope.

Appendix B. Calibration: noisy routing changes global gradient balance

Frozen weights / credit	Total norm	Clip scale	Cosine to local	Replacement norm	
Pilot256 / local	3.5519	0.2815	1.0000	0.0000	
Pilot256 / choice	3.5519	0.2815	1.0000	0.0052	
Pilot256 / clock	3.8068	0.2627	0.8413	2.0875	
Full984 / local	3.3901	0.2950	1.0000	0.0000	
Full984 / choice	3.3915	0.2949	1.0000	0.0215	
Full984 / clock	6.4002	0.1562	0.3940	5.9463	

Frozen weights / credit	Context norm	Message norm	Route norm	Time norm	Decoder norm
Pilot256 / local	3.1646	0.6550	0.0929	0.0239	1.4707

Frozen weights / credit	Context norm	Message norm	Route norm	Time norm	Decoder norm
Pilot256 / choice	3.1646	0.6551	0.0928	0.0238	1.4707
Pilot256 / clock	3.1557	1.0448	1.1304	0.0284	1.4707
Full984 / local	2.9646	0.4905	0.1972	0.0238	1.5568
Full984 / choice	2.9656	0.4922	0.2016	0.0239	1.5568
Full984 / clock	5.4179	1.8514	2.3982	0.0729	1.5568

Same selected weights/first16 fitting prefixes/draw314159/event9 layer0 head0 for each local, actual-write choice+native timing and joint-clock intervention. Factual logits and losses are identical. All gradients summed then divided by16 before computing the native norm1 clip multiplier. No optimizer, fitting selection or new quality score.

For Full984, the noisy joint-clock correction raises total gradient norm3.3901 to 6.4002, rotates its direction to cosine.3940 and reduces the global clip multiplier .2950 to.1562. The decoder gradient is unchanged before clipping, so the global multiplier also attenuates its useful supervised update. This directly measures cross-path balance on one draw; it is not attribution of the held-out regression.

Write-choice credit keeps total norm3.3915, cosine.99998 and clip multiplier.2949 while correcting actual legal memory utility. Positive scaling alone cannot repair a wrong direction or manufacture useful information. Per-sample norm equalization can bias a zero-mean signal; a prefix-only clock baseline or independent calibration requires a separate matched protocol. Shared weights retain both message and timing credit.

Completed auditlocal_dvs_credit_gradient_balance_20261002T191600Z.json;31.254s/356.6MiB. Every intervention's forward/replay, backward and normalization operator coverage passes. No measured expected variance, batch covariance, Adam trajectory or fitted benefit from gradient normalization follows.

Appendix B. Exact probabilities do not remove winner-dependent credit bias

Conditional gradient	Candidate0 score	Candidate1 score
True hard-outcome risk	+.018750	-.018750
Original teacher expectation	+.018750	-.018750
Exact-pi teacher expectation	-.028125	+.028125

Four completed contracts use actual backward implementations and independent autograd. A convex quadratic loss. $5 \cdot (\text{value} - .6)^2$ with values0/1 and race rates1/3 gives probabilities .25/.75 and legal losses.18/.08. The original local teacher happens to equal the true categorical derivative here; the exact-probability replacement reverses its direction.

Replacing a random rate coefficient by known pi preserves the old expectation when downstream error is fixed. In a nonlinear model that error depends on the winner, its write and subsequent routes. The full gradient is therefore not generally unchanged-expectation or zero variance. In the quadratic witness, replacement optimizes loss at the MEAN value; hard-delivery expected loss also contains a value-variance derivative.

Protocol correction beside theory59 section400: the earlier curie192000Z fidelity audit forces each candidate with its individual arrival time, changing identity AND timing. Conditional first time has mean.25 for both winners; unconditioned individual times have means1/.333333. Its original sign/magnitude numbers remain combined-intervention evidence, not an exact fixed-time route-credit diagnosis. Existing models/results are preserved.

Theory89; contracts204000Z. Exact-pi remains an empirical candidate separately owned by the other host. Neither this counterexample nor the earlier diagnostic establishes general superiority of either local teacher. Correct counterfactuals match actual sparse delivery, legal writes, time law and

downstream utility.

Appendix B. Evolution offset preserves signal-time coupling

Schedule	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
alternating	25.00	2.3294	0.107416	2.237827	0.596655
joint	25.00	2.3032	0.107803	2.245888	0.596655

Phase=frequency*physical_age+beta, beta=pi*tanh(raw_offset). Physical-age damping, emission delays, receiver readiness and stored timestamps remain. The32 offsets calibrate reception phase; they do not create a separate signal clock. Time still drives representation evolution and receives timing derivatives. Initial zero reproduces the original model.


Both24fit/eightdev/two passes/48 presentations/four updates,U16+partialU8; p16/L2/H2/pool2 has15,555 parameters/eight receivers,168 inference keys/84 commits/168 candidate values over21 events,720 persistent-state bytes. Joint fitting NLL2.64095 to2.09915;alternating to2.14262. Both smokes take21.0 seconds. Tiny development scores establish readiness, not quality ranking or advantage against strong controls.

Alternating message/route windows train the offset only with message parameters. Inactive parameters and Adam momentum/steps remain fixed. Shared content/context maps receive the full derivative every window; their updates can still change routes. No content detach or changed credit estimator. Private blocks receive fewer updates and active-gradient clipping differs at the same total presentation count.

Four contracts205000Z: zero-offset state/logit/original-gradient identity; nonzero serial/batched identity; physical-age/offset finite differences and independent directions with damping; phase ownership and actual interrupted recovery/accounting. Theory90. Pure rotation can make offset/time locally redundant; constant phase is a known calibration operator, not a universal capacity or convergence theorem.

Appendix B. Completed coupled-time offset pilot: seed6

Ours = Sleeping Machines



local dvs evolution offset comparison s6 20261002T210100Z learning

Schedule	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
local	54.17	1.3059	2.285696	2.232125	0.591737
joint	57.29	1.2582	2.287775	2.234155	0.596711
alternating	53.65	1.2630	2.281584	2.228110	0.596655

Offset schedule	NLL improvement	Accuracy gain pp	Fit work ratio	Gate
joint	+0.047767	+3.1250	1.000910	Pass

Offset schedule	NLL improvement	Accuracy gain pp	Fit work ratio	Gate
alternating	+0.042929	-0.5208	0.998201	Pass


Fixed256 fit/192 dev/four passes/1,024 presentations/64 Adam updates. Same native initial predictions, data, draws and minimum-devNLL selection. Native15,523 parameters versus15,555 with offsets; eight available receivers and identical selected activity. All offset operators, backward, clipping and active Adam work are charged.

Candidate gate requires $\geq .02$ NLL improvement, ≤ 1 pp accuracy decline, ≤ 1.50 fitting work ratio and $< 900,000$ KiB RSS. Selected schedule: joint. Seed7 joint confirmation fails:1.327347 versus native1.306508NLL and2.0833pp accuracy decline. No unchanged full-data stage or extra epochs is admitted. Two seed6 schedules were declared before either result; all failures remain.

Preserved strong984-fit RBF73.44%/.7065 and compact66.67%/.9030 controls have unequal fitting data; solver FLOPs unmeasured. Alternation also changes private-block update counts/clipping, so this is an algorithm comparison, not isolated interference attribution. Physical time still drives content; no official test or supremacy claim.

Appendix B. Completed coupled-time offset pilot: seed7

Ours = Sleeping Machines



local dvs evolution offset comparison s7 20261002T210800Z learning

Schedule	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
local	55.21	1.3065	2.285696	2.232125	0.591863
joint	53.12	1.3273	2.287775	2.234155	0.596823

Offset schedule	NLL improvement	Accuracy gain pp	Fit work ratio	Gate
joint	-0.020839	-2.0833	1.000910	FAIL

Fixed256 fit/192 dev/four passes/1,024 presentations/64 Adam updates. Same native initial predictions, data, draws and minimum-devNLL selection. Native15,523 parameters versus15,555 with offsets; eight available receivers and identical selected activity. All offset operators, backward, clipping and active Adam work are charged.

Candidate gate requires $\geq .02$ NLL improvement, ≤ 1 pp accuracy decline, ≤ 1.50 fitting work ratio and $< 900,000$ KiB RSS. Selected schedule: none. Seed7 joint confirmation fails:1.327347 versus native1.306508NLL and2.0833pp accuracy decline. No unchanged full-data stage or extra epochs is admitted. Two seed6 schedules were declared before either result; all failures remain.

Preserved strong984-fit RBF73.44%/.7065 and compact66.67%/.9030 controls have unequal fitting data; solver FLOPs unmeasured. Alternation also changes private-block update counts/clipping, so this is an algorithm comparison, not isolated interference attribution. Physical time still drives content; no official test or supremacy claim.

Appendix B. When added representation freedom can help

Theory91 derives a local, weighted score/content fitting criterion. With existing Jacobian A, added offset Jacobian B and desired change b, first fit existing coordinates; let r be the remaining residual. Eliminating those coordinates gives offset sensitivity $h=B\text{-transpose}\cdot r$ and positive definite Schur matrix S. The exact best local improvement is one half $h\text{-transpose}\cdot S\text{-inverse}\cdot h$. It is strict precisely when h is nonzero.

Additional coordinates can address a missing useful direction, or lower a movement penalty along an already available direction. The second case improves conditioning without proving increased representational rank. Shared offsets act at many receptions, whose demands may cancel. A local post-race offset preserves that already selected race while changing future state and routes; the full sequence remains coupled.

For a damped rotating carrier z, the age/phase Jacobian determinant is minus damping rate times squared message norm. The raw bounded offset adds $\pi\cdot\text{sech}(\text{raw-offset})^2$. Nonzero damping gives independent local age and phase directions; weak damping, decayed messages or saturated offsets can still make them poorly conditioned. Physical-time derivatives and actual timestamps remain intact.

Four completed independent matrix contracts check 16 direct-versus-eliminated solves, zero-projection/no-benefit, redundant coordinates with lower penalized movement cost, and the rotor determinant. These are mathematical identities, not a demonstrated learning cure. Exact gradients of one loss can have opposed route/content terms and still yield descent when added; opposition alone does not justify deleting credit.

Contracts211000Z; theory91. Joint offset seed6 gains .047767 NLL at 1.000910 fitting-work ratio; unchanged seed7 loses .020839 and 2.0833 accuracy points. No full-data stage is admitted. The local criterion does not guarantee validation generalization, convergence of surrogate race credit or practical superiority.

Appendix B. Other-host completed native capacity comparisons

Model	Dev accuracy %	Dev NLL	Whole fit GFLOPs est.	Fit MFLOPs / target est.	Infer MFLOPs / prefix est.
Calibrated counts	58.85	1.3734	Unmeasured	Unmeasured	Unmeasured
rbf_C10_g1	73.44	0.7065	Unmeasured	Unmeasured	Unmeasured
Ours original	65.10	0.9632	20.075	2.550	0.592
Compact prototype33	66.67	0.9030	Unmeasured	Unmeasured	Unmeasured
Native p16/L2/pool2	66.15	1.0420	20.075	2.550	0.592
Native p16/L2/pool8	62.50	1.0880	57.298	7.279	0.874
Native p16/L4/pool2	57.29	1.2216	37.869	4.811	1.068

Native capacity	Parameters	Available receivers	Keys / prefix	Commits / prefix	State bytes
Native p16/L2/pool2	15523	8	168	84	720
Native p16/L2/pool8	42091	32	672	84	2376
Native p16/L4/pool2	28539	16	336	168	1296

These completed stages use the same984-fit/192-dev, eight fixed passes/7,872 presentations/496 updates and seed6 as the packet-clock reference. Wider eligible pool and extra depth retain native timing/races/persistent addressed state/key-value separation/counterfactual credit. All scored keys

and losing values remain charged.

Pool8 scores 62.50%/1.0880NLL versus pool2 reference 66.15%/1.0420, at 57.2985 versus 20.0747 whole fitting GFLOPs. Depth4 scores 57.29%/1.2216 at 37.8686 GFLOPs. Neither completed variant improves the baseline. Extra private capacity also adds untied trainable maps; this is not isolated counterfactual-support or useful-capacity evidence. Other-host tied-map/replication work remains pending and separate.

Completed curie_dvs_clock_p16d2pool8 and p16d4pool2 results 172500Z; theory 59§§396/398. Conventional solver work is unmeasured, not zero. Raw controls retained in the full comparison above; selected references here use the same column units. No official-test or supremacy claim.

Appendix B. Strong compact controls rule out an easy storage claim

Lowest-NLL cell / family size	Dev accuracy %	Dev NLL	Model KiB	In budget	CPU ms / prefix
class_prototype_m11_g1_C10	59.38	1.0926	39.2	True	0.210
class_prototype_m22_g1_C10	66.67	0.9549	67.2	True	0.228
class_prototype_m33_g1_C10	66.67	0.9030	95.2	True	0.230
class_prototype_m44_g1_C10	67.19	0.8865	123.2	False	0.230
nystroem_m8_g0.25_C10	55.21	1.3055	32.3	True	0.836
nystroem_m16_g0.25_C10	62.50	1.1175	53.4	True	0.859
nystroem_m24_g0.25_C10	62.50	1.0640	75.1	True	0.852
nystroem_m32_g1_C10	66.15	0.9954	97.2	True	0.854

Native uncompressed saved-model budget 103.9KiB. The fixed 72-cell grid uses random 8/16/24/32 Nyström landmarks or 1/2/3/4 learned prototypes per class, three gamma scales and three logistic C values. Prototypes, landmarks and normalization use fitting data only; each model includes its transform in the same uncompressed joblib serialization. Every cell remains in the completed JSON.

Selected in-budget control class_prototype_m33_g1_C10: 66.67%/0.9030NLL, 95.2KiB, 0.230ms/prefix. This exceeds the original native 65.10%/0.963161 while using less storage and far less CPU time. Thus the original native model loses to the prototype control in this bounded-storage region.

The whole compact grid costs 10.043s/183.9MiB peak RSS; common raw preprocessing still costs 123.976s. Shared prototype construction, kernel features, all fitting/tuning and scoring are paid in workflow wall. Solver FLOPs are unknown, not zero. The table summarizes each family/size by its minimum development NLL, with all gamma/C settings and probabilities retained.

Exploratory development calibration, not independent model confirmation. Single timing passes are not repeated latency claims. Class prototypes are supervised fitting controls, not target-derived inference inputs. Neither the full kernel nor these compact alternatives is omitted when evaluating the next native initialization.

Appendix B (continued). Diagnostic: count receivers over the temporal carrier

Labelled diagnostic, not the integrated native architecture. The input-gated temporal carrier supplies the base predictive to the same escape-race count cascade (Theory §§386–388). It tests whether sufficient-statistic receivers remove the memorization tax: if counts hold the exact local statistics, a small learned base should lose far less than the carrier alone does.


Model	Fit chars / passes	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Infer MFLOPs / char ↓
Minimal base w2/d1 + counts K5 + escape gate	131,072/4	2.178	3.6	0.007	0.0028
Minimal base w2/d1 + counts K5 + gate + message	131,072/4	2.124	20.5	0.039	0.0177
Carrier alone w32	131,072/4	2.848	80.9	0.154	0.0494
Carrier w32 + counts K5	131,072/4	2.313	82.3	0.157	0.0504
Same, untrained base w32	131,072/0	2.371	Not trained	Not trained	0.0504
Carrier w32 + counts K5 + escape gate	131,072/4	2.136	83.8	0.160	0.0518
Carrier w32 + counts K5 + gate + message	131,072/4	2.127	100.8	0.192	0.0667
Carrier alone w128	131,072/4	2.587	1046.7	1.996	0.6389
Carrier w128 + counts K5	131,072/4	2.313	1048.1	1.999	0.6399
Same, untrained base w128	131,072/0	2.370	Not trained	Not trained	0.6399
Carrier w128 + counts K5 + escape gate	131,072/4	2.130	1049.6	2.002	0.6412
Carrier w128 + counts K5 + gate + message	131,072/4	2.129	1066.6	2.034	0.6562
Carrier alone w256	131,072/4	2.572	4025.5	7.678	2.4571
Carrier w256 + counts K5	131,072/4	2.316	4026.9	7.681	2.4580
Same, untrained base w256	131,072/0	2.370	Not trained	Not trained	2.4580
KN counts, frozen o5	131,072/1	2.349	Not FLOPs	Not FLOPs	Not FLOPs
Calibration ceiling: adaptive interpolated KN o8	131,072/1	2.101	Not FLOPs	Not FLOPs	Not FLOPs

Same 8,191 development targets; seed 6, one seed per row; same depth, chunk, learning rate and passes per width. Predeclared: P1 composed w128 < 2.326; P2 composed w32–w256 gap < half the carrier gap; both hold formally for the scalar cascade, but the trained bases alone score 8.17 (w32) / 11.34 (w128) bpc, worse than uniform: the standalone base is trained as a conditional residual, so this alone cannot establish an inert base (Theory §389.1). With the escape gate (+ count message) every width reaches 2.12–2.14, and a minimal 2-wide, one-layer base matches w128 (2.124 vs 2.129) at 1/50 of the work: the gain is learned count smoothing, not the temporal carrier. Count references are near-optimal estimators at this size. The rows above exceed the frozen-KN and Witten-Bell references, and the stronger stream-adaptive interpolated Kneser-Ney is the calibration ceiling here (§393). This is a mechanism diagnostic in a regime where counting is expected to be near-optimal for any learner, not an architecture verdict (§394). Count increments/lookups (5 per target) are integer table work outside FLOPs. Exploratory development evidence; no comparable-quality Transformer claim.

Appendix B (continued). Ours: native temporal-core language

27-character text8 content enters one conversation address in the native event core. Independent receiver heads, persistent state, learned delays, evolving channels and losing-route credit remain; no historical KV attention or dense carrier is added. The receiver maps now learn across all token marks.

Ours = Sleeping Machines



native language quality work


Ours	Fit chars / passes	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Infer MFLOPs / char ↓
d16/p2/s6	2,048/4	3.765	3.778	0.461	0.0974
d16/p2/s6	8,192/4	3.557	15.116	0.461	0.0974

Fixed passes and frozen 8,191-target development selection; official test untouched. Existing stronger receiver and dense-reference evidence is preserved in the common ledger. Architecture, capacity and weight sharing differ from the KV variants. Work uses audited representative optimizer windows including backward, teachers and Adam; special functions have unit weight here, with separate JSON counts. These are exploratory scaling points, not a frontier score or a measured energy advantage. With H2/depth8/pool2, available receivers are 32, selected commits 16/token, scored keys and training proposals 32/token. Parameter counts and occupied state are retained in completed records.

Appendix B (continued). Ours: native event/state capability

Each source writes three signed marks then receives an explicit causal query. Fits use 512 queries / 2,048 input events per pass, eight passes; development has 256 queries. All source populations are occupied. The order task predicts the last two signs; the timing task predicts a two-mode temporal trace. All events, timestamps and query flags are observed; target labels never address a module.

Ours = Sleeping Machines



native event quality work

Ours: task / sources / seed	Credit / time	Dev accuracy ↑	Whole fit GFLOPs ↓	Fit MFLOPs / query ↓	Infer MFLOPs / event ↓
order/S4/s6	CF/observed	100.0%	7.099	1.733	0.0923

Equal declared query/event budgets; available receivers scale with source population while per-event key scores and selected commits are logged explicitly. All fitting arithmetic is summed from executed operators, including complete producer graphs, losing proposals, clipping and Adam; CPU numerical clocks and unit-weight specials are included. Bootstrap intervals by population and raw state-clearing/time/long-gap probes remain in the completed JSON. More addresses change training exposure. Synthetic capability and bounded activity do not establish superiority over timestamp-aware recurrent/attention controls, real-stream generalization or clockless energy savings.

Appendix B (continued). Ours: occupied state and activity

Every source stores three content events and receives a query. These are useful-state tests rather than padded unused capacity. Event budgets, active depth, head count and candidate pool are fixed. More addresses reduce observations per local parameter; addressing is input information that strong controls must also receive.


Ours: task / sources / seed	Available / occupied receivers	Commits / scores per event	Persistent state KiB	Parameters
order/S4/s6 CF/observed	128/122	16/32	9.14	159,716
Ours: task / sources / seed	Whole fit GFLOPs	Fit MFLOPs / input event	Infer MFLOPs / query	
order/S4/s6 CF/observed	7.099	0.433	0.369	
Ours: task / sources / seed	Selected dev accuracy	Clear-history accuracy	Independent confirmation	
order/S4/s6 CF/observed	100.0%	27.7%	Not read	

Occupied receivers and stored float bytes come from the audited inference population, not parameter storage, Python metadata, graphs or host RSS. Inference/query includes all four producer/query events. Full fitting includes every event and optimizer operation. Arithmetic plus unit-weight specials; traffic/RNG/energy separate. Development selects checkpoints; independent confirmation is read once only on fixed promoted seed protocols. Population bootstrap intervals in the quality plot condition on the selected model and do not correct development selection or substitute for independent seed uncertainty.

Appendix B (continued). Ours: compact delayed learning

Available historical content is not automatically learned historical content. After graph detachment, a retrieved old key/value can affect the prediction while its write map receives no later loss credit. Increasing the ordinary credit span from 16 to 64 did not improve the completed matched 2K pilot.

Ours = Sleeping Machines



historical write credit flow

The tested integrated construction saves the normalized feature ϕ of each historical write. Later queries send a compact producer teacher to sealed key/value maps while preserving independent heads, temporal races, incoming content and persistent addressed state.

Ours: added teacher	Factorization	Extra work / query
Old keys	$\alpha q (\sum e_i \phi_i)^T$	$O(C_{old} d + d^2)$
Old winning value	$\alpha a \phi_w^T$	$O(d^2)$
Saved eligibility	One detached d-vector / write	50% extra K/V float storage

A shared offset of all race scores leaves winner probabilities unchanged but rescales arrival time. Content contrasts can teach which value wins; a separate clock signal can teach when it arrives and how persistent state evolves. The new write teacher retains both signals.


The shared query factors all admitted old key teachers into one matrix update. Live writes keep ordinary gradients without duplicate credit. Six read-only numerical checks and guarded full eight-block update/recovery contracts passed; disabling the new teacher exactly reproduces parent outputs, RNG, gradients and two Adam windows.

This is exact for a hypothetical common perturbation of historical write maps with saved features fixed. Transport to current maps is a delayed local surrogate: old weight versions and omitted representation paths remain limitations. Additional learning arithmetic and eligibility traffic are real costs. Existing counterfactual score credit is retained. Fitting work uses representative audited windows, with exact credit-coverage counters logged separately. Matched quality/work results require completed fits; improved language accuracy and native clockless learning are not established by these contracts.

Appendix B (continued). Ours: historical write credit

H2, d32/head, 8 event blocks, pool2; 2,048 fitting characters / 4 passes; 8,191 frozen development targets, seed6. Credit16, Adam U64, lr0.002, warmup512. Forward races/history are preserved; sealed writes gain a compact local producer teacher.

Ours = Sleeping Machines



historical write quality work 003fd093a8

Ours: credit	Dev bpc ↓	Whole fit GFLOPs ↓	Fit MFLOPs / target ↓	Inference MFLOPs / char ↓
parent	3.733	22.753	2.779	0.5059
$\alpha=0.25$	3.722	23.469	2.866	0.5059
$\alpha=1$	3.724	23.464	2.866	0.5058

A historical normalized write feature ϕ is retained alongside each cached key/value. The old key-map teacher factors as $\alpha q (\sum e_i \phi_i)^T$: one weighted feature sum and one matrix outer product per query. The old winning value map receives $\alpha a \phi_w^T$. Live writes keep ordinary credit without duplication; no full-history graph is reopened.

One extra feature vector raises K/V float storage by 50%. The update is exact for a common hypothetical perturbation of historical maps at fixed saved features; transporting it to current maps is a delayed local surrogate with stale weight versions and omitted representation paths. Additional teacher/optimizer arithmetic is counted; feature traffic and physical energy are separate. Parent reuse requires exact forward/RNG/gradient/Adam nesting at $\alpha=0$. One-seed exploratory development selection, not comparable-quality supremacy.

Appendix B (continued). Ours: frozen information-flow diagnosis

Saved selected H2/H4 eight-block 8K checkpoints; 256 targets from the development prefix. Weights remain fixed. Each intervention removes one path only during this short evaluation; these are diagnostic probes rather than refitted architecture comparisons.

Ours	Frozen intervention	Window bpc ↓	Context RMS mean
H2	unaltered	3.316	1.770
H2	source off	4.013	1.676
H2	kv off	3.312	1.750
H2	channel identity	3.773	1.488
H4	unaltered	3.331	2.144
H4	source off	4.001	2.066
H4	kv off	3.375	2.073
H4	channel identity	3.786	1.477

Removing source-message carry and learned channel mixing worsens both saved checkpoints in this window. These content/state paths are useful here; short-window removal probes do not warrant discarding them or demonstrate a refitted architecture advantage.

A route-credit audit replays each admitted historical value at three depths/head0 for the two-head model. It holds one realized race time and continuation seed fixed, then compares the centered local content teacher with the conditional categorical loss gradient.


Ours: local audit	Probes	Mean cosine	Opposed directions	Mean oracle gap (nats)
H2 / head0	24	0.717	1 / 24	0.071

Higher cosine means better alignment in this narrow replay. The oracle gap is realized loss minus the best candidate replay, not achieved improvement. Different routing can change subsequent candidate/RNG paths. The audit conditions away time derivatives and does not estimate the full expected gradient, long-history utility or discovery coverage. Interventions have no refitting and no confidence intervals; do not use these window scores as promotion metrics or dense-model superiority evidence.

Appendix B (continued). Ours: separate online neural learning

Both arms start from the same selected integrated-model checkpoint and maintain persistent event memory on the same new development stream. The frozen arm retains its parameters. The online arm updates the complete neural backbone after making the causal predictions in each 16-character block, with no replay.

Ours = Sleeping Machines



integrated online language

Ours: mode	Stream bpc ↓	Whole stream MFLOPs ↓	Parameter change L2	Updates
Ours: frozen	3.191	210.903	0.000	0
Ours: online	3.096	2,264.163	1.551	512

The 8,191 targets come from [90,065,536, 90,073,728). Learning rate 0.0001 was fixed before this stream; Adam starts with fresh moments. Both arms use paired block race noise. Parameter updates precede only future blocks; the current target cannot change its own prediction. Their contexts can diverge after learning.

This measures block-delayed online adaptation, not instant per-character updates or a frozen official-test score. One inherited model/window/rate. All neural parameters may learn, including content, keys, clocks and retention; this differs from the earlier statistical expert-mixing-only ablation. Inherited fitting is additional and identical in both arms. Work estimates observe the actual first/middle/last blocks and include counterfactual learning, backward, clipping and Adam; RNG, indexing and physical traffic remain additional. Lower online loss, if observed, is evidence only for this protocol.

Appendix B (continued). Separate statistical language baseline

This count/copy predictor does not use the learned Sleeping Machines event backbone. It is a separate statistical system: order-0 through order-6 counts, backoff probabilities and a bounded causal copy cache, combined by learned mixing weights. Its quality/work results must not be attributed to the event architecture.

The statistical predictor and saved neural references score the same 999,999 character targets, starting from a cold context. Parameters are frozen during testing. Earlier observed test characters can supply causal context, including the mixture's bounded 256-character copy cache. Lower bits per character means better prediction.

Predictor	Test bpc ↓	Fitting and selection budget	Estimated training FLOPs ↓
Ours: separate statistical count/copy baseline	1.727	10M count fitting + three 1M mixing-rate trials	4.47G + integer count construction
Ours: count/copy plus causal word context	1.719	10M count fitting + three 1M mixing-rate trials	4.80G + integer count construction

The neural totals charge all original optimizer steps that produced the inherited E174 checkpoints: forward, estimated 2x-forward backward, clipping and Adam. Alignment evaluation is excluded. The statistical floating estimate charges expert probability preparation and all three mixing-rate trials, including local gradients and weight updates. Exp/log/root evaluations count as one operation in

these language estimates. Count construction additionally uses about 80M integer count presentations, plus sorting, lookup and hashing; that work is not quantified as FLOPs. The floating totals alone cannot establish total compute, runtime or energy savings.

The count/copy mixture improves on LSTM by 0.072 bpc and Transformer by 0.181 bpc. These results establish useful specialized prediction; generic learned representations are assessed in the separate language screen.

The statistical baseline's count arrays occupy 66.55 MB. Vocabulary, capacities, optimization and fitting budgets differ from the neural references. The comparison does not measure total training energy or a matched-capacity advantage.


All three use the same historical 27-character alphabet. Modern shared subword tokenization is a separate comparison gate for the learned event architecture, described later in this appendix.

One exploratory seed. Text8 offsets: count fitting [0,10M), mixing-weight validation [90M,91M), test [95M,96M); test index zero is excluded for all three predictors. The mixture selects its update rate on validation. Saved neural weights are unchanged. Results: E173/E174; stream contract: E175. At 90M training characters, reference test scores are 1.661 for the LSTM, 1.604 for the four-layer Transformer. Estimated training work (forward, backward, Adam and gradient clipping): LSTM: 3.89 PFLOP, TF: 8.00 PFLOP. Shape-based estimates count multiply-add as two operations; backward is approximated as twice forward. Validation/test inference is excluded. These are single-seed comparisons; capacities and fitting budgets are not matched. Calibration (Theory §381): untuned modified Kneser-Ney counts of the same 90M characters score 1.652 on the same test targets, so these controls sit near count level and are not frontier bars. Target-leaked historical mixtures have been quarantined; a corrected 90M mixture comparison remains open.

Appendix B (continued). Learned language and depth

A bounded screen trains the common event backbone to predict the next character. Both configurations use width 32, the same 8,192 training characters, four passes, 32-character contexts and 1,024 validation predictions. All eight layers' value, route and memory-time parameters update. Lower bits per character means better prediction.

Ours = Sleeping Machines



e133 generic language

Ours: depth	Learned parameters	Validation bpc: lower is better	Total CPU wall time
1	7,671	3.464	27.4 s
8	53,430	3.395	170.5 s

Eight layers improve validation loss from 4.752 to 3.395 bpc, versus 3.464 with one layer. Shuffling preceding characters while preserving the last character, count and timestamps increases the deeper model's loss to 3.805; replacing preceding context raises it to 3.777. These frozen input probes show context sensitivity, not a retrained baseline comparison.


The deeper model has more parameters and takes more CPU time. The quality/time panel includes fitting and evaluation, with backpropagation and optimizer updates executed during fitting. These bounded-query models replay preceding context. The following persistent implementation consumes each character once. Physical memory traffic and joules remain unmeasured.

One seed; different parameter counts. This establishes a generic learned-language foothold and a small depth gain, not competitive large-scale representation, a matched tuned dense-model advantage or a scaling law. The statistical 10M-character mixture remains a separate result. Official test data are untouched. E133 preserves commands, source/data hashes, layer diagnostics and work coverage.

Appendix B (continued). Persistent learned language

The event-state language model consumes each character once and retains local modal memories and its delayed-message queue. Chunk boundaries truncate learning credit without discarding the observed history. Only actual event arrivals evaluate layers; text time is measured in token intervals.

Ours = Sleeping Machines



e176 stream language learning

Ours: representation	Parameters	Fitting bpc ↓	Validation bpc ↓	Deliveries/pass
Characters	28,403	3.142	3.351	65,784
Causal prefix tokens	35,163	3.085	3.517	45,816

Validation loss falls from 5.329 to 3.351 bpc. All eight layer teachers are nonzero in every fitting pass. Each pass consumes 8,223 characters including warmup and makes 65,784 block deliveries. The result establishes learning with persistent causal state and no prefix replay.

A train-only 131-token prefix dictionary reduces layer deliveries by 30.4% and observed CPU time by 26.2%, while validation bpc is 3.517. Exact partial-token marginalization scores identical raw targets. Its larger vocabulary fits better but generalizes less well in this small screen: compression alone does not explain or resolve the quality gap.

One seed; 28,403 parameters, width 32, sixteen temporal modes per block. Four passes, 128 Adam steps/pass, credit truncated every 64 characters, 31 warm characters and 1,024 validation targets. Target offsets match the bounded E133 screen; topology, capacity, history and update counts differ, so this is not a matched intervention. Total CPU wall time 364.9 s including fitting/evaluation; peak RSS 364.8 MiB. These are event counts and observed resources, not total arithmetic, physical memory traffic or energy. No official test or large-corpus claim. E176.

Appendix B (continued). Learned event language: training and inference work

These counts belong to the learned eight-layer persistent Sleeping Machines event model that reaches 3.351 validation bpc. It has no count/copy/word experts. Primary counts describe the logical event algorithm; simulator dispatch/allocation is excluded.

Ours: fitting stage	Estimated arithmetic FLOPs
Forward and loss	1.96G
Backward	4.77G
Gradient clipping	43.52M
Optimizer	173.95M
Warmup	7.42M
Total	6.96G

The entire completed budget includes 32,768 fitting targets, 512 Adam/clipping steps and four stream warmups. Special functions add 85.431M evaluations, reported separately from arithmetic FLOPs.

Ours: inference boundary	Per character
Prediction plus NLL-scoring arithmetic	59,741 FLOPs
Additional special functions	1,212 evaluations

A 64-character saved-checkpoint trace measures forward/loss, backward, clipping and warm Adam; the fitting ledger scales those stages by the original 512 steps and separately charges warmup. This is a representative arithmetic estimate, not a whole-run trace. The traced chunk has complete floating-operator formula coverage. Index/queue work, comparisons, memory traffic and evaluation passes are outside the fitting arithmetic boundary.

The 10M-character LSTM/Transformer checkpoints have different data, capacity and achieved quality. Dividing their full fitting budgets by this small run would not establish a fair training advantage. Full learned-event runs must complete before a larger aligned quality/work comparison.

Evidence: `event_language_work/local_event_language_work_20260930T124646Z.json`; quality: E176. Two arithmetic FLOPs per multiply-add. Additional unit-weight special functions would give a different logical-operation total; no event-device runtime or energy has been measured.

Appendix B (continued). Statistical online-learning pilot

The official language comparison freezes parameters during evaluation. A new development-only pilot asks whether causal local learning improves prediction: score each character first, reveal it, then update only the expert mixing weights. Count experts stay frozen and both arms use identical causal copy-cache behavior.

Ours: statistical expert set	Frozen bpc ↓	Online bpc ↓	Extra update FLOPs
Without word	2.605	2.564	2.75M
With causal word	2.537	2.487	3.21M

This pilot reuses a 100,000-character count checkpoint and its learning rate selected on an earlier 2,048-character validation window. It scores 8,191 targets in a fresh 8,192-character development window [90,032,768,90,040,960), excluding the first target. Each adaptive arm performs 8,192 updates, including first-position warmup. Official test data are untouched.

One checkpoint and one window; a specialized adaptive readout, not deep learned TTT or a frontier result. The extra arithmetic column counts only local gradient/weight updates, on top of shared

prediction and inherited fitting costs. Total pilot CPU wall time is 0.587 s. Evidence: [online_language/local_online_language_20260930T121741Z.json](#).

Compute allocation is the next architectural hypothesis

Independent budgets for active width/depth, dormant capacity, temporal memory, retrieval, credit and adaptation may let extra work buy more prediction quality. The new theory derives conditional marginal-value allocation and retrieval-error bounds, and identifies exposure, routing overhead and hardware utilization as possible limits. A small-scale win does not prove a better scaling exponent or a widening frontier advantage.

Next learned-language comparisons need shared modern subword tokenization, Unicode/byte coverage, suitable rotary/relative position, packed optimized Transformer kernels, sparse MoE and modern recurrent/attention-hybrid controls. Existing event memory already uses relative-time rotations. Token coordinates must remain distinct from learned scheduling delays.

Known sequences can be fitted with input-known affine scans and then generated sequentially with persistent state. Sequence-parallel fitting also permits Transformer adaptation between tokens/chunks. Compare frozen and online arms at identical observations and adaptation budgets; charge cache consistency and update work.

[Theory §§280–286](#) and the [frontier compute protocol](#) specify these tests. Modern architecture and larger-scale comparison arms remain proposed, not completed.

Appendix C. Preserved historical results and revisions

Earlier result files remain part of the research record. The following numbers explain older report headlines and why their interpretation changed. They are preserved here with the identified protocol errors; they are not current valid benchmark comparisons.

Earlier statistical language results

The target-dependent partial-word mixtures have been removed from the active results tree and all numerical comparisons. Raw records are quarantined for audit only. The corrected E17310M results are 1.727 without word context and 1.719 with causal word context; a corrected 90M mixture comparison remains open.

Earlier event world-model results

Historical predictor	Day 6 log-likelihood ↑	Day 7 log-likelihood ↑
Ours: event hazard + rate/flow state	-2.184	-1.999
Ours: event hazard + per-type state	-2.159	-1.973
Ours: event hazard + per-type state; finer gap bank	-1.912	-1.734
Saved Transformer Hawkes reference	-1.971	-1.816

The event hazard models use sparse conditional memories and local rate/flow state. Their frozen test parameters and causal state updates are useful mechanisms. However, the event-size threshold was fitted across all seven pilot days, including the evaluation days, for both the event and neural references. Day resets and warmup exclusions also differ. The recorded gap needs fitting-only preprocessing and aligned rescoring before it supports a held-day advantage.

[Source and numerical review](#); [E57 world-model records](#). Preserved timing, composition, retrieval and modular results remain in the main report. New learned-model benchmarks add evidence; they do not erase these earlier runs.

Completed AWS 90M native language comparisons

Native pool4 reaches 1.998 test BPC at T256, versus 2.045 for pool2. Both use payload32/depth4, two heads, linear local value-informed route credit and eight selected addressed writes per character. Pool4 doubles scored keys from16 to32; available receiver capacity increases. These are single-seed results, not replicated advantage.

Model	Params	Test BPC	Fit GFLOPs	Fit MF/target
p32/D4/pool2	108875	2.045356	65,219.6	0.724684
p32/D4/pool4	177019	1.998416	107,606.9	1.195668
p32/D8/skip2	210043	1.983227	127,430.3	1.415934
p64/D4/pool2	422475	1.857306	241,218.9	2.680290
LSTM (6 passes)	1199323	1.661015	3,893,396.0	7.210099
Transformer (4 passes)	1839003	1.704384	2,069,929.8	13.142462

All four native fits use FIT[0,90M), 89,997,312 presentations, 10,986 Adam windows, 64 lanes of128-character reset segments, cosine lr0.004, seed6. DEV[90M,91M) and test[95M,96M), 999,936 targets each; T256 is an additional scorer of the same final weights. Random segments with replacement constitute a pass-equivalent budget, not full unique coverage.

Pool4 improves its matched-window10M score2.345157 to1.998416, a0.346742BPC scaling gain. At90M its0.046941BPC gain over pool2 costs1.649917x estimated fitting work. This is useful capacity with unchanged selected activity, not unchanged total work or an iso-FLOP gain.

All work columns share units and divide by each model's actual fitting presentations. Native estimates extrapolate two complete traced optimizer windows; controls use shape-based forward, 2x-forward backward and approximate clip/Adam. Validation/test, traffic and energy are excluded. Unequal parameters/passes/quality and different estimate conventions prevent comparable-quality supremacy claims.

Completed public archive benchmark campaign

No public benchmark win is established by this campaign. Three fixed TRAIN-only development screens per dataset selected minimum DEV negative log likelihood. Selected configurations were refitted on full official TRAIN at the selected epoch budget, then scored once on official TEST for each seed6/7/8. All nine refits completed; negative results and their checkpoints remain preserved.

Dataset	Test %	Seed SD	Fit GF/run	Fit MF/series
ECG200	72.33	1.53	149.12	114.704
JapaneseVowels	93.96	2.73	173.15	29.150
PenDigits	95.96	0.82	257.73	0.955


ECG200 and JapaneseVowels selected payload32/depth4/pool4 with two heads, at13 and22 epochs respectively; PenDigits selected payload16/depth2/pool2 at36 epochs. These models use temporal races, separate keys/values, sparse addressed private-state writes and linear local message route credit. Full counterfactual suffix replay is absent. Inputs use ordinal synchronous steps; PenDigits coordinates are spatial resampling, not physical event times.

Seed SD measures variation across training seeds on the same TEST examples, not a generalization confidence interval. Matched nearest-neighbor controls were DEV-only and do not supply matched TEST or resource comparisons. Published protocol variants require verification before frontier comparisons. Inference work and energy remain unmeasured; these fitting estimates establish no resource advantage.

Fit GF/run is mean whole-fit work per final seed; MF/series divides by fitting presentations, not unique examples. Completed pilots, screens and final refits together account for 7,162.53GF estimated fitting work. Failed/contract/evaluation/compilation/preprocessing work is additional. Estimates extrapolate one eager optimizer window per epoch; variable padding prevents exact whole-fit accounting. Sources: experiments/results/public_benchmarks/*.json.

Appendix. Native language at 10M: the integrated core, segment-batched

Model (10M)	Params	Steps	Test bpc T128/T256	Whole fit TF est.	Fit MF/char	Infer MF/pos. emulator	Infer MF/pos. winner-only
Ours p16/d8, v1 (610 updates)	54,907	610	2.899 / —	4.04	0.40	0.13	0.09
Ours p16/d8, skip2	54,907	1,220	2.719 / 2.719	4.05	0.41	0.13	0.09
Ours p32/d4	108,875	1,220	2.507 / 2.506	7.22	0.72	0.24	0.16
Ours p32/d8, skip2	210,043	1,220	2.456 / 2.456	14.11	1.41	0.46	0.31
Ours p32/d8/pool4, skip2	346,331	1,220	2.498 / 2.498	23.60	2.36	0.77	0.31
Ours p32/d4/pool1 (control: no selection)	74,803	1,220	2.439 / 2.439	4.88	0.49	0.16	0.16
Ours p32/d4 + route credit	108,875	1,220	2.370 / 2.371	7.24	0.72	0.24	0.16
Ours p32/d4/pool4 + route credit	177,019	1,220	2.343 / 2.345	11.95	1.20	0.39	0.16
Ours p32/d4 + read and write credit	108,875	1,220	2.384 / 2.384	7.28	0.73	0.24	0.16
Ours p64/d4 + route credit	422,475	1,220	2.184 / 2.183	26.79	2.68	0.89	0.60
Ours p32/d8, skip2 + route credit	210,043	1,220	2.326 / 2.326	14.15	1.42	0.46	0.31
Ours p64/d4/pool4 + route credit	689,787	1,220	2.180 / 2.179	44.07	4.41	1.46	0.61
Ours p96/d4 + route credit	940,875	1,220	2.163 / 2.162	58.65	5.87	1.95	1.32
Ours p64/d4 + route credit, 4 passes	422,475	4,882	1.955 / 1.955	107.19	2.68	0.89	0.60
Ours p96/d4 + route credit, 6 passes	940,875	7,324	1.889 / 1.888	352.08	5.87	1.95	1.32
E64 LSTM-256	338,395	1,220	— / 2.171	20.3	2.03	0.68	0.68
E64 Transformer-256x2	1,658,907	1,220	— / 2.427	111.3	11.13	3.71	3.71
E64 Transformer-256x4, 4 passes	3,238,427	4,882	— / 1.908	888.8	22.22	7.41	7.41
E64 LSTM-512, 6 passes	1,199,323	7,324	— / 1.799	432.6	7.21	2.40	2.40



native language frontier

Matched-compute comparisons at 10M (the native row uses no more than the reference's compute)

Reference (budget)	Reference bpc	Best native within budget	Native bpc	Outcome
LSTM-256 (training, 20.3 TF)	2.171	p32/d8, skip2 + route credit	2.326	reference better
LSTM-256 (inference, 0.7 MF/pos)	2.171	p64/d4 + route credit, 4 passes	1.955	native better
Transformer-256x2 (training, 111.3 TF)	2.427	p64/d4 + route credit, 4 passes	1.955	native better
Transformer-256x2 (inference, 3.7 MF/pos)	2.427	p96/d4 + route credit, 6 passes	1.888	native better
Transformer-256x4, 4 passes (training, 888.8 TF)	1.908	p96/d4 + route credit, 6 passes	1.888	native better
Transformer-256x4, 4 passes (inference, 7.4 MF/pos)	1.908	p96/d4 + route credit, 6 passes	1.888	native better
LSTM-512, 6 passes (training, 432.6 TF)	1.799	p96/d4 + route credit, 6 passes	1.888	reference better
LSTM-512, 6 passes (inference, 2.4 MF/pos)	1.799	p96/d4 + route credit, 6 passes	1.888	reference better

Training budget: whole-fit FLOPs estimate; inference budget: per-position FLOPs (native exact winner-only trace, controls shape estimate). Conventions differ (traced vs estimated); single seeds; the native multi-pass rows use more optimizer updates than one-pass references. A native row qualifies for a budget only if its estimate does not exceed the reference's. Native/Transformer: reset T256 windows, 999,936 targets. Saved LSTM: continuous state, 999,999 targets; same test interval, different context. Identical-context LSTM rescoring is pending.

Model (90M)	Params	Updates	Test bpc T128/T256	Whole fit TF est.	Fit MF/char
Ours p32/d4/pool2 + route credit (AWS, 1 passes, seed 6)	108,875	10,986	2.045 / 2.045	65.2	0.72
Ours p32/d4/pool4 + route credit (AWS, 1 passes, seed 6)	177,019	10,986	1.997 / 1.998	107.6	1.20
Ours p32/d8/pool2 + route credit (AWS, 1 passes, seed 6)	210,043	10,986	1.984 / 1.983	127.4	1.42
Ours p64/d4/pool2 + route credit (AWS, 1 passes, seed 6)	422,475	10,986	1.858 / 1.857	241.2	2.68

Model (90M)	Params	Updates	Test bpc T128/T256	Whole fit TF est.	Fit MF/char
Ours p64/d4/pool2 + route credit (AWS, 4 passes, seed 6)	422,475	43,945	1.801 / 1.800	964.9	2.68
Ours p96/d4/pool2 + route credit (AWS, 4 passes, seed 6)	940,875	43,945	1.783 / 1.783	2112.5	5.87
E64 LSTM-512, 6 passes (AWS)	1,199,323	65,917	— / 1.661	3893	7.21
E64 LSTM-512, 1.4 passes (AWS)	1,199,323	15,380	— / 1.729	908	7.21
E64 Transformer-256x4, 4 passes (AWS)	3,238,427	43,945	— / 1.604	8000	22.22
E64 Transformer-192x4, 0.8 passes (AWS)	1,839,003	8,789	— / 1.780	946	13.14
E64 Transformer-192x4, 1.75 passes (AWS)	1,839,003	19,226	— / 1.704	2070	13.14
E64 Transformer-256x4, 1 passes (AWS)	3,238,427	10,986	— / 1.811	2000	22.22

90M rows: text8[0:90M], same test interval. Native and Transformer use reset T256 windows; LSTM uses continuous state. The native rows state their actual pass counts for the segment-batched protocol on AWS (compiled, 64 x 128 windows, lr .004 cosine); the references are multi-pass with larger models and are listed for scale, not as matched comparisons.

Native model	Slots	Memory scalars	Writes	Keys	Emulator values	Winner values
p16/d8, v1 (610 updates)	32	512	16	32	32	16
p16/d8, skip2	32	512	16	32	32	16
p32/d4	16	512	8	16	16	8
p32/d8, skip2	32	1024	16	32	32	16
p32/d8/pool4, skip2	64	2048	16	64	64	16
p32/d4/pool1 (control: no selection)	8	256	8	8	8	8
p32/d4 + route credit	16	512	8	16	16	8
p32/d4/pool4 + route credit	32	1024	8	32	32	8
p32/d4 + read and write credit	16	512	8	16	16	8
p64/d4 + route credit	16	1024	8	16	16	8
p32/d8, skip2 + route credit	32	1024	16	32	32	16
p64/d4/pool4 + route credit	32	2048	8	32	32	8
p96/d4 + route credit	16	1536	8	16	16	8
p64/d4 + route credit, 4 passes	16	1024	8	16	16	8
p96/d4 + route credit, 6 passes	16	1536	8	16	16	8

Native mechanism counts per input position. Memory scalars are available unit-value storage per lane; timestamps, readiness bits and source context are additional. One value is delivered per selected head/layer write. Every candidate key and proposal value is computed before selection in the emulator; the winner-only evaluator computes only selected proposals and also caches one key-read vector per slot. Every key is still scored. These are shape counts, not traffic or energy measurements.

Integrated native core only: temporal races (factorized law), sparse addressed writes into persistent rotating memories, transport between layers; no dense carrier and no count statistics. One pass over text8[0:10M] in 64 lanes x 128 characters (state reset per segment, exact credit within it), lr .004 with cosine annealing, compiled layer steps (contract-tested against the batched path). The v1 row used 128 lanes, about 610 updates and a constant lr; it is kept as measured. Test text8[95M:96M] on E64 windows, scored at the training length and at the controls' 256 with the same weights. Single

seed per row; exploratory, not a benchmark claim.

The E64 rows are the matched one-pass controls (1,220 steps of 32 x 256, cosine). Work: ours traced unit/special operations (fitting extrapolated from traced windows). Inference is traced twice: the batched emulator, which computes every proposal, and the exact winner-only evaluator (THEORY §414), whose small random float64 fixtures match emulator logits within 1e-10. Direct winner/state/cache contracts on the actual trained float32 weights are prepared and pending (note 143); the saved test scores use the compiled training evaluator. For fixed weights, cached stored-memory key reads need refreshing only on a slot's write: arithmetic scales as U.P plus winner maps per race. The present implementation also stacks every unit's matrices at each call, an $O(U.P^2)$ parameter-copy/allocation cost outside these FLOP counts. Cache storage, traffic and wall time must be measured before claiming total-resource scaling. Every key is scored and counted. Control columns repeat their single shape estimate. The conventions differ, so work comparisons are estimates.

Reading: update calibration took p16/d8 from 2.899 to 2.719. Width beat depth (p32/d4 2.507), and depth then helped at width 64 (p32/d8 2.456). Without route credit the fast path trains the race address only through first-time clock credit (THEORY §413), and more units then cost quality: pool 4 is worse than pool 2 (2.498 vs 2.456), and the no-selection pool-1 control beats pool 2 at depth 4 (2.439 vs 2.507). With the linearized local-expectation route credit (forward values unchanged, about 0.3% more counted fitting work) the same p32/d4 pool-2 model scores 2.370: .137 better than without it, .069 better than the control, and .057 better than the one-pass Transformer; at the matched T256 window it scores 2.371 versus 2.427, a .0554 bpc advantage, with about 1/15 of its parameters and estimated fitting work. It remains .199 behind the one-pass LSTM. With credit, pool 4 at the same 8 selected writes per character scores 2.343 (T256 2.345): more stored units now improve quality instead of costing it. Width is the strongest lever: p64/d4 with credit scores 2.184 (T256 2.183), .012 behind the one-pass LSTM, with exact winner-only inference of 0.60 MFLOPs per position against the LSTM estimate of 0.68 and more estimated fitting work (2.68 vs 2.03 MFLOPs per character). A write-address credit on stored coordinates diverged; the corrected variant trained stably at pool 2 without improving on value credit (2.384 vs 2.370) and diverged at pool 4, so write-address credit is withdrawn. Single seeds; pending arms are not filled.

Appendix. Tokenized256K: bounded learning-rate comparison

P24seed6, same data/source/initialization, 524272 fitting targets over two passes, 1024 updates, four-checkpoint cadence and 2040 development targets. Only constant learning rate changes; all core mechanisms retained. Public validation untouched. This tuning comparison is separate from the fixed-recipe scaling packet.


LR	Selected NLL	Selected step	Final NLL	Final – selected	Fit targets
0.003	7.741714	512	7.867910	0.126195	524272
0.001	7.740815	512	7.801747	0.060933	524272

Lower rate improves selected NLL 0.000900 and final NLL 0.066162. Both select step 512, the first-pass boundary; late deterioration falls 0.126195→0.060933. The bounded comparison reduces late deterioration without materially changing selected quality in this seed. Retain the fixed 0.003 recipe for the crossed scaling comparisons; no expanded rate grid or second-pass cure is inferred.

Both fits charge all 524272 fitting targets. Actual complete 256K work is pending for each trajectory; no cost is borrowed from 64K or from the other rate. Ordinary measured throughput 384.66/495.85 targets/s reflects separate host executions and is not a causal speed effect of learning rate.

Appendix. Measured language quality, capacity and compute

Sleeping Machines pursues a general-purpose substrate spanning language/reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. Completed integrated language measurements show quality improving with data and a useful intermediate capacity choice.



token language measured scaling 20261006 v4

P16 selected NLL improves 8.297491→8.099440 from 8K to 64K TRAIN tokens; P24 improves 8.033311→7.741714→7.251503 from 64K to 256K to 1M. The 256K to 1M gain is 0.490211 NLL, seed6 at fixed recipe. At 64K/P24 beats P16/P32 by 0.066129/0.057108 NLL, seed6. The independent P24 seed7 point is 8.078991. At 1M/P24 beats P32 by 0.012258 NLL; both select step 2,048. P24 seed7 at 1M selects 7.263557 at final step 4,096; two-seed mean 7.257530. Each plotted point is a completed fit.

Same 2040-target DEV population and two-pass fits, initial-inclusive selection; TRAIN frequency prior and checkpoint cadence depend on budget. Seed6 primary; P24 seed7 repeats at 64K and 1M. Width changes core/input/readout together. No raw points connected or extrapolated. Audited cost points differ in data and width; no iso-quality/iso-FLOP or exponent claim. Arithmetic fit boundary includes discovery/replay/backward/optimizer; special functions separate, random-sampling work unquantified. 256K and 1M work absent until actual audits. 4M remains reserved.

Appendix. Tokenized language: 1M capacity/data comparisons

Sleeping Machines pursues a general-purpose substrate for language and reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. This stage measures integrated tokenized learning with more data and capacity.

Seed6, development only. GPT-2 FineWeb: 1,048,576 admitted training tokens, two passes/2,097,136 fitting targets, 4096 updates, 2,040 scored development targets. Payload varies; depth2/heads2/pool4, batch64, credit16 and uniform-site K4 actual alternative-write credit are fixed. Every temporal, sparse and persistent-state mechanism is retained. Initialization is eligible for selection, with four evaluation checkpoints over the two passes. Public validation untouched.

Payload	Initial NLL	Selected NLL	Gain	Selected step
24	8.022527	7.251503	0.771024	2048
32	8.022527	7.263761	0.758766	2048

Payload	All parameters	Core + input parameters	Readout parameters	Targets/s	RSS KiB
24	3164842	2462880	701962	479.88	508712

Payload	All parameters	Core + input parameters	Readout parameters	Targets/s	RSS KiB
32	4225420	3305328	920092	448.06	560592

Member	Fit targets	Whole-fit GFLOPs	Fit MFLOPs/target	Eval MFLOPs/target
24	2097136	pending	pending	pending
32	2097136	pending	pending	pending
TF head $\geq$	695992320	≥ 161331598.841610	≥ 231.800832	≥ 77.266944

TF head $\geq$ is a source-derived arithmetic lower bound: forward output projection plus its two explicit backward matrix contractions; evaluation has one projection. Width768 and padded50,304classes give231.800832MFLOPs/fitting target and77.266944MFLOPs/evaluation target, at two FLOPs per multiply-add. Other reference computation and optimizer work are excluded. Its695,992,320training targets and public NLL3.2774 use a different data/quality population from these native development fits; these columns show raw work, not a matched-quality or iso-FLOP win. FP8 GPU arithmetic and native CPUfloat32 have different hardware costs.

Work cells require a complete replay with trajectory parity and full arithmetic formula coverage. Fitting includes discovery, actual alternative-write replay, readout, backward, clipping, optimizer and in-step diagnostics; preprocessing, evaluation and serialization are separate. Special functions are counted separately, random sampling work remains unquantified. Pending cells contain no extrapolation from a smaller fit.

All parameters include the token interface and readout; the core/input column includes lexical input parameters. Selected activity remains four writes and sixteen scored keys per token, while vector width grows. Equal data and passes are not equal fitting FLOPs. These ordinary throughput measurements exclude instrumented arithmetic tracing.

The 8K, 64K, 256K and 1M stages score the same development population and use two passes. Train-frequency priors and evaluation cadence differ; report absolute loss and within-fit learning separately. The learning gate requires a 0.02 NLL improvement over initialization. No scaling exponent or matched-compute Transformer win is inferred from these cells.

Payload	Context gain	Memory erase delta	Message erase delta	Both erase delta
24	0.942828	0.010388	0.436172	0.398140
32	0.944286	0.013540	0.440229	0.395295

Context gain is constant TRAIN-mean feature NLL minus intact NLL through the same frozen readout. Erasure deltas are intervention NLL minus intact NLL: positive means erasure hurts prediction, negative means it helps. Source, matched-RNG and partition checks pass. The constant-feature control is not an optimally refitted unigram; full-message erasure removes payload, arrival metadata and presence together, changing the normalization branch and read-clock policy. Frozen erasures are not retrained architecture comparisons. The completed payload-only diagnostic is shown separately.

Appendix. Tokenized language: 256K capacity/data comparisons

Sleeping Machines pursues a general-purpose substrate for language and reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. This stage measures integrated tokenized learning with more data and capacity.

Seed6, development only. GPT-2 FineWeb: 262,144 admitted training tokens, two passes/524,272 fitting targets, 1024 updates, 2,040 scored development targets. Payload varies; depth2/heads2/pool4, batch64, credit16 and uniform-site K4 actual alternative-write credit are fixed. Every temporal, sparse and persistent-state mechanism is retained. Initialization is eligible for selection, with four evaluation checkpoints over the two passes. Public validation untouched.

Payload	Initial NLL	Selected NLL	Gain	Selected step	
16	8.078820	7.744329	0.334492	512	
24	8.078820	7.741714	0.337106	512	
32	8.078820	7.753460	0.325360	512	

Payload	All parameters	Core + input parameters	Readout parameters	Targets/s	RSS KiB
16	2115464	1631184	484280	551.76	455084
24	3164842	2462880	701962	384.66	500896
32	4225420	3305328	920092	448.35	551536

Member	Fit targets	Whole-fit GFLOPs	Fit MFLOPs/target	Eval MFLOPs/target
16	524272	pending	pending	pending
24	524272	pending	pending	pending
32	524272	pending	pending	pending
TF head $\geq$	695992320	≥ 161331598.841610	≥ 231.800832	≥ 77.266944

TF head $\geq$ is a source-derived arithmetic lower bound: forward output projection plus its two explicit backward matrix contractions; evaluation has one projection. Width768 and padded50,304classes give231.800832MFLOPs/fitting target and77.266944MFLOPs/evaluation target, at two FLOPs per multiply-add. Other reference computation and optimizer work are excluded. Its695,992,320training targets and public NLL3.2774 use a different data/quality population from these native development fits; these columns show raw work, not a matched-quality or iso-FLOP win. FP8 GPU arithmetic and native CPUfloat32 have different hardware costs.

Work cells require a complete replay with trajectory parity and full arithmetic formula coverage. Fitting includes discovery, actual alternative-write replay, readout, backward, clipping, optimizer and in-step diagnostics; preprocessing, evaluation and serialization are separate. Special functions are counted separately, random sampling work remains unquantified. Pending cells contain no extrapolation from a smaller fit.

All parameters include the token interface and readout; the core/input column includes lexical input parameters. Selected activity remains four writes and sixteen scored keys per token, while vector width grows. Equal data and passes are not equal fitting FLOPs. These ordinary throughput measurements exclude instrumented arithmetic tracing.

The 8K, 64K, 256K and 1M stages score the same development population and use two passes. Train-frequency priors and evaluation cadence differ; report absolute loss and within-fit learning separately. The learning gate requires a 0.02 NLL improvement over initialization. No scaling exponent or matched-compute Transformer win is inferred from these cells.

Payload	Context gain	Memory erase delta	Message erase delta	Both erase delta
24	0.524467	0.019070	0.078129	0.064475
32	0.506384	0.022304	0.077059	0.049682

Context gain is constant TRAIN-mean feature NLL minus intact NLL through the same frozen readout. Erasure deltas are intervention NLL minus intact NLL: positive means erasure hurts prediction, negative means it helps. Source, matched-RNG and partition checks pass. The constant-feature control is not an optimally refitted unigram; full-message erasure removes payload, arrival metadata and presence together, changing the normalization branch and read-clock policy. Frozen erasures are not retrained architecture comparisons. The completed payload-only diagnostic is shown separately. Selected-checkpoint utility pending for payload 16; no utility value is predicted.

Appendix. Tokenized language: 64K capacity/data comparisons

Sleeping Machines pursues a general-purpose substrate for language and reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. This stage measures integrated tokenized learning with more data and capacity.

Seed6, development only. GPT-2 FineWeb: 65,536 admitted training tokens, two passes/131,056 fitting targets, 256 updates, 2,040 scored development targets. Payload varies; depth2/heads2/pool4, batch64, credit16 and uniform-site K4 actual alternative-write credit are fixed. Every temporal, sparse and persistent-state mechanism is retained. Initialization is eligible for selection, with four evaluation checkpoints over the two passes. Public validation untouched.

Payload	Initial NLL	Selected NLL	Gain	Selected step	
16	8.162280	8.099440	0.062840	128	
24	8.162280	8.033311	0.128969	128	
32	8.162280	8.090419	0.071861	64	

Payload	All parameters	Core + input parameters	Readout parameters	Targets/s	RSS KiB
16	2115464	1631184	484280	409.80	448368
24	3164842	2462880	701962	374.87	496232
32	4225420	3305328	920092	343.93	546508

Member	Fit targets	Whole-fit GFLOPs	Fit MFLOPs/target	Eval MFLOPs/target
16	131056	pending	pending	pending
24	131056	192.211765	1.466638	0.403428
32	131056	pending	pending	pending
TF head $\geq$	695992320	≥ 161331598.841610	≥ 231.800832	≥ 77.266944

TF head $\geq$ is a source-derived arithmetic lower bound: forward output projection plus its two explicit backward matrix contractions; evaluation has one projection. Width768 and padded50,304classes give231.800832MFLOPs/fitting target and77.266944MFLOPs/evaluation target, at two FLOPs per multiply-add. Other reference computation and optimizer work are excluded. Its695,992,320training targets and public NLL3.2774 use a different data/quality population from these native development fits; these columns show raw work, not a matched-quality or iso-FLOP win. FP8 GPU arithmetic and native CPUfloat32 have different hardware costs.

Work cells require a complete replay with trajectory parity and full arithmetic formula coverage. Fitting includes discovery, actual alternative-write replay, readout, backward, clipping, optimizer and in-step diagnostics; preprocessing, evaluation and serialization are separate. Special functions are counted separately, random sampling work remains unquantified. Pending cells contain no extrapolation from a smaller fit.

All parameters include the token interface and readout; the core/input column includes lexical input parameters. Selected activity remains four writes and sixteen scored keys per token, while vector width grows. Equal data and passes are not equal fitting FLOPs. These ordinary throughput measurements exclude instrumented arithmetic tracing.

The 8K, 64K, 256K and 1M stages score the same development population and use two passes. Train-frequency priors and evaluation cadence differ; report absolute loss and within-fit learning separately. The learning gate requires a 0.02 NLL improvement over initialization. No scaling exponent or matched-compute Transformer win is inferred from these cells.

Payload	Context gain	Memory erase delta	Message erase delta	Both erase delta
16	0.208993	0.006892	-0.024694	-0.028288
24	0.252343	0.007313	0.022623	0.014234
32	0.166731	0.002895	-0.009512	-0.011456

Context gain is constant TRAIN-mean feature NLL minus intact NLL through the same frozen readout. Erasure deltas are intervention NLL minus intact NLL: positive means erasure hurts prediction, negative means it helps. Source, matched-RNG and partition checks pass. The constant-feature control is not an optimally refitted unigram; full-message erasure removes payload, arrival metadata and presence together, changing the normalization branch and read-clock policy. Frozen erasures are not retrained architecture comparisons. The completed payload-only diagnostic is shown separately.

Appendix. Independent seed for selected 64K capacity

Payload24, selected from the three completed seed6 capacity fits. Seeds6/7 share GPT-2 FineWeb, 65,536 admitted training tokens, 131,056 fitting targets/two passes and 2,040 development targets; initialization-inclusive selection every64 updates. Public validation untouched.

Seed	Initial NLL	Selected NLL	Gain	Selected step
6	8.162280	8.033311	0.128969	128
7	8.162280	8.078991	0.083288	128

This repeat measures selected-member learning reliability and seed variation. The quality gain over P16 is a seed6 capacity comparison; it is not a paired two-seed capacity win.

Context gain	Memory erase delta	Message erase delta	Both erase delta
0.234159	0.005551	-0.032675	-0.039134

Frozen interventions use the same readout and matched RNG, with intact partition parity. Positive erasure delta means the intervention hurts prediction. Full-message erasure bundles payload, arrival metadata and presence, including changes to normalization and the read clock. Constant TRAIN-mean features are not an optimally refitted unigram; no retrained ablation claim.

Appendix. Tokenized messages: replicated payload contribution

P24 selected64Kcheckpoints, two seeds, 2,040development targets. Erasure delta is intervention NLL minus intact NLL. Payload-only erasure hurts prediction in both seeds by0.105206/0.101270NLL; the message information path contributes under this intervention.

Seed	Intact NLL	Payload erase delta	Full message erase delta
6	8.033311	0.105206	0.022621
7	8.078991	0.101270	-0.032673

Frozen selected checkpoints, development only. Payload intervention zeros ctx_vals while preserving ctx_arr and has_ctx before each event; full-message intervention zeros all three. Payload intervention preserves the normalization branch and incoming arrival metadata at that event, but future routes/timing/states can change. Matched route RNG and intact chunk-partition/selected-score parity. Neither intervention is a retrained architecture or a causal decomposition of total gains.

Appendix. Tokenized8K: replicated learning and memory-use decisions

Sleeping Machines pursues a general-purpose substrate for language and reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. This integrated stage advances the Transformer-leading language program; tiny fits are engineering diagnostics.

Both horizons beat initialization in both seeds. GPT-2 FineWeb:8192admitted training tokens,16368fitting target presentations/two passes,32AdamWupdates,2040development targets. P16/D2/H2/U4,batch64; credit16/64 share data/updates/cadence. Initial-inclusive development selection every8updates; public validation untouched. Four selected writes,16scored keys and256persistent memory scalars per lane. The bounded local fits preserve the immutable AWSqueues.

Seed	Credit	Initial NLL	Selected NLL	Gain	Context gain	Memory gain
6	16	8.340313	8.297491	0.042822	0.017759	0.000381
6	64	8.340313	8.300131	0.040182	0.001898	-0.000268
7	16	8.340313	8.286427	0.053886	0.027515	0.000980
7	64	8.340313	8.290236	0.050078	0.017711	0.000449

Credit16mean selected8.291959 beats credit64mean8.295183, and is better in both paired seeds; the0.003224mean gap is smaller than seed spread. Ordinary measured fitting throughput is395–422targets/s, peakRSS452456–453252KiB. Retain both completed records. Numeric state carries across credit boundaries; truncating derivatives does not reset state.

Context gain uses the same frozen readout with a constant mean of8184causal TRAINfeatures. Memory gain is NLL after per-token addressed-state/arrival/seen erasure minus intact NLL; negative means the intervention improves loss. Recurrent-message erasure costs0.017843–0.021955NLL. Matched route RNG, source binding and token partition parity pass. These are frozen interventions, not retrained ablations; the constant-feature control is not an optimally refitted unigram.

Alternative	Seed	Selected NLL	Gain vs credit16	Parameters
Full-width tails	6	8.302969	-0.005478	3292836
Memory gain2	6	8.310055	-0.012563	2115464
Full-width tails	7	8.286345	0.000082	3292836

Full-width decoder tails lose at seed6 and are effectively tied at seed7. Retain the narrower decoder. First-tail training exposure is695targets/data pass, so this comparison exercises learned tail projections. The memory-gain intervention doubles the existing unit output contribution while preserving all temporal/state/routing/learning mechanisms and parameter count; it is selected only from completed fits. See TOKEN_MEMORY_COUPLING_20261005.md for derivation and contracts.

Member	Fit targets	Whole fit GFLOPs	Fit MFLOPs/target	Eval MFLOPs/target
Credit16seed6	16368	13.534669	0.826898	0.272633
Credit64seed6	16368	pending	pending	pending
TF head $\geq$	695992320	≥ 161331598.841610	≥ 231.800832	≥ 77.266944

TF head $\geq$ is a source-derived arithmetic lower bound: forward output projection plus its two explicit backward matrix contractions; evaluation has one projection. Width768 and padded50,304classes give231.800832MFLOPs/fitting target and77.266944MFLOPs/evaluation target, at two FLOPs per multiply-add. Other reference computation and optimizer work are excluded. Its695,992,320training targets and public NLL3.2774 use a different data/quality population from these native development fits; these columns show raw work, not a matched-quality or iso-FLOP win. FP8 GPU arithmetic and

native CPUfloat32 have different hardware costs.

Actual whole-fit traces must reproduce each control trajectory. Arithmetic includes factual computation, all-key discovery, suffix replay, exact target readout, backward, clipping, optimizer and in-step diagnostics; initialization/frequency counting, evaluation and serialization excluded. Special functions separate; random-sampling work unquantified. Evaluation includes scorer reductions. Pending cells contain no predicted work or quality. Equal updates/data here are not an iso-FLOP claim. Selected Transformer-reference quality and resource protocol remains the larger benchmark target.

Appendix. Language strengths: measured decisions and complete work

The target data region deliberately favors strong Transformer references over LSTMs; tiny native fits diagnose learning before that larger comparison. The native construction chooses temporal computation, persistent memory lifetime and sparse addressed activity independently of Transformer topology or fixed windows. Matched-history and longer-native-history comparisons are labelled separately; causal inputs and scored targets stay explicit. Capacity beyond activity earns promotion through better held-out prediction with discovery and learning fully charged.

Controlled learning-rate test: seed6, GPT-2 FineWeb2K, P16/D2/H2/U4, batch64/credit64, uniform K4 actual alternative-write utility, 16 updates and 1,016 development targets. Initial weights eligible. Lower learning rate loses selected quality by 0.046604 NLL; retain .003. Its better final loss does not replace the better selected incumbent. These are development comparisons.

Recipe	Selected NLL	Final NLL	Fit targets	Targets/s	RSS KiB
Retained lr .003	8.819470	9.111889	8160	404.6	451728
Alternative lr .001	8.866074	8.866074	8160	423.2	451308

Decoder exposure uses train-only frequency rank and eight contiguous lanes. At 2K every target reaches the head: identical default/full-width development curves do not test learned tail rank. The existing 8K comparison reaches the first tail. Branch counts below are targets per data pass; head cutoff 2,000, tail boundaries 10,000/30,000/50,257.

Train tokens	Distinct	Head	Tail1	Tail2	Tail3
2048	796	2040	0	0	0
8192	2696	7489	695	0	0
32768	6778	26915	5845	0	0
65536	10914	51272	13343	913	0
262144	23144	196953	47809	17374	0

Actual native work: multiply-add=2 FLOPs; arithmetic excludes separately counted special functions and unquantified random sampling. Whole fitting includes factual core, all-key scoring, alternative discovery, suffix replay, exact readout, backward, clipping, optimizer and in-step diagnostics. Initialization/frequency counts, evaluation and serialization are outside that fitting column. Selected evaluation includes scorer reductions, with no gradient/optimizer. Audit instrumentation wall time is not ordinary throughput. Public Transformer work remains in its separate audited protocol; no matched-compute win is inferred here.

Boundary	Targets	Total GFLOPs	MFLOPs/target	Coverage
Whole fit	8160	6.183337	0.757762	complete
Selected evaluation	1016	0.296939	0.292263	complete

Next: existing AWS 8K paired seeds/horizons, trained memory utility, informative decoder comparison, then measured width/data and fixed-activity pool scaling. Public validation remains reserved. See

experiments/TOKEN_STRENGTH_EXECUTION_20261005.md for the ordered strength tests.

Appendix. Tokenized language: confirmed small-fit learning

Sleeping Machines pursues a general-purpose substrate for language and reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. These diagnostics select constructions within that family.

Both integrated members beat initialization in both seeds. GPT-2 FineWeb:2,048 admitted training tokens,8,160 training targets over four passes and16updates,1,016 scored development targets. Optimizer batch64 per lane; credit16/64 share data, updates and evaluation cadence. Natural-log NLL, lower is better; initialization is eligible for selection. Public validation is reserved.

Seed	Credit	Initial dev	Selected dev	Gain	Context gain
6	16	8.906910	8.812527	0.094383	0.005029
6	64	8.906910	8.819470	0.087440	0.015265
7	16	8.906910	8.813838	0.093072	0.011278
7	64	8.906910	8.812853	0.094057	0.011398

Context gain compares actual features with constant mean causal TRAIN features through the same frozen learned readout. All four contributions are positive. This intervention does not isolate persistent memory from current-token features or refit an optimal unigram control. Much of the total gain is marginal recalibration.

Frozen history interventions: erasing the recurrent event message worsens NLL0.04618-0.06734 across all four checkpoints; erasing addressed memory changes it by-0.001080 to+0.001388. Matched RNG and intact partition parity pass. Messages can carry multi-event history. These out-of-distribution interventions target addressed-memory utilization for diagnosis; they are not retrained ablations.

Credit16 wins selected quality at seed6; credit64 wins at seed7 and has better final loss in both seeds. Retain both for the matched AWS8K comparison. Changing optimizer batch16 to64 also reduced updates64 to16 at equal targets: a recipe improvement, not an isolated horizon effect.

P16/D2/H2/U4 supplies256 persistent memory scalars per lane, four selected writes and16 scored keys per target. Mechanisms include race clocks, addressed persistent state, separate keys/values, small messages, depth and actual counterfactual write credit. Decoder rank and credit boundaries are selectable design choices. These are confirmed development learning gains; the public benchmark comparison is the next larger test.

Appendix. Tokenized language: integrated learning development

Sleeping Machines pursues a general-purpose substrate for language and reasoning, multimodal world models, embodiment, event-native analytics, continual learning, communication, self-design and hardware. These language diagnostics test constructions within that family.

Historical optimizer-batch16 recipe, seed6; GPT-2 FineWeb; 2,048 admitted training tokens, 8,160 presented targets over64updates, 1,024 development tokens yielding1,016 scored targets after8lane starts. Earlier notes called these1,024targets; saved losses are unchanged.

P16/D2/H2/U4:256persistent memory scalars per lane, four selected writes and16scored keys per target. Public validation remains reserved. Loss is natural-log NLL; lower is better. Throughput includes training work, excludes evaluation.

Member	Initial dev	Best trained dev	Final dev	Targets/s
Local value	8.906910	9.131615	9.888922	404.09

Member	Initial dev	Best trained dev	Final dev	Targets/s
First-site full suffix	8.906910	9.130419	9.865284	342.83
First-site K4 suffix	8.906910	9.130868	9.844923	351.84
Uniform-site K4 suffix	8.906910	9.063122	10.354817	386.82

Uniform-site K4 wins the trained development comparison against first-site K4 by 0.067745NLL, single seed at equal data/presentations. It combines broader site exposure, utility weighting and removal of immediate-only route credit.

All trained members in this historical batch lose to initialization. Uniform-site late dev 10.354817 loses to first-site K4 late 9.844873. Initial weights now remain eligible for selection; historical trained-only fields stay preserved with the correction beside them.

Mechanisms: persistent state, race clocks/transport, sparse hard writes, separate selection/value roles, small messages, depth and actual counterfactual writes. Credit remains chunk-bounded; all keys are scored. The uniform-site learner uses winner-only factual/replay proposals. Memory occupancy alone is not predictive capacity.

FAS native frozen-recipe replication

Replicated win against all six anonymous generic controls at 256 events. Completed native fitting seeds: 6, 7, 8.

Native seed	AUROC N256	Gap vs controls	AUROC N512	Gap vs controls
6	0.599794	+0.041126	0.742185	+0.014952
7	0.591049	+0.032381	0.735964	+0.008730
8	0.586373	+0.027705	0.732869	+0.005636

Best generic controls: 0.558669 AUROC at N=256 and 0.727233 at N=512. The fixed P32/D4/H2/U2 recipe uses two passes over 10,000 clean FIT runs; selection uses validation-clean NLL only. Test: 2,000 clean and 2,000 faulty runs. Seeds 7/8 score artifacts each reproduce all 18 aggregate/type/prefix AUROCs exactly and retain the selected weights.

FAS v1 replication, not sealed v2 confirmation. Generic-control comparison uses saved references; stronger neural benchmarks are pending. FIFO/timing de-interleaving trained on hidden item identities remain oracle-assisted diagnostics, excluded here. Original seed 6 has aggregate scores, not the new per-run artifact; paired multi-seed bootstrap cannot be reconstructed from those aggregates. Current frozen dependency hashes are checked; historical unrecorded dependency identity is not independently verified.

Appendix. Public benchmarks: NeuroBench and SHD (status from completed result files)

Targets and rules: experiments/SOTA_TARGETS.md. Official protocols are fixed before they run and use the vendored official neurobench 2.3.0 loaders/metrics (slices proven identical by contract tests). Development results select settings on other series (Mackey-Glass tau 18/19) or on validation splits (primate).

NeuroBench chaotic function prediction (Mackey-Glass tau 17, sMAPE, lower is better)

Model	sMAPE	Footprint (bytes)	Inference ops / step
Ours: race member, mix8 (primary, complete)	14.84	57,572 (float32)	93,264 (traced)
Ours: same weights, argmax (reporting only)	18.43	57,572	11,658 (traced)
Ours: same weights, sampled (reporting only)	17.65	57,572	11,658 (traced)
LSTM (NeuroBench)	13.37	490,000	60,300 (effective MACs)
ESN (NeuroBench)	14.79	281,000	4,370 (effective MACs)

Official protocol: 30 start offsets, 750 teacher-forced training points, 750 autonomous predictions; per-repeat mix8 range 6.2–25.2. Our model has 14,393 parameters; the primary inference mode (averaging 8 race-noise streams) was fixed on tau 18 development before tau 17 was scored. Our operation counts come from the operator tracer (multiply-adds); NeuroBench counts hooked layer MACs; the conventions differ.

Mackey-Glass development (tau 18/19, three official-protocol repeats per arm)

Arm	Member	Mean SMAPE by inference mode	Repeats
tau 18, p16/d2/pool2, closed-loop	race	sampled 24.5	3
tau 18, p16/d2/pool2	race	sampled 22.5	3
tau 18, p32/d2/pool2, closed-loop	race	sampled 26.3	3
tau 18, p16/d2/pool2, increments	race	sampled 25.3	3
tau 18, p16/d2/pool2, increments	race	argmax 18.5	3
tau 18, p16/d2/pool1, increments	race	sampled 16.4	3
tau 18, p32/d2/pool2, increments, 3000 steps	race	argmax 19.1	3
tau 18, p16/d2/pool2, increments, det. training	race	argmax 18.4	3
tau 18, p16/d2/pool1, increments, 3000 steps	race	sampled 16.8	3
tau 18, p32/d2/pool2, increments, det. training, 3000 steps	race	argmax 20.9	3
tau 18, p16/d2/pool2, increments	race	mix8 17.2, argmax 18.5, sampled 25.3	3
tau 18, p16/d2/pool2, increments, 3000 steps	race	mix8 21.0, argmax 18.4, sampled 15.1	3
tau 18, p16/d2/pool4, increments	race	mix8 23.0, argmax 18.3, sampled 19.8	3
tau 18, p16/d2/pool2, 1000 steps	race	sampled 23.0	3
tau 18, p16/d2/pool2, increments	expected reception	sampled 23.1	3
tau 18, p16/d2/pool4, increments	expected reception	sampled 17.1	3
tau 18, p16/d2/pool2, increments	expected reception	sampled 21.4	3
tau 18, p32/d2/pool2, increments	expected reception	sampled 19.8	3
tau 19, p16/d2/pool2, increments	expected reception	sampled 23.5	3
tau 19, p16/d2/pool4, increments	expected reception	sampled 30.4	3

Findings (THEORY §§417–418): sampled races inject output noise; deterministic or averaged inference helps; closed-loop training did not; the exact-expected-reception member (deterministic delivery, hard writes) is under development. Three repeats cannot rank close arms (per-repeat spread is large).

NeuroBench primate reaching (development session indy_20170131_02, R², higher is better)

Arm	Test R ²	Validation R ²	Params
p32/d2/pool4 tied, 4000 steps, leaky readout, 4 route samples	0.729	0.709	44,826
p32/d2/pool4 tied, wd 0.0001, 4000 steps, leaky readout, 4 route samples	0.754	0.715	44,826
p32/d2/pool8 tied, 4000 steps, leaky readout, 4 route samples	0.720	0.715	45,866
p32/d2/pool2, 4000 steps, leaky readout, 4 route samples	0.738	0.737	61,082
p32/d2/pool4 tied, wd 0.0001, 4000 steps, leaky readout, 4 route samples	0.747	0.744 (4 blocks)	44,826
p32/d2/pool4 tied, wd 0.0001, 4000 steps, leaky readout, 4 route samples, traces	0.710	0.749 (4 blocks)	57,114
p32/d2/pool2, 4000 steps, leaky readout, 4 route samples, traces	0.724	0.750 (4 blocks)	73,370
p32/d2/pool2, 2000 steps	0.725	0.710	61,082
p32/d2/pool4 tied, 2000 steps	0.741	0.711	44,826
p64/d2/pool2, 2000 steps	0.714	0.704	228,634
bigRSNN (published, this session)	0.772	—	—
tinyRSNN (published, this session)	0.746	—	—

Leaderboard (mean of six official sessions): the six-session protocol run with a configuration fixed on validation is prepared on AWS; no leaderboard cell is filled from one development session.

Leaderboard entry	Six-session R ²	Footprint (bytes)
AEGRU	0.710	45,500
GRU-t1	0.707	352,904
bigSNN (bigRSNN)	0.698	4,833,360
tinyRSNN	0.660	27,144

Validation and test rank the development arms differently; the protocol selects by validation only. Test on the development session was observed during development; the six-session report will also give the mean of the five untouched sessions.

SHD and other public benchmarks

SHD (official files, speaker-held-out validation) development runs are queued. The completed public time-series campaign (ECG200, JapaneseVowels, PenDigits) is negative and reported in its own appendix.

Appendix B. AWS coarse temporal screens: two seeds pass

Native packet/clock	Accuracy %	Dev NLL	Whole fit GF est.	Fit MF / target	Infer MF / prefix
s6 20 bins/.05	48.96	1.4242	2.285696	2.232125	0.591709
s6 4 bins/.05	57.81	1.2221	0.548517	0.535661	0.138119
s6 4 bins/.25	55.73	1.2749	0.548517	0.535661	0.138133
s7 20 bins/.05	49.48	1.4347	2.285696	2.232125	0.591695
s7 4 bins/.05	53.12	1.3233	0.548517	0.535661	0.137993
s7 4 bins/.25	55.21	1.2208	0.548517	0.535661	0.137965

Same integrated p16/L2/H2/pool2 core, 256 fitting gestures, 192 subject-disjoint development gestures, four passes, 1,024 fitting presentations and U16. Coalescing five 50ms count packets into

each 250ms packet preserves total causal counts and the 1s query but removes timing within each quarter second. Learned temporal computation, races, separate keys/values, sparse addressed writes and counterfactual credit remain. The two coarse arms differ only in clock initialization.

Both coarse arms pass the prespecified screen in seeds6 and7: at least .02 NLL improvement, at most 1 percentage point accuracy decline and at most .50 fitting work ratio. Work ratio .239978 means 76.002% less counted fitting work. Events 21,504 to 5,120; key scores 172,032 to 40,960; writes 86,016 to 20,480. Eight available receivers and four selected writes per event in every arm.

Seed6 one-thread AWS job walls are 42.846/14.031/14.046s, under the authorized three-slot scheduler. These are observations, not universal hardware speedups. The subsequent full-data matrix is complete on the next page; the earlier pending replication and full-data statuses are superseded by completed evidence.

Exploratory reused development set; no independent confirmation or control advantage. All passes, candidate values, backward/Adam/clip are charged in consistent arithmetic plus unit-special estimates. Whole-job wall includes preprocessing/evaluation; NumPy preprocessing FLOPs, traffic and energy remain unknown. Both completed analysis files and the verbatim pre-integration AWS appendix are retained in version control.

Appendix B. Full coarse matrix: positive means, failed all-seed gates

Native packet/clock	Accuracy %	Dev NLL	Whole fit GF est.	Fit MF / target	Infer MF / prefix
s6 20 bins/.05	57.81	1.1798	17.573520	2.232409	0.591779
s6 4 bins/.05	61.46	1.0718	4.218015	0.535825	0.138035
s6 4 bins/.25	66.67	0.8910	4.218015	0.535825	0.138119
s7 20 bins/.05	55.73	1.1726	17.573520	2.232409	0.591737
s7 4 bins/.05	70.31	0.9333	4.218015	0.535825	0.137923
s7 4 bins/.25	61.98	1.0569	4.218015	0.535825	0.138007
s8 20 bins/.05	68.75	0.8986	17.573520	2.232409	0.591751
s8 4 bins/.05	60.42	1.0286	4.218015	0.535825	0.138007
s8 4 bins/.25	69.27	0.9207	4.218015	0.535825	0.138035

Nine completed fits: same984 fitting/192 development gestures, eight passes, 7,872 fitting presentations per run. Mean fine accuracy/NLL 60.764%/1.083660; 4-bin/.05 64.062%/1.011241; 4-bin/.25 65.972%/1.956220. Coarse fitting work is 4.218015 versus 17.573520 GFLOPs per fit: ratio .240021. Every per-target column uses the same denominator across models; no whole-fit/per-target unit mixing.

Both frozen all-seed gates FAIL. Seed8 fine 68.75%/.898591 has lower NLL than 4-bin/.05 60.417%/1.028562 and 4-bin/.25 69.271%/.920744. Diagnostic crossed seed/development-user intervals include zero. Preserve positive mean quality and work savings alongside this failure; no unchanged extension follows.

Original stronger local references remain valid: native65.10%/.963161 and clock66.15%/1.041987. They have different fitting implementations and work totals and are not replaced by the weaker matched fine arm. Same8 receivers, four writes per event, 21 versus5 events per query. Saved full-coarse analysis retains every curve, activity ledger, result digest, wall/RSS and uncertainty scope. No official test.

Appendix B. Strong coarse controls and frozen readout evidence

Native packet/clock	Accuracy %	Dev NLL	Whole fit GF est.	Fit MF / target	Infer MF / prefix
Raw RBF 1 bins	68.75	0.8068	Unknown	Unknown	Unknown
Raw RBF 4 bins	77.60	0.6867	Unknown	Unknown	Unknown
Raw RBF 20 bins	74.48	0.7080	Unknown	Unknown	Unknown

Frozen encoder/access	Mean accuracy %	Mean dev NLL	Nomination gate
Initial/query32	61.632	1.036019	Diagnostic
Initial/all resident state	68.750	.886478	Diagnostic
Trained/query32	69.097	.856456	Pass
Trained/all resident state	72.743	.829107	FAIL

Raw controls use 984 fitting/192 development examples and fitting-user GroupKFold selection. The 4-bin RBF result 77.604%/.686661 is a strong completed reference under that protocol; the earlier local RBF 73.44%/.706478 uses different selection and remains visible. Solver FLOPs and total fit/inference work are unknown, not zero.

All three matched-clock native encoders and their initial reservoirs are frozen. Query RBF probes improve native mean 65.972%/.956220 by 3.125 points/.099764 NLL; individual NLL gains .056455/.163228/.079610 pass the context nomination gate. All resident reads add 128 memory values, eight ages and eight occupancy flags. Their extra mean .027349 NLL gain misses the .05 resident gate.

These are information/readout diagnostics. Dense resident access does not demonstrate sparse dormant-value retrieval. Conditional fitting-only decoder selection shares label-trained encoders; it is not producer-cross-fitted validation. The next integrated quadratic head fails, as recorded on the next page.

Original native fits 4.218015 GF each, six 1,176-prefix replays, 72 fold fits and 12 refits retained. Probe 53.605s/516,444 KiB; state/winner/probability contracts pass. Replay solver/materialization/traffic/energy totals remain unknown. Full per-seed readout outcomes and costs are preserved in AWS_COARSE_READOUT_FINDINGS_20261002.md.

Appendix B. Integrated quadratic head: contracts pass, quality fails

Native packet/clock	Accuracy %	Dev NLL	Whole fit GF est.	Fit MF / target	Infer MF / prefix
affine	55.73	1.2749	0.548517	0.535661	0.138133
quadratic	55.21	1.4415	0.594382	0.580452	0.201506

Same 4-bin/.25-clock native p16/L2/H2/pool2, 256 fitting/192 development gestures, seed 6, four passes and 1,024 fitting presentations. A standard local degree 2 query residual starts at zero: initial logits/RNG and every old gradient nest the affine model exactly. Six numerical/deep-gradient/recovery/accounting contracts and 24 fit/8 dev/two-pass learning smoke pass before this screen.

The quality gate FAILS: NLL worsens .166618 and accuracy declines .521 percentage point. Fitting ratio 1.083617 and inference ratio 1.458782 satisfy resource admission but do not rescue prediction quality. There is no unchanged seed 7/full-data/epoch extension and no replacement of the leading result.

The 5808 added head weights bring parameters from 15,523 to 21,331. Eight available receivers, candidate keys/values and four selected writes per event are unchanged. Every one of the five

sequential readout calls is charged. Physical time, hard races, persistent addressed state, separate keys/values and counterfactual learning remain; positive frozen RBF information alone does not prove this coupled encoder/head optimizer can exploit it.

Full-coarse seed8 failure and initial-reservoir/readout controls remain beside this negative result. The AWS team owns the separately frozen affine versus polynomial convex-fit diagnostic. No pending cell is reported as evidence; solver work, preprocessing FLOPs, traffic and energy remain separately scoped.

Appendix B. Receiving losing arrivals requires boundary credit

Constructed gradient	Complete conditional reference	Ordinary history-cell autograd
Width dL/dH	-27.662490	0.656618
Zero-width right derivative	-34.323670	Birth credit required

For a common-start exponential race, winner W and first time T factor as $W \sim \text{Categorical}(\lambda/\sum \lambda)$, $T = E_0/\sum \lambda$. Conditional losing residual arrivals are independent $\text{Exp}(\lambda_j)$. Winner-only replay omits those residuals when a layer receives several messages. Under the native increasing bounded delay map, a deadline H after the first arrival gives an explicit raw residual cutoff $R(T, H)$ and losing membership $q_j = 1 - \exp(-\lambda_j R)$.

The complete expected gradient decomposes into ordinary fixed-history derivatives, winner choice and paired boundary flux. At a losing arrival crossing the deadline, dq_j multiplies the downstream loss difference between heard and unheard outcomes, including the actual separate receiver write. This produces faster/slower emitter and reception-width credit while preserving coupled contents and physical times.

Seven numerical contracts verify joint arrival density, physical cutoff/cap, membership normalization/activity, every score/content/decay/width derivative, independent finite differences, zero-width birth and actual-write effects. Three candidates, decaying messages and a later addressed-memory read; quadrature agrees with the boundary/winner/history formula. Maximum width finite-difference error $4.41e-9$. The constructed sign reversal shows ordinary autograd can widen or narrow in the wrong direction; it does not identify a benchmark failure.

Enumerating membership histories grows exponentially. Paired boundary sampling can avoid enumeration, but every candidate discovery, actual timed counterfactual and future replay is paid and variance remains. Conditional inverse-CDF derivatives already redistribute heard arrival times; adding the full unconditional boundary term to them double counts credit. Neither primitive supplies a free whole-model gradient.

Theory97, guarded .743s/295,904KiB reference. No optimizer or fitted native window. Fixed-first deadline differs from silence-reset popcorn, whose scheduler and merge/split credit remain separate. FLOPs/traffic/energy unmeasured, not zero. The next pages test real native information/access before any training admission.

Appendix B. Native reception intervention, seed6: no smoke admission

H ms	One-site mode	Initial NLL	Trained NLL	Initial acc %	Trained acc %	Extra arrivals
0	Winner only	2.820885	1.071110	6.25	50.00	0/16
1	Heard sum + writes	2.817558	1.070983	6.25	50.00	0/16
1	Heard mean + writes	2.824706	1.070983	6.25	50.00	0/16
1	Winner x heard count	2.811214	1.070983	6.25	50.00	0/16
1	Winner + same wait	2.820885	1.070983	6.25	50.00	0/16
1	Heard sum, winner write	2.817760	1.070983	6.25	50.00	0/16
3	Heard sum + writes	2.799750	1.104131	6.25	50.00	4/16
3	Heard mean + writes	2.819896	1.069441	6.25	50.00	4/16
3	Winner x heard count	2.800960	1.105012	6.25	50.00	4/16
3	Winner + same wait	2.820911	1.070549	6.25	50.00	4/16
3	Heard sum, winner write	2.800910	1.104000	6.25	50.00	4/16
Mean/winner H ms	First-prefix core MF est.	Minimum ready time s		Maximum ready time s		
0	0.591695	1.004001		1.012672		
1	0.591874	1.005001		1.012672		
3	0.591874	1.007001		1.014160		

Frozen fixed-four-pass producer and its initial reservoir; 16 prespecified evenly spaced FIT examples unused by the256-label producer. No optimizer, decoder refit, development or official-test evaluation. One actual site: observed packet event19/layer0/head0; the following deeper computation/query sees real changed receiver memories. Transport runs to the actual deadline.

Prespecified1ms heard-mean versus same-wait winner gate FAILS in both trained seeds: zero extra arrivals and zero NLL difference. At3ms seed6 hears extras4/16, seed7 7/16; mean-over-wait NLL gains .001108/.002312. Seed6 sum worsens loss. Those diagnostic settings cannot replace the frozen gate. Failure concerns this site/width/sample, not all learned windows.

Theory98. Eight available receivers, four original writes/event and 168 scored keys/prefix; only this site can add one actual write. Original trained fit2.285696GF/2.232125MF per fitting presentation retained. Table inference MF is the first selected prefix only, not full-audit FLOPs. All44 model/configurations104.222s/396,720KiB; twelve zero-nesting/gradient/causality/write/target-invariance/refused-training contracts pass. Frozen weights remain bitwise equal. Unknown audit/scheduler/traffic/energy work is not zero. Processing readiness is charged beyond the1s observation cutoff.

Appendix B. Native reception intervention, seed7: no smoke admission

H ms	One-site mode	Initial NLL	Trained NLL	Initial acc %	Trained acc %	Extra arrivals
0	Winner only	2.461633	1.025841	12.50	56.25	0/16
1	Heard sum + writes	2.459791	1.025891	12.50	56.25	0/16
1	Heard mean + writes	2.460627	1.025891	12.50	56.25	0/16
1	Winner x heard count	2.461396	1.025891	12.50	56.25	0/16
1	Winner + same wait	2.461622	1.025891	12.50	56.25	0/16
1	Heard sum, winner write	2.459954	1.025891	12.50	56.25	0/16
3	Heard sum + writes	2.458680	1.022673	12.50	56.25	7/16
3	Heard mean + writes	2.455532	1.023701	12.50	56.25	7/16
3	Winner x heard count	2.454758	1.024533	12.50	56.25	7/16
3	Winner + same wait	2.452995	1.026013	12.50	56.25	7/16
3	Heard sum, winner write	2.458298	1.022674	12.50	56.25	7/16
Mean/winner H ms	First-prefix core MF est.	Minimum ready time s		Maximum ready time s		
0	0.591695	1.004017		1.015801		
1	0.591874	1.005011		1.016597		
3	0.591874	1.007011		1.019004		

Frozen fixed-four-pass producer and its initial reservoir; 16 prespecified evenly spaced FIT examples unused by the256-label producer. No optimizer, decoder refit, development or official-test evaluation. One actual site: observed packet event19/layer0/head0; the following deeper computation/query sees real changed receiver memories. Transport runs to the actual deadline.

Prespecified1ms heard-mean versus same-wait winner gate FAILS in both trained seeds: zero extra arrivals and zero NLL difference. At3ms seed6 hears extras4/16, seed7 7/16; mean-over-wait NLL gains .001108/.002312. Seed6 sum worsens loss. Those diagnostic settings cannot replace the frozen gate. Failure concerns this site/width/sample, not all learned windows.

Theory98. Eight available receivers, four original writes/event and 168 scored keys/prefix; only this site can add one actual write. Original trained fit2.285696GF/2.232125MF per fitting presentation retained. Table inference MF is the first selected prefix only, not full-audit FLOPs. All44 model/configurations104.222s/396,720KiB; twelve zero-nesting/gradient/causality/write/target-invariance/refused-training contracts pass. Frozen weights remain bitwise equal. Unknown audit/scheduler/traffic/energy work is not zero. Processing readiness is charged beyond the1s observation cutoff.

Appendix B. AWS frozen polynomial decoder: both nomination gates fail

Encoder	Seed	Degree	Accuracy %	Dev NLL	Native infer MF/query
Initial	6	1	57.81	1.162621	Unmeasured
Initial	6	2	64.06	1.065509	Unmeasured
Fitted	6	1	68.75	0.835175	0.138119
Fitted	6	2	68.75	0.870350	0.201534
Initial	7	1	56.77	1.154207	Unmeasured
Initial	7	2	59.90	0.994128	Unmeasured
Fitted	7	1	66.67	0.939479	0.138007
Fitted	7	2	69.79	0.899185	0.201422
Initial	8	1	60.94	1.097227	Unmeasured
Initial	8	2	64.58	0.976572	Unmeasured
Fitted	8	1	68.75	0.936226	0.138035
Fitted	8	2	71.35	0.885890	0.201450

All three coarse trained encoders and exact initial reservoirs,984 fitting/ 192 development examples. Local affine/quadratic heads selected by fitting-user GroupKFold over C(.01,.1,1);108 fold fits,12 final fits and six feature replays. No encoder retraining, dense prefix carrier or resident-memory read. Algebraic folding reproduces actual native predictions/state; every sequential head is charged.

Fitted affine mean68.056%/.903627; quadratic69.965%/.885142. Only .018485 mean NLL gain and1.910 accuracy points, with seed6 NLL regression .035176, fail the polynomial nomination. Gains over parent native .020633/.157749/ .034854 miss the each-seed .05 requirement. Initial affine58.507%/1.138019 and quadratic62.847%/1.012070 remain visible. Both gates FAIL; no unchanged scaling.

Original trained encoder fit4.218015GF each, .535825MF per7,872 fitting presentations; initial optimizer work zero. Combined fitting and solver FLOPs UNKNOWN, so parent fit alone is not total work. Whole study126.813s/608,044KiB includes every fit/replay/evaluation/export; one convergence warning retained. Tensor bytes exclude shared input normalization and metadata. Conditional decoder folds reuse label-trained encoders; no pipeline cross-fit or test. RBF context evidence remains stronger; raw coarse77.604%/.686661 and historical compact66.667%/.902951 controls preserved. AWS owns the next compact context head.

Appendix B. Race-scaled reception: available arrival support

Seed	Encoder	Entropy	Top p>=.99 %	1ms heard	3ms heard	c4 heard
6	initial	0.750	5.13	1.197	1.523	1.510
6	fixed_pass4	0.335	44.47	1.120	1.341	1.228
7	initial	0.726	7.02	1.202	1.513	1.493
7	fixed_pass4	0.399	34.59	1.144	1.394	1.272

A relative raw deadline $(1+c)T$ maps through the actual bounded temporal delay. Its heard set is unchanged under a common entering-score shift: all raw clocks scale together. Physical computational times still change. The exact expected receiver count is $1 + \sum_w \pi_w \sum_{j \neq w} c \cdot \pi_j / (1 + c \cdot \pi_j)$. The maximum additional physical waiting is $.010 \cdot (\sqrt{1+c} - 1) / (\sqrt{1+c} + 1)$:1.716ms at c1,3.820ms at c4. The local11ms delay bound remains.

This fixes speed-dependent membership, not utility.

43,008 actual native races across all 84 sites, 16 producer-unseen FIT inputs, eight noise histories and four encoders. The earlier event 19/layer 0/head 0 site is more concentrated than the whole model: top probability $\geq .99$ in 71.88%/57.03% of trained races. First events have the most alternative support. A previous one-history 1ms lack of reception is not proof that all alternatives are absent.

Trained relative c4 mean additional waiting 1.406/1.675ms. These are frozen arrival profiles, with no changed deliveries/writes, new loss, optimizer, DEV or test evaluation. No learned window or quality nomination follows from counts.

Theory 99: six numerical contracts and original logit/state checks pass. 59.553s/335,536KiB; exact candidate/clock arrays and checksums retained. Theory 100 independently verifies collector end-RNG preservation. Prior trained core fits 2.285696GF each remain paid. Whole audit FLOPs, traffic and energy unknown, not zero.

Appendix B. Memory-conditioned keys explain routing confidence

Seed	Encoder	Static gap	Memory gap	Static entropy	Full entropy
6	initial	0.070	1.350	0.999	0.744
6	fixed_pass4	0.110	5.322	0.996	0.333
7	initial	0.105	1.450	0.998	0.732
7	fixed_pass4	0.128	4.340	0.996	0.402

Native entering score = query dot static key /sqrt(payload) + clock bias + query dot key_read(persistent memory)/sqrt(payload), then the existing clamp. Exact hooks reconstruct all candidate scores within $1.91e-6$. Trained memory gaps exceed static gaps in 93.75%/92.86% of races. First-event memory is zero and routing is nearly uniform. Persistent-memory reads are the observed confidence source.

Component-only entropies hold the current query fixed; they are algebraic local counterfactuals, not predictions from a modified model history. Large memory scores may express useful specialization or brittle commitment. This decomposition does not establish harmful confidence or a representation-learning regression.

Conventional router z-loss penalizes $\log \sum \exp(\text{scores})$ for numerical stability (ST-MoE, Zoph et al., 2022). Here that quantity is log total rate, so it controls actual first-arrival time. Subtracting it forces total rate 1 and changes computation. Choice uncertainty and common speed must be distinguished before importing normalization.

Theory 100, 5,376 race score pairs / 10,752 candidate scalars; four frozen encoders x 16 unused FIT prefixes x one noise history. Eight exact logit/state/end-RNG contracts; 32.992s/328,876KiB. No loss, fitting, DEV/test or causal harm claim. All weights fixed; prior fits and diagnostic work retained. Full audit FLOPs unknown.

Appendix B. Clock-preserving route calibration, seed6

Encoder	Temp.	Scope	FIT NLL	FIT accuracy %	Prefix MF est.
initial	1	all	2.779015	0.00	0.591695
initial	2	all	2.773072	0.00	0.594131
initial	4	all	2.771472	0.00	0.594131
initial	2	layer0	2.769006	0.00	0.592913
fixed_pass4	1	all	1.376994	31.25	0.591695
fixed_pass4	2	all	1.365840	34.38	0.594131
fixed_pass4	4	all	1.369226	32.81	0.594131
fixed_pass4	2	layer0	1.362028	32.81	0.592913

At each entering state, keep the original first raw time T and total rate Λ . A losing exponential residual supplies a uniform independent of T ; conditional inverse CDF changes the categorical winner to $\text{softmax}(\text{score}/\text{temperature})$. No extra random draws, losing value deliveries or receiver updates. The actual selected receiver changes, so subsequent memories and later clocks may change.

Separate16 unused FIT examples, four prespecified noisy histories; every initial/fixed-four-pass model and setting retained. This is mean per-history loss/accuracy, not ensemble inference. No optimizer, new decoder, DEV or test. Temperature2/all NLL gain 0.011155, accuracy change 3.125 points; fixed both-seed smoke gate FAILS.

Theory101. Distribution screens and exact tau1 output/state/all-gradient/RNG contracts pass29.010s/354,328KiB. Positive-temperature training is refused until correct choice/common-clock gradients and optimizer/recovery/accounting exist. Each prefix scores168 keys, writes84 receivers, has8 available units;576..720 live-state bytes. Original trained fit2.285696GF/2.232125MF per presentation; table work is one traced inference prefix including calibration, not full-audit cost or benchmark advantage.

Accounting clarification,2 October:720bytes is the observed maximum, not every prefix. Completed cases retain576,648 or720 persistent tensor bytes as6,7 or8 units become occupied. Repeated receiver writes and available capacity do not imply all available units have live memory. Parameters, Python metadata and graphs are separate. Original report retained.

Appendix B. Clock-preserving route calibration, seed7

Encoder	Temp.	Scope	FIT NLL	FIT accuracy %	Prefix MF est.
initial	1	all	2.444056	18.75	0.591695
initial	2	all	2.445819	18.75	0.594131
initial	4	all	2.451157	18.75	0.594131
initial	2	layer0	2.441686	18.75	0.592913
fixed_pass4	1	all	1.189652	43.75	0.591849
fixed_pass4	2	all	1.187395	43.75	0.594131
fixed_pass4	4	all	1.205503	43.75	0.594131
fixed_pass4	2	layer0	1.193236	43.75	0.593067

At each entering state, keep the original first raw time T and total rate Λ . A losing exponential residual supplies a uniform independent of T ; conditional inverse CDF changes the categorical winner to $\text{softmax}(\text{score}/\text{temperature})$. No extra random draws, losing value deliveries or receiver updates. The actual selected receiver changes, so subsequent memories and later clocks may change.

Separate 16 unused FIT examples, four prespecified noisy histories; every initial/fixed-four-pass model and setting retained. This is mean per-history loss/accuracy, not ensemble inference. No optimizer, new decoder, DEV or test. Temperature 2/all NLL gain 0.002257, accuracy change 0.000 points; fixed both-seed smoke gate FAILS.

Theory 101. Distribution screens and exact tau1 output/state/all-gradient/RNG contracts pass 29.010s/354,328KiB. Positive-temperature training is refused until correct choice/common-clock gradients and optimizer/recovery/accounting exist. Each prefix scores 168 keys, writes 84 receivers, has 8 available units; 576..720 live-state bytes. Original trained fit 2.285696GF/2.232125MF per presentation; table work is one traced inference prefix including calibration, not full-audit cost or benchmark advantage.

Accounting clarification, 2 October: 720bytes is the observed maximum, not every prefix. Completed cases retain 576,648 or 720 persistent tensor bytes as 6, 7 or 8 units become occupied. Repeated receiver writes and available capacity do not imply all available units have live memory. Parameters, Python metadata and graphs are separate. Original report retained.

Appendix B. AWS frozen replay critics and parameter calibration

Frozen conditional score study	Seed	k2/plain k4 MSE	Gate
Initial fine norm	7	2.083742	FAIL
Initial fine norm	8	2.081475	FAIL
Trained coarse norm	7	2.429824	FAIL
Trained coarse norm	8	2.247774	FAIL
Signed/label critic	7	2.143301	FAIL
Signed/label critic	8	2.493344	FAIL
Training-only shrinkage	7	2.149378	FAIL
Training-only shrinkage	8	2.493344	FAIL

Correct first-time-preserving actual-write replay and critics frozen before subset sampling retain unbiased credit. Every reduced-replay nomination fails. The initial fine screen uses 84 races; trained coarse uses 20. Signed-message and label-aware critic inputs are detached learning features, with 32 critic-training and 32 heldout FIT prefixes. Sixteen original numerical contracts precede screens.

Score-coordinate MSE is insufficient for shared parameters: cross-site covariance contributes. Actual training-only calibrated parameter variance ratios 2.383887/2.189517 in seed 7 and 37.462736/1.674462 in seed 8 also fail. Exhaustive finite-population and cached factual-probability contracts pass. 160 new VJPs paid.

All outcomes, targets, critics and original costs retained. No new producer fit, DEV/test quality or reduced-work learning claim. Diagnostic FLOPs unmeasured, not zero. This rejects unchanged critic reduction rather than counterfactual credit in general. AWS_REPLAY_VARIANCE_FINDINGS and AWS_REPLAY_CALIBRATION_FINDINGS carry lineage.

Appendix B. Replay allocation: fresh score gains, parameter failure

Proposal	Seed	Score variance/k4	Two parameter ratios
Ordinal k2	7	2.074656	Not measured
Ordinal k2	8	3.053016	Not measured
Feature k2	7	1.206815	Not measured
Feature k2	8	1.519582	Not measured
Weighted distinct k3	7	0.768144	0.600428/4.011873
Weighted distinct k3	8	0.693763	0.639390/1.009399

Ordinal and feature-conditioned k2 proposals sample WITH replacement, using exact1/(kp) correction and a positive floor. Baseline uniform k4 is WITHOUT replacement; uniform k2 with replacement is 2.375x. Feature-conditioned priorities improve over that reference but fail reduced-budget gates. Diagnostic exact-return oracle ratios .631115/.519480 need all costly utilities and are not deployable.

The frozen feature predictor then selects three distinct weighted sites on fresh FIT64..95, with exact marginal/pair inclusion probabilities and 1/inclusion correction. Conditional score variance falls 23.19%/30.62% versus uniform four; uniform three would increase it 41.67%. Preserve this positive allocation result. The combined gate fails on actual parameter cases shown above.

Exact expectation/shared-vector variance/uniform nesting contracts pass. 2560 actual diagnostic lanes and 80 confirmation VJPs; distinct screen 3.578s/ 518,380KiB. Inclusion computation .742/.747ms per prefix, plus tree/discovery/VJP costs. Proposed 6 versus 8 execution lanes is projected, not measured reduced-work training. Full diagnostic FLOPs unknown; no quality/supremacy claim or unchanged fit.

Appendix B. AWS compact context head: practical controls remain stronger

Predictor	Dev accuracy %	Dev NLL	Combined fit GF	Fit MF/target
s6_initial	58.85	1.268551	Unknown	Unknown
s6_fitted	67.71	0.901193	Unknown	Unknown
s7_initial	62.50	1.145877	Unknown	Unknown
s7_fitted	64.58	0.938003	Unknown	Unknown
s8_initial	55.73	1.170969	Unknown	Unknown
s8_fitted	68.75	0.915083	Unknown	Unknown
raw4	65.62	0.888167	Unknown	Unknown
raw20	66.15	0.902387	Unknown	Unknown

33-anchor dense local RBF readout uses all distances/basis outputs on frozen native query contexts, plus mandatory raw4/raw20 prototype controls. 984 FIT/ 192 DEV, three native trained/initial pairs, fitting-user three-fold regularization selection. 72 fold fits, 8 refits and 352 class-specific KMeans fits are paid. Six Torch and 48 query/state contracts pass before decoder fitting.

Both fixed stage/storage gates FAIL. Learned-over-initial mean NLL gain .277040 is supported. Native standalone exports ~115KB versus raw4 40.966KB and raw20 184.334KB; fine-control byte savings do not overcome stronger coarse quality/storage. Raw4 SVC 77.604%/.686661 and earlier valid leaders remain.

Whole/per-target combined fitting and inference FLOPs unknown for EVERY arm because solver/replay work is unmeasured. Trained native parent fit 4.218015GF/ .535825MF per 7872

presentations retained, not total fit. Initial/raw producer fit0 is not zero solver work. Study23.296s/690,896KiB. Portable prediction/roundtrip and three sequential latency repeats pass; exports include normalization/core/ anchors/decoder metadata. This is a dense readout diagnostic, not sparse race attention.

Appendix B. Integrated exact replay: first seed does not improve quality

Native seed7 method	Dev acc %	Dev NLL	Whole fit GF	Fit MF/presentation	Infer MF/prefix
Original route teacher	57.81	1.105326	20.074	2.550	0.591709
Factorized clock only	56.77	1.116189	18.244	2.318	0.591737
All-race exact replay	58.33	1.141794	1040.397	132.164	0.591779

Same984 fitting/192 development inputs, eight passes/7872 fitting presentations, p16/layer2/head2/pool2. Every method retains computational delays, temporal races, persistent sparse writes and separate keys/values. All-race replay forces alternatives at the factual first time with real writes; vectorized shadow histories include unrealized downstream effects. Factorized-clock-only removes the original local route teacher.

All-race replay58.333%/1.141794 versus factorized56.771%/1.116189 and original57.812%/1.105326 is not a quality improvement under the prespecified loss criterion. Approximately1040GF counted fit versus18.24GF factorized is paid despite batched CPU execution. A large fit/dev gap motivates independent generalization diagnostics; it does not prove routing credit cannot help.

Whole fitting and per-presentation fitting columns share units and7872 denominator. Inference mean first11 DEV prefixes;2FLOPs/MAC plus unit specials. Eight available receivers,168 scored keys/84 writes per21-event prefix. Candidate discovery, replay, backward and Adam charged; preprocessing, traffic, RNG and energy separate. Source/class/estimator variants and one-seed exploratory scope retained. Other host owns remaining replay/regularization comparisons.

Appendix B. Native conditional choice-clock parameter geometry

Seed	Encoder	Median	cos	Min info ratio	Max info ratio	Capped sites
6	initial	0.7708	2.17e-05	4.24e-01	0/8	
6	fixed_pass4	0.9850	1.68e-07	1.33e-01	0/8	
7	initial	0.7826	4.47e-02	4.91e-01	0/8	
7	fixed_pass4	1.0000	4.12e-21	3.58e-01	5/8	

At a fixed observed earlier winner/time history, differentiate entering score contrast $u=\text{grad}(s_0-s_1)$ and common log-rate $v=\text{grad}(\log\text{sumexp}(s))$. The local winner/time metric is $\pi_0\pi_1u*u^T+v*v^T$; its two eigenvalues come from a 2x2 Gram matrix, avoiding a dense15523-coordinate matrix. Past physical times are held fixed for conditional likelihood geometry; real earlier message/memory maps and sparse writes still receive derivatives.


Four prespecified sites, two unused FIT inputs258/981, four frozen initial/ fixed-four-pass encoders:32 Jacobian pairs/64 VJPs. Native primal outputs, ALLstate and end-RNG reproduce exactly. Four-block direction arrays retain clock-bias, key/query, message/state/transport and classifier contributions. Directional finite differences use eps .002/.001, relative2%/absolute.005 tolerance; every error is retained rather than reported as exact algebra.

Five of sixteen selected trained cases pin a score at12 and leave local choice/clock sensitivities collinear. Other confident trained cases are poorly conditioned without clipping. Negative cosine alone is not harmful credit: no target utility or proposed update is evaluated. Single-site bias control also does not imply independent control over all shared histories.

Theory103;36.536s/348,112KiB. Exact originals/weights restored, no optimizer/DEV/test or replaced time-learning path. Original core fit2.285696GF each /2.232125MF per1024 presentations retained. Conditional diagnostic is not full expected-risk credit, parameter covariance or advantage. Full audit FLOPs, traffic and energy unknown, not zero.

Appendix B. Bounded emitter bridge restores the missing direction

FIT index	Site/head0	Raw max	Min slope	Hard info ratio	Bridge info ratio
258	e19/L0	13.9539	0.02772	5.48e-20	1.45e-06
258	e20/L0	15.0771	0.02415	8.31e-20	1.70e-06
258	e20/L1	15.9095	0.02184	4.12e-21	1.92e-05
981	e19/L0	14.6482	0.02545	5.35e-20	2.09e-06
981	e20/L0	17.8833	0.01730	9.25e-20	3.28e-06



local score bound evidence 20261002T235800Z geometry

Fixed alpha.1 bridge: $.9*\text{clip}(\text{raw}, -12, 12) + .1*\text{raw}/\sqrt{1+(\text{raw}/12)^2}$. It remains bounded, odd and monotone, with positive mathematical sensitivity beyond the hard clamp. On these same conditioned histories it restores the missing local direction. Ratios remain small because choices are rare; this is not measured noise reduction, useful-route credit or model-quality improvement. The hard-map near-zero ratios are numerical roundoff around theoretical rank1, not meaningful residual information.

Theory104; five primitive contracts and eight native comparisons pass 37.961s/371,992KiB. Alpha0 outputs/state/RNG and all parameter gradients exactly nest the old model. Positive complete native inference retains coupled clocks/ routes/real writes and target invariance; training refuses. Each prefix scores 168 logical keys/writes84; diagnostic hooks RECOMPUTE168 extra dot products, paid work. Live state576..720bytes,8available units. Prior fits paid; audit FLOPs unknown. No automatic fit, free scoring or hardware/quality advantage.

Appendix B. Completed fine-packet replay confirmations

Seed	Credit/depth	Dev acc %	Dev NLL	Whole fit GF	Fit MF/presentation	Infer MF/prefix
6	L2/all	61.46	1.119163	1040.397	132.164	0.591737
7	L2/all	58.33	1.141794	1040.397	132.164	0.591779
8	L2/all	66.67	0.941821	1040.397	132.164	0.591849
7	L4/k8	55.73	1.330390	214.394	27.235	1.068297
8	L4/k8	57.29	1.101878	214.394	27.235	1.068297

Full984 FIT/192 DEV,20 observed packets plus query,8passes/7872 fitting presentations,U16/lr.003. All-race depth2 now has three completed seeds; depth4 sampled8 has seeds7/8. Original/factorized controls remain on the preceding replay ledger; missing matched deeper controls are not filled with predictions. Seed8 depth2 reaches66.667%/.941821; preserve it as a completed result with its1040.397GF fitting cost, not an isolated practical advantage.

Depth4 sampled8 is worse than same-seed depth2 on development loss in these runs. This is a restricted eight-site credit variant, not a test of all-race deeper credit or proof that depth cannot learn. The owner has queued matched depth4 all-race, depth6 sampled and growth-by-nesting comparisons.

Every column uses consistent units/denominators; core whole-fit and per-presentation work include replay, backward and Adam. Inference mean first11 DEV prefixes,2FLOPs/MAC plus unit specials. L2:8available units/168key scores/ 84writes per prefix;L4:16/336/168. FIT diagnostic scores use only32 fitting examples, not whole-FIT risk. DEV epoch selection and one-seed comparisons remain exploratory. Raw preprocessing, RNG, traffic and energy separate.

Appendix B. AWS deeper all-race pilot: matched gates fail

Coarse L4 seed7	Dev acc %	Dev NLL	Whole fit GF	Fit MF/presentation	Infer MF/prefix
Original teacher	56.77	1.342018	1.039740	1.015371	0.250529
Factorized control	55.73	1.368549	1.082407	1.057038	0.250487
All40-race replay	56.77	1.374384	28.968062	28.289123	0.250473

Coarse4/.25-clock,p16/L4/H2/pool2,256FIT/192DEV,fourpasses/1024 presentations,U16/lr.003. Every-parameter sequential/forked replay contracts and three actual recovery/accounting smokes precede these pilots. Full40 race sites/80shadow lanes preserve factual first time and actual forced writes. Teacher/factorized/replay share data and observation protocol; no weight decay or input noise is mixed into this comparison.

Replay NLL is worse by.032366 versus teacher and.005835 versus factorized, at27.861x/26.763x counted fitting work. Both fixed gates fail, confirmation is absent and no larger fit is nominated. This is a first small coarse/deep pilot; it does not replace the independently owned full-data fine-depth tests.

Completed summary234200Z,108.202s total campaign. Same common-unit denominator in every fitting column.16available receivers,8writes/event, 40factual races/prefix. Full replay/optimizer paid despite parallel wall execution. Inference first11 DEV prefixes; raw input preprocessing/traffic/ RNG/energy separate. No official test, comparable-quality work supremacy or claim that all deep counterfactual learning must fail.

Appendix B. Parameter-targeted allocation retains a scoped positive gain

Seed	Score variance/k4	Parameter FIT96 ratio	Parameter FIT97 ratio	Each-case gate
7	0.906188	0.599166	0.486158	True
8	0.958262	0.614246	1.137133	False

Frozen priority predictor trains on true shared-parameter route norms on FIT0..31 and confirms on new FIT96..127. Exact distinct weighted sampling and inclusion correction remain. Both aggregate score ratios pass; one seed8 parameter case fails. Combined nomination stays FAIL. The two-case weighted parameter aggregates .516746/.658755 support 48.3%/34.1% reductions and are preserved as diagnostics, not replacements of the every-case gate.

This is materially better allocation evidence than isolated score-space improvement, but four parameter examples cannot establish robustness or quality. No unchanged reduced-replay fit follows. Producer FIT labels were already seen; only priority supervision is held out. Every expensive training utility remains paid.

11.439s/531,560KiB;1280 training plus 80 confirmation VJPs, 2560 diagnostic shadow lanes, saved predictors/proposals. Whole fitting/diagnostic FLOPs unknown, not zero. No inference change, DEV/test quality or 6-versus 8 execution-lane work advantage. Main deeper replay/growth remains independently owned.

Appendix B. AWS private state and shared maps: pilot gates fail

Variant	Params	Available	Dev acc %	Dev NLL	Fit GF	Fit MF/target	Infer MF
Untied pool2	15523	8	55.73	1.274860	0.548517	0.535661	0.138133
Shared pool2	11227	8	53.12	1.276615	0.544113	0.531360	0.138105
Shared pool8	12019	32	54.69	1.233413	1.554905	1.518462	0.204447

Same 256 FIT/192 DEV, four passes/1024 presentations, seed 6, coarse 4/.25-clock, p16/L2/H2. Shared input/output/gate maps retain private keys, clock biases, timescales and receiver memories. Delays perform computation, hard addressed races select four real writes/event and separate keys/values remain. Pool8 adds 32 available receivers while keeping four selected writes.

Both fixed quality/resource nominations FAIL: tied2 loses .001754 NLL and 2.604 accuracy points; tied8 improves NLL .041447 but misses .05 and loses 1.042 points. Lower parameters and positive tied8 loss improvement remain supported. Pool growth also changes total hazard and per-unit exposure, so this is not isolated proof of capacity beyond scored work. No confirmation/full fit.

Every-parameter/state/batch and actual Adam/cursor/RNG/operator contracts pass before both learning smokes and pilots. Historical failed coalesce-adapter contract remains beside corrected admission. Campaign 56.762s; all candidate keys/proposals/losing credit/optimizer work paid, same fitting denominator. Inference first 11 prefixes/2 FLOPs per MAC plus unit specials; raw preprocessing, physical traffic/RNG/energy separate. Deeper replay owns the next integrated comparison; no unchanged shallow expansion.

Appendix B. Restored native sensitivity has small, mixed route utility

FIT	Event/layer/head	Raw capped score	L0 minus L1	Hard emitter credit	Bridge emitter credit
258	19/0/0	13.9539	+0.003911	0.00e+00	+9.07e-10
258	20/0/0	15.0771	-0.001274	0.00e+00	-2.52e-10
258	20/1/0	15.9095	-0.019186	0.00e+00	-6.14e-11
981	19/0/0	14.6482	-0.007041	0.00e+00	-1.57e-09
981	20/0/0	17.8833	-0.037535	0.00e+00	-5.36e-09

Five prespecified strictly saturated seed7 histories from theory104, fixed-pass4 producer, FIT258/981. For each, both actual alternatives write persistent receiver state at the factual FIRST arrival. Prefix winners/times are fixed; suffix clocks, messages and memory evolve with the original hard map. Original-winner replays recover all original logits/state/end RNG. Ten alternative plus five factual native forwards are paid;168additional raw dot products per diagnostic forward. These FIT labels are revealed for utility diagnosis and are no longer label-unused. No optimizer/DEV quality/test selection.

L0-L1 is positive in one case, so its alternative improves NLL; four negative gaps mean the factual winner is better. Every previously blocked emitter has nonzero bridge choice sensitivity, confirmed by double analytic derivatives and finite-difference gradcheck. Absolute restored derivatives are only6.14e-11 to5.36e-9: alternative probabilities remain1.46e-7 to 8.77e-6. Restoring rank does not by itself make those routes useful or their contribution to expected risk substantial.

Conditional choice risk is $\sum p_i \cdot L_i$, so $dR/draw0 = \pi_0 \cdot \pi_1 \cdot (L_0 - L_1) \cdot f'(raw0)$. Utilities and first time are held fixed; this tests choice credit only. The bridge is applied to the current conditional score map in this derivative, not the suffix law. These numbers are not a full positive-bridge risk, joint time/content update or learning result.

Audit18.432s/328716KiB; diagnostic whole arithmetic/traffic/energy unknown, not zero. Conditional two-score Fisher determinant is $\pi_0 \cdot \pi_1 \cdot d_0^2 \cdot d_1^2$ times the raw-gradient Gram determinant. Positive slopes preserve existing raw rank; confidence and nearly parallel raw directions can still suppress useful learning. Theory106/107 retain exact scopes. No automatic bridge scale-up; independently owned integrated deeper replay and reserved AWS comparisons remain the quality priority.

Appendix B. AWS depth8 replay: seed7 positive pilot

Family/credit	Params	Dev acc %	Dev NLL	Whole fit GF	Fit MF/target	Infer MF
private/teacher	54571	48.438	1.369199	2.022431	1.975030	0.475853
private/factorized	54571	47.396	1.399592	2.124006	2.074225	0.475713
private/replay	54571	56.771	1.316790	109.118229	106.560770	0.475769
depth/teacher	22351	52.604	1.410495	1.989397	1.942771	0.475699
depth/factorized	22351	46.875	1.395239	2.090973	2.041966	0.475909
depth/replay	22351	52.083	1.348356	109.085195	106.528511	0.475923

Both families retain native computational delays/races, key/value separation, source carry, sparse writes and private persistent receiver state. Private maps differ perdepth/receiver; depth-shared maps are shared acrossdepth/pool perhead, with private keys/clocks/timescales.

p16/L8/H2/pool2,256FIT/192DEV,fourpasses/1024presentations,U16/lr.003,coarse4/.25clock.32available units,32key scores/16writes per event; five-event prefix80factual races/160full shadow lanes. All163840 shadow lanes and819200 replay events perfit are charged.

Both seed7 families pass the predeclared .03NLL gain versus BOTH matched controls and at most1pp accuracy loss. Private replay improves NLL .052409/.082802 versus teacher/factorized and

accuracy8.333/9.375pp. Depth-shared replay improves .062138/.046883NLL; teacher accuracy declines .521pp within the gate. These positive heldout results are preserved. The unchanged seed8 confirmation below subsequently fails both family gates: nomination is not robust, and no full984FIT promotion follows.

Original teacher, factorized clock/content control and corrected all-race first-time-preserving actual-write replay share data/order/noise/ optimizer budgets. Full every-parameter/recovery/state/shared-alias contracts and allsix learning smokes pass before pilot/confirmation. Replay costs about52–55x the controls in complete fitting arithmetic; equal sparse inference does not remove that expense. No comparable-quality total-resource advantage, officialtest or unchanged fullfit is claimed.

Same fitting denominator and units for every row/column; 2FLOPs/MAC plus unit specials, first11 DEV prefix inference mean. Reused DEV epoch selection, two seeds and fixed nomination scope. Input preprocessing, candidate index/RNG, traffic and measured energy separate. Verbatim shared report evidence preserved in appendices/ aws_depth8_history_20261003T003500Z.md. Source/result hashes checked; historical depth4 failed gate and local bridge evidence retained.

Appendix B. AWS depth8 replay: seed8 confirmation fails

Family/credit	Params	Dev acc %	Dev NLL	Whole fit GF	Fit MF/target	Infer MF
private/teacher	54571	55.729	1.331247	2.022431	1.975030	0.475909
private/factorized	54571	48.438	1.270459	2.124006	2.074225	0.475993
private/replay	54571	54.167	1.282510	109.118229	106.560770	0.475699
depth/teacher	22351	48.438	1.283150	1.989397	1.942771	0.475825
depth/factorized	22351	50.000	1.271398	2.090973	2.041966	0.475783
depth/replay	22351	48.958	1.270943	109.085195	106.528511	0.475811

Both families retain native computational delays/races, key/value separation, source carry, sparse writes and private persistent receiver state. Private maps differ perdepth/receiver; depth-shared maps are shared acrossdepth/pool perhead, with private keys/clocks/timescales.

p16/L8/H2/pool2,256FIT/192DEV,fourpasses/1024presentations,U16/lr.003, coarse4/.25clock.32available units,32key scores/16writes per event; five-event prefix80factual races/160full shadow lanes. All163840 shadow lanes and819200 replay events perfit are charged.

Both unchanged seed8 confirmation gates FAIL. Private replay preserves .048736NLL improvement over teacher, but loses1.5625pp accuracy and is .012051NLL worse than factorized. Depth-shared replay improves only .012207/.000455NLL versus teacher/factorized, below .03, and declines 1.0417pp versus factorized accuracy. Keep the earlier positive seed7 result beside this failure; neither establishes broad failure of deep representation learning or practical advantage.

Original teacher, factorized clock/content control and corrected all-race first-time-preserving actual-write replay share data/order/noise/ optimizer budgets. Full every-parameter/recovery/state/shared-alias contracts and allsix learning smokes pass before pilot/confirmation. Replay costs about52–55x the controls in complete fitting arithmetic; equal sparse inference does not remove that expense. No comparable-quality total-resource advantage, officialtest or unchanged fullfit is claimed.

Same fitting denominator and units for every row/column; 2FLOPs/MAC plus unit specials, first11 DEV prefix inference mean. Reused DEV epoch selection, two seeds and fixed nomination scope. Input preprocessing, candidate index/RNG, traffic and measured energy separate. Verbatim shared report evidence preserved in appendices/ aws_depth8_history_20261003T003500Z.md. Source/result hashes checked; historical depth4 failed gate and local bridge evidence retained.

Appendix B. Persistent representation and missing delayed credit

Credit horizon H	Geometric weighting after H (%)	Absolute geometric tail
16	85.214	85.641161
32	72.615	72.978583
64	52.729	52.993328
128	27.804	27.942980

Exact constructed encoder: observe x , write θx , transport by ρ^{32} with $\rho = \exp(-.01)$, predict delayed label x . At $\theta.3$, full encoder gradient $-.567961$, credit16 with detached future state0; $\theta.31$ reduces loss.305883 to.300230. Race writes x or0; delayed conditional choice gradient $-.114477$, within-chunk return gradient0 despite replaying both choices. Seven analytic/autograd/finite-difference and exhaustive-sampling contracts pass. This proves a training-credit obstruction with representational capacity retained, not a measured text8 failure.

The linear contraction-mode tail fraction ρ^H is large for a 100event decay. Native whole-state contraction is NOT established: fixed-input unit damping/orthogonal rotation do not bound learned source carry, gates, clock history and addressed interactions. Shared parameters can get other later derivatives; that is not automatic recovery of the omitted original-write path. Decoder adaptation is also not proof of learning new long-range representations.

Fixed execution-count alternative: fullT16/L8/H2/pool2 has512 shadow events/target plus one factual; T64/k8 uniform distinct sites has16shadow events/target plus one factual. The longer factual graph and score discovery still cost work. Horvitz scaling $R/k=128$ makes single-informative-route covariance $127 * g * g^T$. Exhaustive12race/k3 subsets independently verify mean and covariance. Unbiasedness can coexist with poor learning; these counts are not FLOPs/wall/quality advantage.

Theory110; .529s/282228KiB, no optimizer/DEV/test. Corrected within-chunk routing support, confidence/sensitivity and credit horizon are separate constraints. Next use exact-source saved language checkpoints for a frozen longer-suffix shared-parameter return/variance audit on unseen FIT inputs before selecting H/k. Existing10M teacher runs remain unchanged; no long sampled-credit fit or predicted benchmark gain is admitted by this synthetic witness.

Appendix B. Native text8 horizon return: four credit reversals

FIT start	Event/L/H	Q0-Q1 H16	Q0-Q1 H32	Late contrast	Param norm16	Param norm32	Flip
8192	0/0/0	+0.119169	-0.009301	-0.128470	0.003241	0.000253	YES
8192	0/7/0	+0.045263	+0.020573	-0.024690	0.026251	0.011932	No
8192	8/0/0	+0.164192	-0.112556	-0.276748	0.043189	0.029606	YES
8192	8/7/0	+0.009205	+0.007439	-0.001766	0.000520	0.000420	No
8320	0/0/0	+0.285896	+0.381399	+0.095503	0.000551	0.000735	No
8320	0/7/0	-0.018730	-0.022666	-0.003936	0.000260	0.000315	No
8320	8/0/0	-0.040859	+0.151710	+0.192568	0.007922	0.029413	YES
8320	8/7/0	-0.009103	+0.004196	+0.013299	0.007070	0.003259	YES

Frozen private native p16/L8/H2/pool2,54907parameters,original 8192-character/fourpass seed6 producer. Fixed final-pass4 ONLINE weights, not DEV-selected state, converted to double for numerical VJPs. Two producer-unseen FIT spans8192/8320,16observed warmup tokens and32 causal next-token targets each. Prespecified sites(event0/8,layer0/7,head0), one fixed common-noise draw. Both receivers actually write private state at the factual FIRST time, then messages/clocks/memory

evolve through all32targets. Original winner recovers ALLfactual logits/state/end RNG.

Four of eight sites reverse the conditional choice-credit sign when losses16..31 are included. Four-site aggregate native parameter-credit cosines -.600832/-.792626 show directional opposition on these contexts; norms .051372/.032006 and.010178/.029475 for16/32returns. None of these chosen scores is capped. This is real omitted downstream utility at this frozen point, beyond theory110 synthetic witnesses, not proof that truncation dominates global learning or longer credit improves quality.

Per-site credit is $\pi_0 \cdot \pi_1 \cdot (Q_0 - Q_1) \cdot \text{grad}(\text{scores}_0 - \text{scores}_1)$, with the SAMErealized-branch native score Jacobian for both horizons. Earlier pathwise clocks remain differentiable; this is not note103 fixed-clock likelihood geometry. The shown contrasts are SUMsuffix NLL, not average prediction quality. Future observed tokens/changed target labels leave earlier factual output unchanged. Allcases/full54907coordinate arrays saved; sampling only four of512sites excludes a full-chunk gradient claim.

Theory111;7.026s/426172KiB.16forced32token suffixes/two factual graphs/eightVJPs plus warmup/causality checks paid; diagnostic totalFLOPs unknown, not zero. Double live state4488..4624bytes,32available/32scores/ 16writes pertoken. No optimizer/DEV/test/changed-parent fit or benchmark advantage. The independent confirmation below FAILS aggregate-opposition replication; retain this original positive mechanism observation beside it.

Appendix B. Horizon confirmation: downstream effects vary by context

Fresh FIT	Draws	Site flips	Mean norm16	Mean norm32	Mean cosine	Opposition gate
8448	3	2/12	0.303282	0.536700	+0.999383	False
8576	3	1/12	0.049228	0.021923	+0.455223	False

Same frozen producer/sites/budgets,NEW FIT8448/8576 spans and three independent evaluated race draws111330/111331/111332. Warmup seed111329 stays fixed across draws.24cases/48actual forced suffix replays,24parameter VJPs,all outcomes and coordinate arrays retained. Before observing results, diagnostic replication required both span mean16/32parameter-credit cosines<0 and at least one site flip per span. Both gates FAIL: mean directions agree,despite three of24 individual site reversals. Do not replace this gate with a larger-gradient criterion.

First fresh span preserves direction across all draws,with mean longer-return norm.536700 versus.303282. Second has more variable perdraw norms/angles and mean cosine+.455223. These are supported context-dependent downstream effects; universal destructive short credit is not established. Neither frozen study measures expected fullstream risk, semantic feature content or a fitted benchmark gain. Original first-audit reversals remain valid within their scope.

The current mechanism priorities stay separate: corrected actual language shadow return is numerically installed; hard-score sensitivity repair has no demonstrated quality gain; full replay work is expensive; horizon changes utility but robust allocation/generalization remain open. Next broader frozen utility/covariance accounting can inform H/k; no unchanged long sampled-horizon training follows this failed diagnostic replication. Existing10M original-teacher controls continue independently.

Theory112;20.355s/470848KiB,one guarded CPU job/no optimizer/ DEV/test. Earlier-loss/factual-first-time/ALLoriginal-winner-state/RNG/ future-token/target invariance and unchanged parameters pass on all three draws. Same original fit work paid; whole audit arithmetic, physical traffic and energy unknown. More counterfactual support cannot automatically repair omitted future loss, while a longer return is also not automatically a better finite-data update. No supremacy claim.

Appendix B. Deep fitting diagnosis: looser clipping fails its prediction

Depth/credit/cap	Seed	FIT32 NLL	DEV NLL	DEV %	Whole fit GF	Fit MF/target	Infer MF/target
D2 factorized	7	0.576049	1.116189	56.771	18.244	2.318	0.592
D2 clip4	7	0.683112	1.150943	56.250	18.244	2.318	0.592
D4 factorized	7	0.718192	1.130272	55.208	34.348	4.363	1.068
D4 clip4	7	0.945086	1.172124	57.292	34.348	4.363	1.068
D4 all-race	7	0.763093	1.152624	55.208	3815.207	484.655	1.068
D4 factorized	8	0.587894	1.093247	59.896	34.348	4.363	1.068
D4 clip4	8	0.711417	1.035432	58.854	34.348	4.363	1.068
D4 all-race	8	0.553204	1.020627	64.583	3815.207	484.655	1.068
D4 lr.006	7	0.959557	1.325388	54.167	34.348	4.363	1.068

Same p16/H2/pool2 fine20 temporal packets,984 FIT/192 subject-disjoint DEV gestures, eight passes/7872 fitting presentations/496 Adam updates, U16/lr.003 except marked lr.006. Factorized/all-race cap1, clip4 changes only the cap. All rows use the SAME whole-fit and per-target denominators; work includes all candidates/shadows/backward/normalization/clipping/Adam, first/last-window estimates. No official test or measured energy.

Correction beside historical §408: clip4 WORSENS the saved FIT32 NLL in BOTH seeds; development accuracy/NLL change in mixed directions. The simple fitting-repair prediction fails on this diagnostic. FIT32 is the first32 fitting gestures at DEV-selected weights, not whole-training loss or a common final epoch. Newly completed D4 lr.006 also worsens FIT32 to.960, and D2 clip4 to.683 versus.576. These rule against the tested larger-step repairs, not every functional-step or optimizer-history explanation.

All-race credit gives a positive seed8 result: DEV accuracy4.6875 points higher and NLL.072620 lower, with better FIT32. Seed7 fails to improve. Thus the old categorical statement that routing credit cannot explain deep fitting is too strong. Sampled8 is a distinct noisy estimator: with21events D4 has168 races, each sampled contribution scaled21; D6 scales31.5. Full replay costs roughly 111 times factorized fitting here; useful occasional quality is not an efficient-learning advantage.

Three independent audits: completed evidence/protocol, optimizer mechanics, and architecture/transport. Theory122 records the ranked hypotheses, primary-paper analogies and discriminating tests. Small elementary transport decay does not establish good complete recurrent/message Jacobian conditioning; large norms do not establish useful features.

Appendix B. Deep race sampling: preserve mixed quality and complete work

Model/credit	Seed	FIT32 NLL	DEV NLL	DEV %	Whole fit GF	Fit MF/target	Infer MF/target
D4 factor	7	0.718192	1.130272	55.208	34.348	4.363	1.068
D4 sampled8	7	0.715109	1.330390	55.729	214.394	27.235	1.068
D4 full	7	0.763093	1.152624	55.208	3815.207	484.655	1.068
D4 factor	8	0.587894	1.093247	59.896	34.348	4.363	1.068
D4 sampled8	8	0.630686	1.101878	57.292	214.394	27.235	1.068
D4 full	8	0.553204	1.020627	64.583	3815.207	484.655	1.068
D6 factor	7	0.658653	1.085895	59.375	50.721	6.443	1.545
D6 sampled8	7	0.952533	1.336483	47.917	313.463	39.820	1.545

Completed native p16/H2/pool2 fine20-packet models;984 FIT/192 subject-disjoint DEV gestures, eight passes/7872 presentations/496 Adam updates, U16/lr.003/clip1. All columns use identical units and target denominators. Work includes discovery, losing-value replay, backward and optimizer; first/last-window estimates, not energy. FIT32 is the first32 fitting examples at DEV-selected weights.

D6 factorized learns and improves over D4 seed7, while D6 sampled8 loses11.458 accuracy points and worsens both saved losses. Thus failure of this sampled-credit variant is not a universal depth-capacity failure. D4 full credit improves seed8 by4.6875 accuracy points and.072620 NLL, but seed7 fails. Preserve this positive result with replication limits and roughly111-fold factorized fitting cost. No practical superiority established.

With21 events, D2/D4/D6 have84/168/252 races per episode. Sampling eight multiplies every chosen contribution by10.5/21/31.5. That unbiased rescaling can increase variance; these counts alone do not quantify interference after the parameter Jacobian, global clipping or Adam. Theory124 computes that conditional parameter covariance directly, without a new quality fit.

Common versus independent race noise is a different question. The existing trained D2 covariance audit fails both improvement gates, with ratios near one. Earlier priority experiments also show score-variance improvements can worsen parameter variance. Neither a new noise policy nor priority-allocation fit is nominated from general variance intuition.

Sparse addressed persistent state, computational clocks and hard routes remain in every row. Counterfactual losses teach route choice through detached returns; they do not directly backpropagate through losing message contents. Language credit horizons and streaming/reset-segment protocols are separate from these full DVS episode graphs. No official test or pending result used.

Appendix B. Live counterfactual messages: conditional content and route credit

The prior replay learns route utility from detached alternative losses, while the realized winner supplies content gradients. This need not be biased, but losing message functions receive no direct payload derivative on that realization. The new training-only sibling differentiates through every alternative of ONE sampled race, including complete live prefixes and private writes.

At a uniformly sampled legal race r , $\text{objective} = \sum \text{stopgrad}(\pi_i) * L_i + R * \sum \pi_i * \text{stopgrad}(L_i)$, averaged over actual examples. The branch-content average REPLACES factual loss gradient; adding both would double-count. Only sampled choice credit gets legal-race scaling R . Native computational delays, factorized first-time gradients, sparse private state, key/value separation and inference remain unchanged.

For ONE episode at fixed entering history/first time and independent future draws, the branch average is the conditional mean of the native pathwise estimator. It removes current-winner variance at that site. Conditioning on the entire candidate noise would reveal the winner. Shared noise across episodes means this is not a lower BATCH-variance theorem. Global clipping and Adam are nonlinear; their expected updates need not be preserved.

Native one-event losing output map	Gradient L2
Factual payload gradient	0
Joint conditional gradient	0.0116820573
Loser selection probability	0.481761694

Eight contract groups cover EVERY parameter against explicit live branches/decomposition, independent pure-Torch factorized clock algebra, factual-forced-winner identity, actual first-time preservation, existing BLk1 choice credit, relabeling-invariant factual predictions, real normalized clip1 Adam, byte-serialized next-update recovery and complete operator coverage. Early/late and unequal-length native episodes are included.

The independent reference holds each recorded branch topology and clock latent $Z = \text{Lambda} * T$ fixed, then differentiates $T = Z / \text{Lambda}$. It does not differentiate a fixed-candidate-noise minimum or detach physical time. Future hard-choice boundaries remain scoped native derivatives; this is not exact whole-risk differentiation.

Theory129;2.644s/343656KiB. Known conditional averaging and stochastic-computation-graph primitives (Schulman et al.2015); the native coupled-clock/content application is implemented and contracted, not a priority claim. Source/kernel/RNG preserved; total contract-campaign work unknown, not zero.